






K - Nearest Neighbor Method – An Instance based method

kNN






The Nearest Neighbor Method



- ***Observation:*** similar examples tend to belong to the same class.
- ***Idea:*** assign a test input with the class label from the example in the training data set that is the most similar to the test input.
执行
- ***Implementation:*** calculate the distances between the test input and the examples in the data set and find the one with the shortest distance.

The K - Nearest Neighbor Method



- **Observation:** similar examples tend to belong to the same class.
- **Idea:** assign a test input with the dominating class label from the examples in the training data set that are within the neighborhood of the test input.
- **Implementation:** calculate the distances between the test input and the examples in the training data set and find the K ones that are the most closest to the test input and find the dominating class label within these K neighboring examples.


Histogram density model

In the case of a single continuous variable x . Given N sample values from x .

- Partition x into distinct bins of width Δ_i
- Count the number of observations of x falling in bin i as n_i
- Probability density value for bin i is given by

$$p_i = \frac{n_i}{N\Delta_i}$$

Hence, we have, $\sum_{i=1}^m p_i \Delta_i = 1$



- Let's suppose we have a data set comprising N_k points in class \mathcal{C}_k with N points in total so that $\sum_{k=1}^C N_k = N$.
- If we wish to classify a new point x , we draw a sphere centered on x containing precisely K points irrespective of their class. Suppose this sphere has volume V and contains K_k points from class \mathcal{C}_k .
- We want to classify the new point x to class \mathcal{C}_k that has the highest ***posterior probability of class membership:***

$$p(\mathcal{C}_k|x) = \frac{p(x|\mathcal{C}_k)p(\mathcal{C}_k)}{p(x)}$$

- Then the likelihood can be estimated as:

$$p(x|\mathcal{C}_k) = \frac{K_k}{N_k V}$$

- The unconditional density can be estimated as:

$$p(x) = \frac{K}{NV}$$

- The class prior can be calculated as:

$$p(\mathcal{C}_k) = \frac{N_k}{N}$$

- Applying **Bayes' theorem** gives the **posterior probability of class membership**:

$$p(\mathcal{C}_k | \mathbf{x}) = \frac{p(\mathbf{x} | \mathcal{C}_k) p(\mathcal{C}_k)}{p(\mathbf{x})} = \frac{K_k}{K}$$

- To minimize the probability of misclassification, we assign the test point \mathbf{x} to the class having the largest posterior probability, i.e., the largest value of K_k/K
- **To classify a new point, we identify the K nearest points from the training data set and then assign the new point to the class having the greatest number of representatives amongst the K nearest neighbor set.**



The KNN Classification Algorithm

Begin

load the training data set $\{x_i, y_i\}$, and the test data point x

Choose the value of K , set $\mathcal{D} = \{\}$

For $i = 1$ to N

 calculate $d_i = \|x - x_i\|$

 append $\{y_i, d_i\}$ to \mathcal{D}

End


sort \mathcal{D} in the ascend order with respect to d value

pick the first K entries $\{y_n, d_n\}, n = 1, \dots, K$ from \mathcal{D}

$y = mode(\{y_n, d_n\}, n = 1, \dots, K)$

Return y

End



Metrics used in the KNN method

- ***Distance functions*** are used in KNN method as measure of similarity of instances
- Commonly used distance functions in machine learning include the following:

Let $\mathbf{x} = [x_1 \quad \cdots \quad x_m]^T$ and $\mathbf{y} = [y_1 \quad \cdots \quad y_m]^T$ be two vectors, then the distances between \mathbf{x} and \mathbf{y} are:

- *Euclidean distance*

$$d(\mathbf{x}, \mathbf{y}) = \left(\sum_{i=1}^m (x_i - y_i)^2 \right)^{\frac{1}{2}} \quad L_2 \text{ norm}$$

- *Minkowsky distance*

$$d(\mathbf{x}, \mathbf{y}) = \left(\sum_{i=1}^m |x_i - y_i|^p \right)^{\frac{1}{p}} \quad L_p \text{ norm}$$

- *Manhattan* distance

$$d(\mathbf{x}, \mathbf{y}) = \sum_{i=1}^m |x_i - y_i| \quad L1 \text{ norm}$$

- *Chebyshev* distance

$$d(\mathbf{x}, \mathbf{y}) = \max_i |x_i - y_i| \quad L\infty \text{ norm}$$

- *Hamming* distance:

$$d_i = \begin{cases} 1, & \text{if } x_i \neq y_i \\ 0, & \text{otherwise} \end{cases}; \quad d(\mathbf{x}, \mathbf{y}) = \sum_{i=1}^m d_i$$

Let $\mathbf{x} = [x_1 \quad \cdots \quad x_m]^T$ and $\mathbf{y} = [y_1 \quad \cdots \quad y_m]^T$ be two vectors, we have,

- **Cosine similarity measure:**

$$sim(\mathbf{x}, \mathbf{y}) = \frac{\langle \mathbf{x}, \mathbf{y} \rangle}{\|\mathbf{x}\| \|\mathbf{y}\|} = \frac{\sum_{i=1}^m x_i y_i}{\sqrt{\sum_{i=1}^m x_i^2} \sqrt{\sum_{i=1}^m y_i^2}}$$

- **Correlation:**

$$corr(\mathbf{x}, \mathbf{y}) = \frac{\sum_{i=1}^m (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum (x_i - \bar{x})^2} \sqrt{\sum (y_i - \bar{y})^2}} = \frac{\langle \mathbf{x} - \bar{\mathbf{x}}, \mathbf{y} - \bar{\mathbf{y}} \rangle}{\|\mathbf{x} - \bar{\mathbf{x}}\| \|\mathbf{y} - \bar{\mathbf{y}}\|} = sim(\mathbf{x} - \bar{\mathbf{x}}, \mathbf{y} - \bar{\mathbf{y}})$$

Questions about the KNN method

Q1: How do we train a KNN algorithm?

Q2: Does a KNN algorithm need a training set/test set? *Yes, validation set → k, distance*

Q3: Can the KNN method be applied to the regression problem? *Yes, average*

Q4: Is KNN a good candidate for online (real-time) classification? *Not*

The curse of dimensionality

维度诅咒


- The loss of locality in *higher dimensions*

- Consider N data points uniformly distributed in a d dimensional unit hypercube.
- Let's find out a hypercube of edge l that contains K nearest neighbors of a test point. That is, we want to find a cube of edge l to contain a fraction $r = \frac{K}{N}$ of the entire data set.
- Since the data points are uniformly distributed, this means the volume of the cube is fraction r of the unit cube, i.e.,

$$l^d = r \Rightarrow l = r^{\frac{1}{d}}$$

- When $d \rightarrow \infty$, we have $l \rightarrow 1$. i.e., the hypercube is close to the entire unit hypercube where the data is uniformly distributed!!! This is true even for very small r value.


$$0 < l < 1$$



For example, let $N = 1000, K = 10$, we have, $r = 0.01$

$$l^d = r \Rightarrow l = r^{\frac{1}{d}}$$

| d | l |
|------|-------|
| 1 | 0.01 |
| 2 | 0.1 |
| 3 | 0.21 |
| 10 | 0.63 |
| 100 | 0.96 |
| 1000 | 0.995 |



To find the 10 – nearest neighbor, we need almost the entire input space in high dimensions!!!


- To find the 10 – nearest neighbor, we need almost the entire input space in high dimensions!!!
- This breaks down the KNN assumptions because in high dimensions, the points in K nearest neighbor are not particularly closer (similar) than any other data points in the training set.
- Can we increase the number of training data points until the K nearest neighbors are truly closer to the test point?
- How many data points do we need such that l , the size of the neighborhood become truly small?

For example,

$$\text{let } l = 0.1. \text{ then, } r = \frac{K}{N} = l^d \Rightarrow N = \frac{K}{l^d} = K \times 10^d$$

When $d = 10$, N becomes a huge number!!!

Dimension reduction (PCA) may provide a solution to alleviate the problem!



Quiz

The Naïve Bayes Classifier – A generative approach

- The training data set is given as $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^N$, where, an instance x_i belongs to class \mathcal{C}_k if $y_i = k$, $k = 1, \dots, K$.
- Given a new instance, $x = [x_1 \quad x_2 \quad \cdots \quad x_d]^T$, we wish to predict the class of this instance.
- The **Bayesian approach** to classify the new instance is to assign the most probable target value, y_{MAP} , given the attribute values x_1, \dots, x_d . That is,

属性值

$$y_{MAP} = \underset{k}{\operatorname{argmax}} P(\mathcal{C}_k | x_1, \dots, x_d)$$

↑
maximum a posteriori (MAP)

$\underset{x}{\operatorname{argmax}} f(x)$: 使 $f(x)$ 最大的 x 值

最大后验

- Using **Bayes' theorem**, this probability can be written as: 求的是最大K值

$$\arg\max_k P(C_k | x_1, \dots, x_d) = y_{MAP} = \arg\max_k \frac{P(x_1, \dots, x_d | C_k)P(C_k)}{P(x_1, \dots, x_d)}$$

不依赖于类别 C_k
对于不同 C_k , $P(X)$ 相同
数据集大时会很小

$$= \arg\max_k P(x_1, \dots, x_d | C_k)P(C_k)$$

- Assume the attributes are **conditionally independent given the target value**, we have

$$P(x_1, \dots, x_d | C_k) = P(x_1 | C_k)P(x_2 | C_k) \cdots P(x_d | C_k) = \prod_{j=1}^d P(x_j | C_k)$$

- Then we have,

✓ $y_{MAP} = \arg\max_k P(C_k) \prod_{j=1}^d P(x_j | C_k)$

需要计算每个类的该值，
选最大值，就是该类

- This is why the method is called "**naïve**" Baye's method

先根据已知数据计算各 $P(C_k)$


再计算各 $P(x_1, \dots, x_d | C_k) = P(x_1 | C_k) \cdots P(x_d | C_k)$

The Naïve Bayes Classifier – A Toy Example


根据以上两数值可得各 $P(C_k | x_1, \dots, x_d) = \frac{P(x_1, \dots, x_d | C_k) P(C_k)}{P(x_1, \dots, x_d)}$

对比大小选出最大的 $P(C_k | x_1, \dots, x_d)$


- The training data set is given as:



- $y_i \in \{+, -\}, i = 1, 2, \dots, 12$ are the class labels.
- $x_i = [x_{i1} \quad x_{i2}]^T, x_{i1} \in \{\text{blue, green, red, yellow}\}, x_{i2} \in \{\text{square, circle}\}, i = 1, \dots, 12$ are the feature vectors.



- Given the following new input, let's use *naïve Bayes method* to predict its class:



- The feature vector is $\mathbf{x} = [x_1, x_2]^T = [\text{blue}, \text{square}]^T$

$$\text{Step 1: } P(+)=\frac{7}{12} \quad . \quad P(-)=\frac{5}{12}$$

$$\begin{aligned}\text{Step 2: } P(\vec{x}|+) &= P(x_1=\text{blue}, x_2=\text{square}|+) \\ &= P(x_1=\text{blue}|+)\cdot P(x_2=\text{square}|+) \\ \text{小心数据} &\Rightarrow \frac{3}{7} \times \frac{5}{7} = 0.31\end{aligned}$$

$$\begin{aligned}P(\vec{x}|-) &= P(x_1=\text{blue}, x_2=\text{square}|-) \\ &= P(x_1=\text{blue}|-)\cdot P(x_2=\text{square}|-) \\ &= \frac{3}{5} \times \frac{3}{5} = 0.36\end{aligned}$$

$$\begin{aligned}\text{Step 3: } P(+|\vec{x}) &\stackrel{\text{def}}{=} P(\vec{x}|+)P(+)=0.31 \times \frac{7}{12}=0.18 \\ P(-|\vec{x}) &= P(\vec{x}|-)P(-)=0.36 \times \frac{5}{12}=0.15\end{aligned}$$

$$\text{Step 4: } P(+|\vec{x}) > P(-|\vec{x})$$

Thus, blue square is classified into class +

- Under the assumption that the training samples are *i. i. d*, we have the class priors as:

$$P(+) = \frac{7}{12} = 0.58; \quad P(-) = \frac{5}{12} = 0.42 \quad P(C_k)$$

- Assuming features are conditionally independent for given class**, the class conditional probabilities can be calculated as:

$$P(x|+) = P(x_1 = \text{blue}, x_2 = \text{square}|+) = P(\text{blue}|+)P(\text{square}|+) = \frac{3}{7} \times \frac{5}{7} = 0.31$$

$$P(x|-) = P(x_1 = \text{blue}, x_2 = \text{square}|-) = P(\text{blue}|-)P(\text{square}|-) = \frac{3}{5} \times \frac{3}{5} = 0.36$$

$$P(+|x) = \frac{P(x|+) P(+)}{P(x)}$$


- Then the posterior can be calculated as:

$$P(+|x) = P(x|+)P(+) = 0.58 \times 0.31 = 0.18$$

$$P(-|x) = P(x|-)P(-) = 0.42 \times 0.36 = 0.15$$

- Since $P(+|x) = 0.18 > P(-|x) = 0.15$, the sample point $x =$  is classified as +.

- Suppose a new sample as the following is given:



- The feature vector is $x = [x_1, x_2]^T = [\text{yellow}, \text{square}]^T$, the color attribute value “yellow” is not present in the training data set. This causes the class conditional probability to become:

$$P(x|+) = P(x_1 = \text{yellow}, x_2 = \text{square}|+) = P(\text{yellow}|+)P(\text{square}|+) = 0 \times \frac{5}{7} = 0$$

$$P(x|-) = P(x_1 = \text{yellow}, x_2 = \text{square}|-) = P(\text{yellow}|-)P(\text{square}|-) = 0 \times \frac{3}{5} = 0$$

And the posteriors will hence be all calculated as 0s.

- This problem can be solved by using the **Laplace smoothing (additive smoothing)** technique to calculate the class conditional probability:

$$P(x_i | \mathcal{C}_k) = \frac{N_{x_i, \mathcal{C}_k} + 1}{N_{\mathcal{C}_k} + d}$$

特征 x_i 在类别 \mathcal{C}_k 训练集中出现的次数

where, N_{x_i, \mathcal{C}_k} is the number of time feature x_i appears in training data set from class \mathcal{C}_k ;

\mathcal{C}_k 类训练集中所有特征出现的总数

$N_{\mathcal{C}_k}$ is the total count of all features appears in training data set from class \mathcal{C}_k ; d is the dimensionality of the feature vector

特征向量 \vec{x} 的维度

- Then, using this new technique, the class conditional probability of the new sample x can be calculated as:

$$P(x|+) = P(x_1 = \text{yellow}, x_2 = \text{square}|+) = P(\text{yellow}|+)P(\text{square}|+)$$

$\Rightarrow d=2$

$$= \frac{1}{7+2} \times \frac{5}{7} = 0.08$$

yellow $\frac{1}{7+2}$
 square $\frac{5}{7}$
 7+2

$$P(x|-) = P(x_1 = \text{yellow}, x_2 = \text{square}|-) = P(\text{yellow}|-)P(\text{square}|-)$$

$$= \frac{1}{5+2} \times \frac{3}{5} = 0.09$$

yellow $\frac{1}{5+2}$
 square $\frac{3}{5}$
 5+2

- Note that the Laplace smoothing formula is used to calculate the likelihood of the features that do not present in the training data set.

- The posterior probability of the classes can be calculated as:

$$P(+|x) = P(x|+)P(+) = 0.08 \times 0.58 = 0.0464$$

$$P(-|x) = P(x|-)P(-) = 0.09 \times 0.42 = 0.0378$$

Therefore, the new sample  should be classified as +

The Naïve Bayes Classifier – Text Documents Classification

- ***Inputs:***

- a document x
- A fixed set of classes $\{\mathcal{C}_1, \dots, \mathcal{C}_K\}$ e.g., {"spam", "ham"}, {"computer science", "biology", "statistics", "economics", "politics"}
- A training set of N hand-labeled documents $\{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$ (training corpus)

- ***Output:***

- The class of the document x

Example: Email

Preprocessing of text documents

- **Tokenization:** breaks down a text corpus into individual words and removes punctuations

语料库

A swimmer likes swimming; thus, he swims.



| | | | | | | |
|---|---------|-------|----------|------|----|-------|
| a | swimmer | likes | swimming | thus | he | swims |
|---|---------|-------|----------|------|----|-------|

Preprocessing of text documents

- **Remove stop words:** stop words are particularly common in a text corpus and thus considered as rather uninformative (e.g., words such as *so, and, or, the*, etc.) a stop list can be created and used for stop words removal.

A swimmer likes swimming; thus, he swims.

| | | | | |
|---------|-------|----------|----|-------|
| swimmer | likes | swimming | he | swims |
|---------|-------|----------|----|-------|

The Bag of Words Model

- **Creation of the vocabulary:** after preprocessing, we can create the collection of all different words that occur in the training data set and each word is associated with a count of how it occurs. **This is a set of non-redundant items where the order does not matter (bag of words).** 多余的

Let D_1 and D_2 be two documents in a training dataset (corpus):

- D_1 : “Each state has its own laws.”
- D_2 : “Every country has its own culture.”

Based on these two documents, after proper preprocessing, the vocabulary can be created as:



$$V = \{each: 1, state: 1, has: 2, its: 2, own: 2, laws: 1, every: 1, country: 1, culture: 1\}$$

- **Vectorization:** The vocabulary can then be used to construct the d -dimensional feature vector for the individual documents. $d = |V|$ is the number of different words in the vocabulary.

- The following are the ***bag of words*** representation of two sample documents D_1 and D_2 : (word position information does not matter)

2 documents to 2 vectors



| V | each | state | has | its | own | laws | every | country | culture |
|-----------|------|-------|-----|-----|-----|------|-------|---------|---------|
| x_{D_1} | 1 | 1 | 1 | 1 | 1 | 1 | 0 | 0 | 0 |
| x_{D_2} | 0 | 0 | 1 | 1 | 1 | 0 | 1 | 1 | 1 |
| Σ | 1 | 1 | 2 | 2 | 2 | 1 | 1 | 1 | 1 |

- After transforming documents into numbers, popular supervised learning methods can be applied to do document classification.

Text Documents Classification using The Naïve Bayes Classifier

- Given a test document $x = [x_1, x_2, \dots, x_d]^T$

- Prior probability estimate:**

$$P(\mathcal{C}_k) = \frac{\text{\# of documents with } \mathcal{C}_k}{\text{Total \# of documents in the training data set}}, \quad k = 1, 2, \dots, K$$

- Class conditional probability estimate: (Laplace smoothing if necessary)**

$$P(x_i | \mathcal{C}_k) = \frac{N(x_i, \mathcal{C}_k) + 1}{N(\mathcal{C}_k) + |V|}$$

Where, $N(x_i, \mathcal{C}_k)$ is the total number of occurrences of term x_i in all documents in the training data set with class \mathcal{C}_k ; $N(\mathcal{C}_k)$ is the number of occurrences of all the terms in all documents in the training data set with class \mathcal{C}_k .

- ***Class conditional probability estimate:***

$$P(x_i|\mathcal{C}_k) = \frac{N(x_i, \mathcal{C}_k) + 1}{N(\mathcal{C}_k) + |V|}$$

Where, $N(x_i, \mathcal{C}_k)$ is the total number of occurrences of term x_i in all documents in the training data set with class \mathcal{C}_k ; $N(\mathcal{C}_k)$ is the number of occurrences of all the terms in all documents in the training data set with class \mathcal{C}_k .

- Then, assuming features are conditionally independent given class, we have, (“naïve”)

$$P(x|\mathcal{C}_k) = P(x_1, x_2, \dots, x_d|\mathcal{C}_k) = \prod_{i=1}^d P(x_i|\mathcal{C}_k), k = 1, \dots, K$$

The chance of occurrence of value x_i is not affected by the values of other attributes.

- ***Calculation of the posterior probability:***

$$P(\mathcal{C}_k|x) \propto P(x|\mathcal{C}_k)P(\mathcal{C}_k)$$