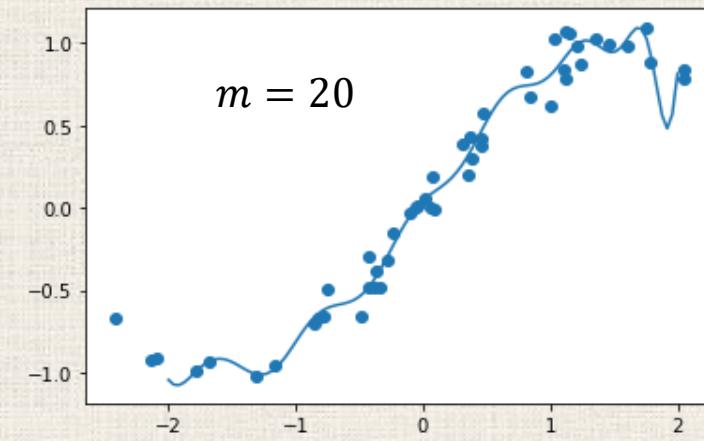
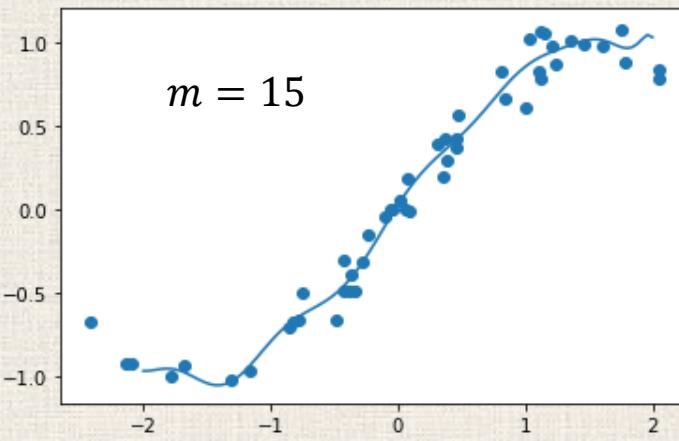
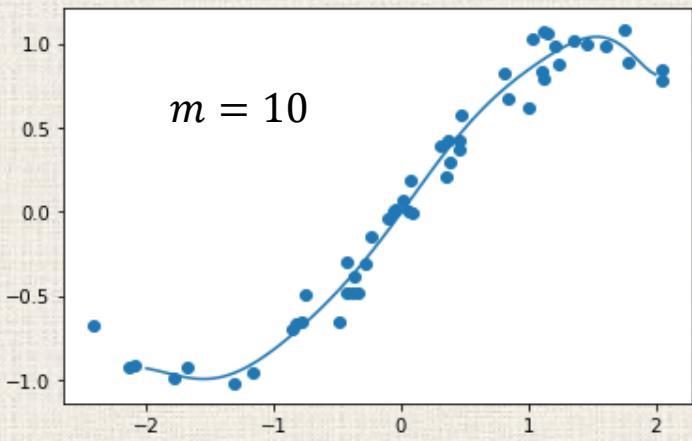
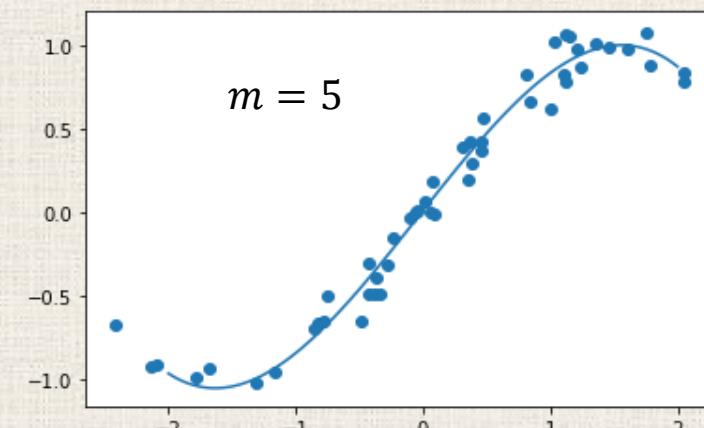
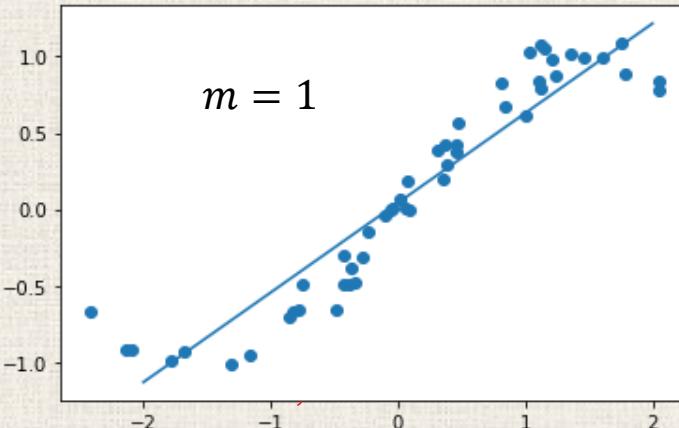
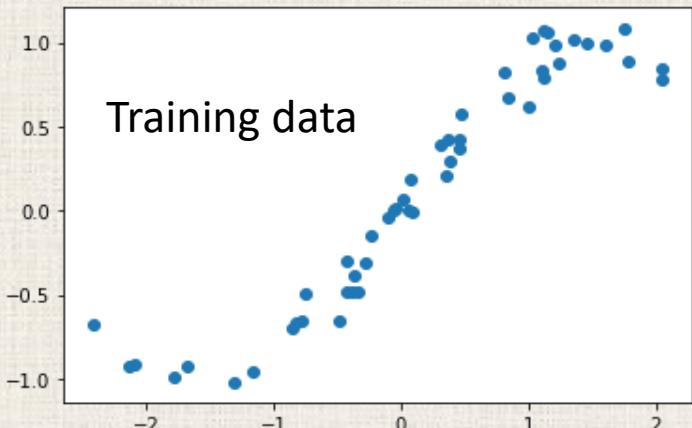


# **Performance Evaluation and Model Selection**

# Underfitting and Overfitting Models



- Underfitting: the model used is too simple to catch the trend behind the given data set;
- Overfitting: the model is so powerful; it not only catch the trend but also the noise in the data set;

- **Model selection:** which model should we use to achieve the best performance?
- **Performance evaluation:** Can we use the training error (SSE in the example) to measure the performance of the trained models (fitted curves)?
- Do we have a procedure to find the “best” model for a given data set?

# Parameters in the supervised machine learning problems

$$\mathbf{w}_{Ridge} = \underset{\mathbf{w}}{\operatorname{argmin}} \left\{ \|\Phi \mathbf{w} - \mathbf{y}\|_2^2 + \lambda \|\mathbf{w}\|_2^2 \right\} = (\Phi^T \Phi + \lambda I)^{-1} \Phi^T \mathbf{y}$$

Assuming polynomial model is used in this ridge regression problem.

- The weight vector,  $\mathbf{w}$ , is called the ***model parameter***, it is the ***decision variable***, a parameter to be optimized.
- The order of polynomial model,  $m$ , and the regularization weight,  $\lambda$ , are not decision variables, they are called ***hyperparameters***.
  - $m$  is a ***model hyperparameter***, since it determines the structure of the model.
  - $\lambda$  is not part of the model. Rather, it is an aspect of the optimization process used to fit the model. It is called an ***optimization hyperparameter***.

## Types of Error (true error vs training error)

- Assume that the data are distributed according to some (unknown) distribution  $p(x, y)$ , and that we have a ***loss function***  $l$ , which is to measure the error between the true output and our estimate  $\hat{y} = h_w(x)$ .
- The ***true error*** of a particular hypothesis  $h \in H$  is the ***expected loss over the whole data distribution*** (over the entire ***population***):

$$R(w) = E_{x,y}[l(h_w(x), y)]$$

- Ideally, we would find the hypothesis that minimizes the ***true error***, i.e.,

$$w^* = \underset{w}{\operatorname{argmin}} R(w)$$

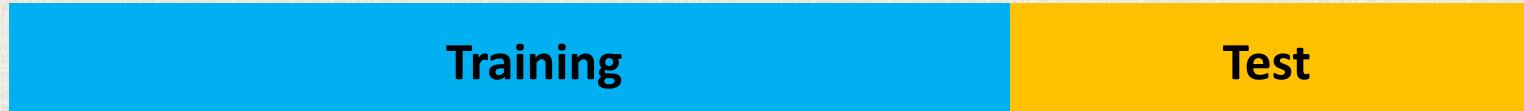
- However, computing this expectation is impossible because we do not have access to the true data distribution.
- Rather, we have access to samples  $(x_i, y_i)$  that follows the same distribution.
- These enable us to approximate the real problem by minimizing the ***training error (in-sample error)***:

$$\hat{R}_{Train}(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N l(h_{\mathbf{w}}(x_i), y_i)$$

- But since we have a finite number of samples, ***the hypothesis that performs the best on the training data is not necessarily the best on the whole data distribution.***

## Training set vs Test set

- A common solution is to set aside some portion (say 30%) of the data (*held-out samples*), to be called the *test set*, which is *disjoint* from the training set and *not allowed to be used when fitting the model*:



- We can use this test set to estimate the **true error** by the *test error (out-of-sample error)* :

$$\hat{R}_{test}(\mathbf{w}) = \frac{1}{M} \sum_{i=1}^M l(h_{\mathbf{w}}(\mathbf{x}_i), y_i)$$

- *Never touch the test set until the final fitted model is decided.*

For example, let's solve the OLS problem on the Iris data set.

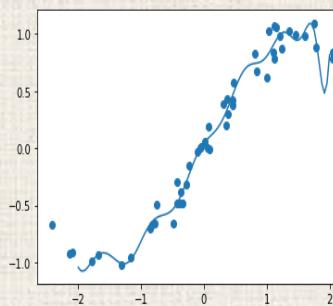
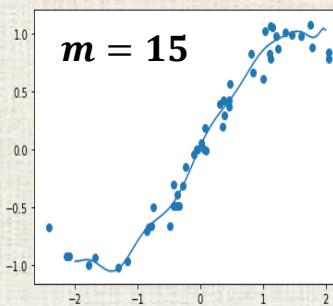
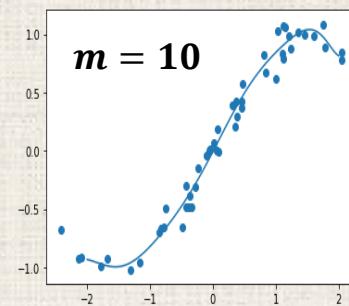
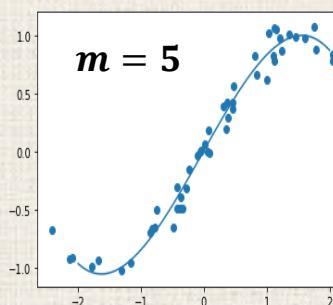
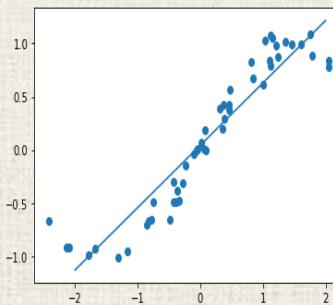
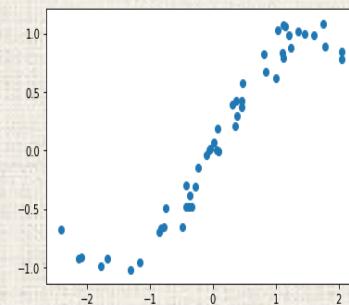
|     | sepal length (cm) | sepal width (cm) | petal length (cm) | petal width (cm) | species |
|-----|-------------------|------------------|-------------------|------------------|---------|
| 0   | 5.1               | 3.5              | 1.4               | 0.2              | 0       |
| 1   | 4.9               | 3                | 1.4               | 0.2              | 0       |
| 2   | 4.7               | 3.2              | 1.3               | 0.2              | 0       |
| 3   | 4.6               | 3.1              | 1.5               | 0.2              | 0       |
| 4   | 5                 | 3.6              | 1.4               | 0.2              | 0       |
| ... |                   |                  |                   |                  |         |
| 50  | 7                 | 3.2              | 4.7               | 1.4              | 1       |
| 51  | 6.4               | 3.2              | 4.5               | 1.5              | 1       |
| 52  | 6.9               | 3.1              | 4.9               | 1.5              | 1       |
| 53  | 5.5               | 2.3              | 4                 | 1.3              | 1       |
| 54  | 6.5               | 2.8              | 4.6               | 1.5              | 1       |
| ... |                   |                  |                   |                  |         |
| 144 | 6.7               | 3.3              | 5.7               | 2.5              | 2       |
| 145 | 6.7               | 3                | 5.2               | 2.3              | 2       |
| 146 | 6.3               | 2.5              | 5                 | 1.9              | 2       |
| 147 | 6.5               | 3                | 5.2               | 2                | 2       |
| 148 | 6.2               | 3.4              | 5.4               | 2.3              | 2       |
| 149 | 5.9               | 3                | 5.1               | 1.8              | 2       |

$$\begin{aligned}\mathbf{w}^* &= \underset{\mathbf{w}}{\operatorname{argmin}}\{L(\mathbf{w}) = \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2\} \\ &= (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}\end{aligned}$$

- The entire data set is split into training set (100 data points) and test set (50 data points)
- The training set is used to find  $\mathbf{w}^*$
- The test set is used to evaluate the performance (true error) of  $\mathbf{w}^*$

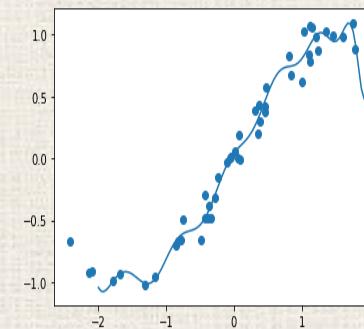
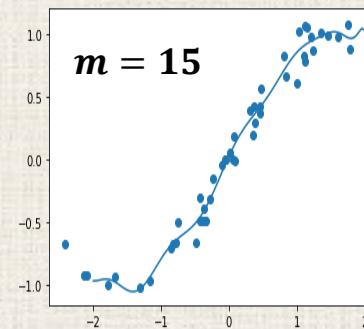
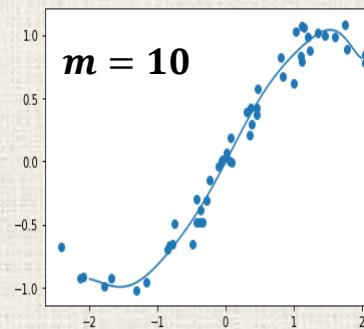
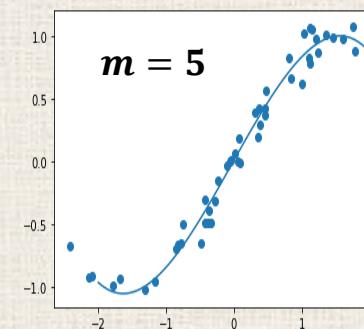
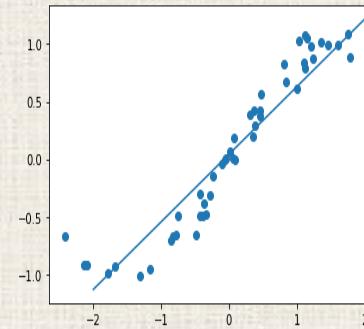
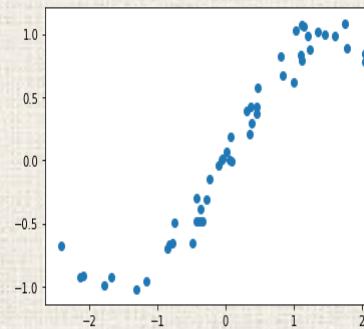
# Training set vs validation set

- To tune hyperparameters, we need to evaluate the performance of the fitted model.
- For example, we want to find out which  $m$  value (out of 5, 10, 15) in the polynomial models fit the given data set the best?



- First, the given data set is split into two parts: the training set (70%), and the test set (30%).
- Then, the following is carried out:
  - The polynomial model with  $m = 5$  is trained on the training set and find the optimal weight vector  $\mathbf{w}_1^*$  (fitted model #1)
  - The polynomial model with  $m = 10$  is trained on the training set and find the optimal weight vector  $\mathbf{w}_2^*$  (fitted model #2)
  - The polynomial model with  $m = 15$  is trained on the training set and find the optimal weight vector  $\mathbf{w}_3^*$  (fitted model #3)
- What are the differences among  $\mathbf{w}_1^*$ ,  $\mathbf{w}_2^*$ , and  $\mathbf{w}_3^*$  ?

- To compare the performance of these three models, corresponding ***training errors*** can not be used.
- The ***test set*** is reserved to evaluate the performance of the finalized model.
- What should we do?



- Understanding that the test set will be only used to evaluate the finalized model, this evaluation can only be carried out on the training set.
- Hence, in many cases, ***we need to split the training set into two parts:***
  - one part (called training set) will be used to training the fitting model
  - the other part (called ***validation set***) will be used for evaluating the fitted model with certain hyperparameter values.
  - We can use this validation set to estimate the true error by the ***validation error***:

$$\hat{R}_{VAL}(\mathbf{w}) = \frac{1}{K} \sum_{i=1}^K l(h_{\mathbf{w}}(\mathbf{x}_i), y_i)$$

- With this estimate, we have a simple method for choosing hyperparameter values: ***try a bunch of configurations of the hyperparameters and choose the one that yields the lowest validation error***

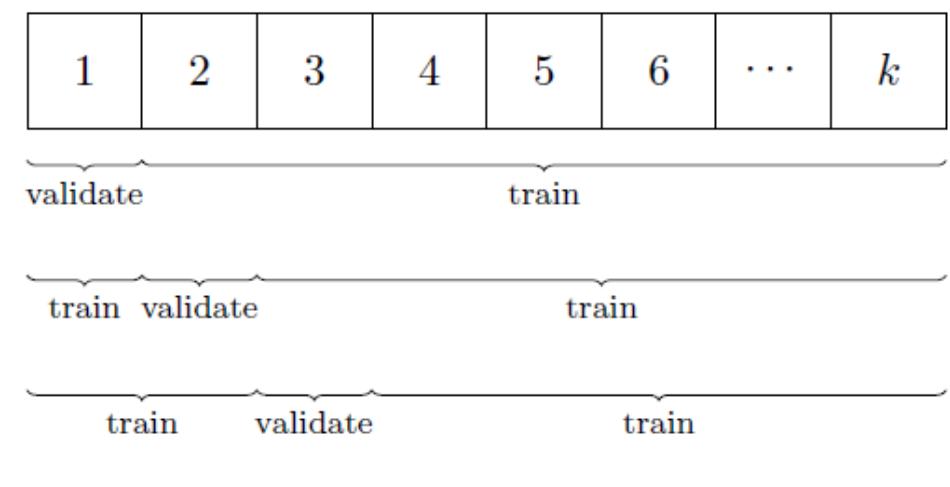
- Note that the ***purpose of validation set is to determine the best values of the hyperparameters! (used for hyper-parameters tuning)***
- Then, the tuned model (the finalized model) will be tested using the test data set.
- Hence, in general the original data is split into the following:



- ***Setting aside a validation set works at a cost!***
- We can not use the validation data for training.
- This is especially important when we have little data to begin with or collecting more data is expensive.

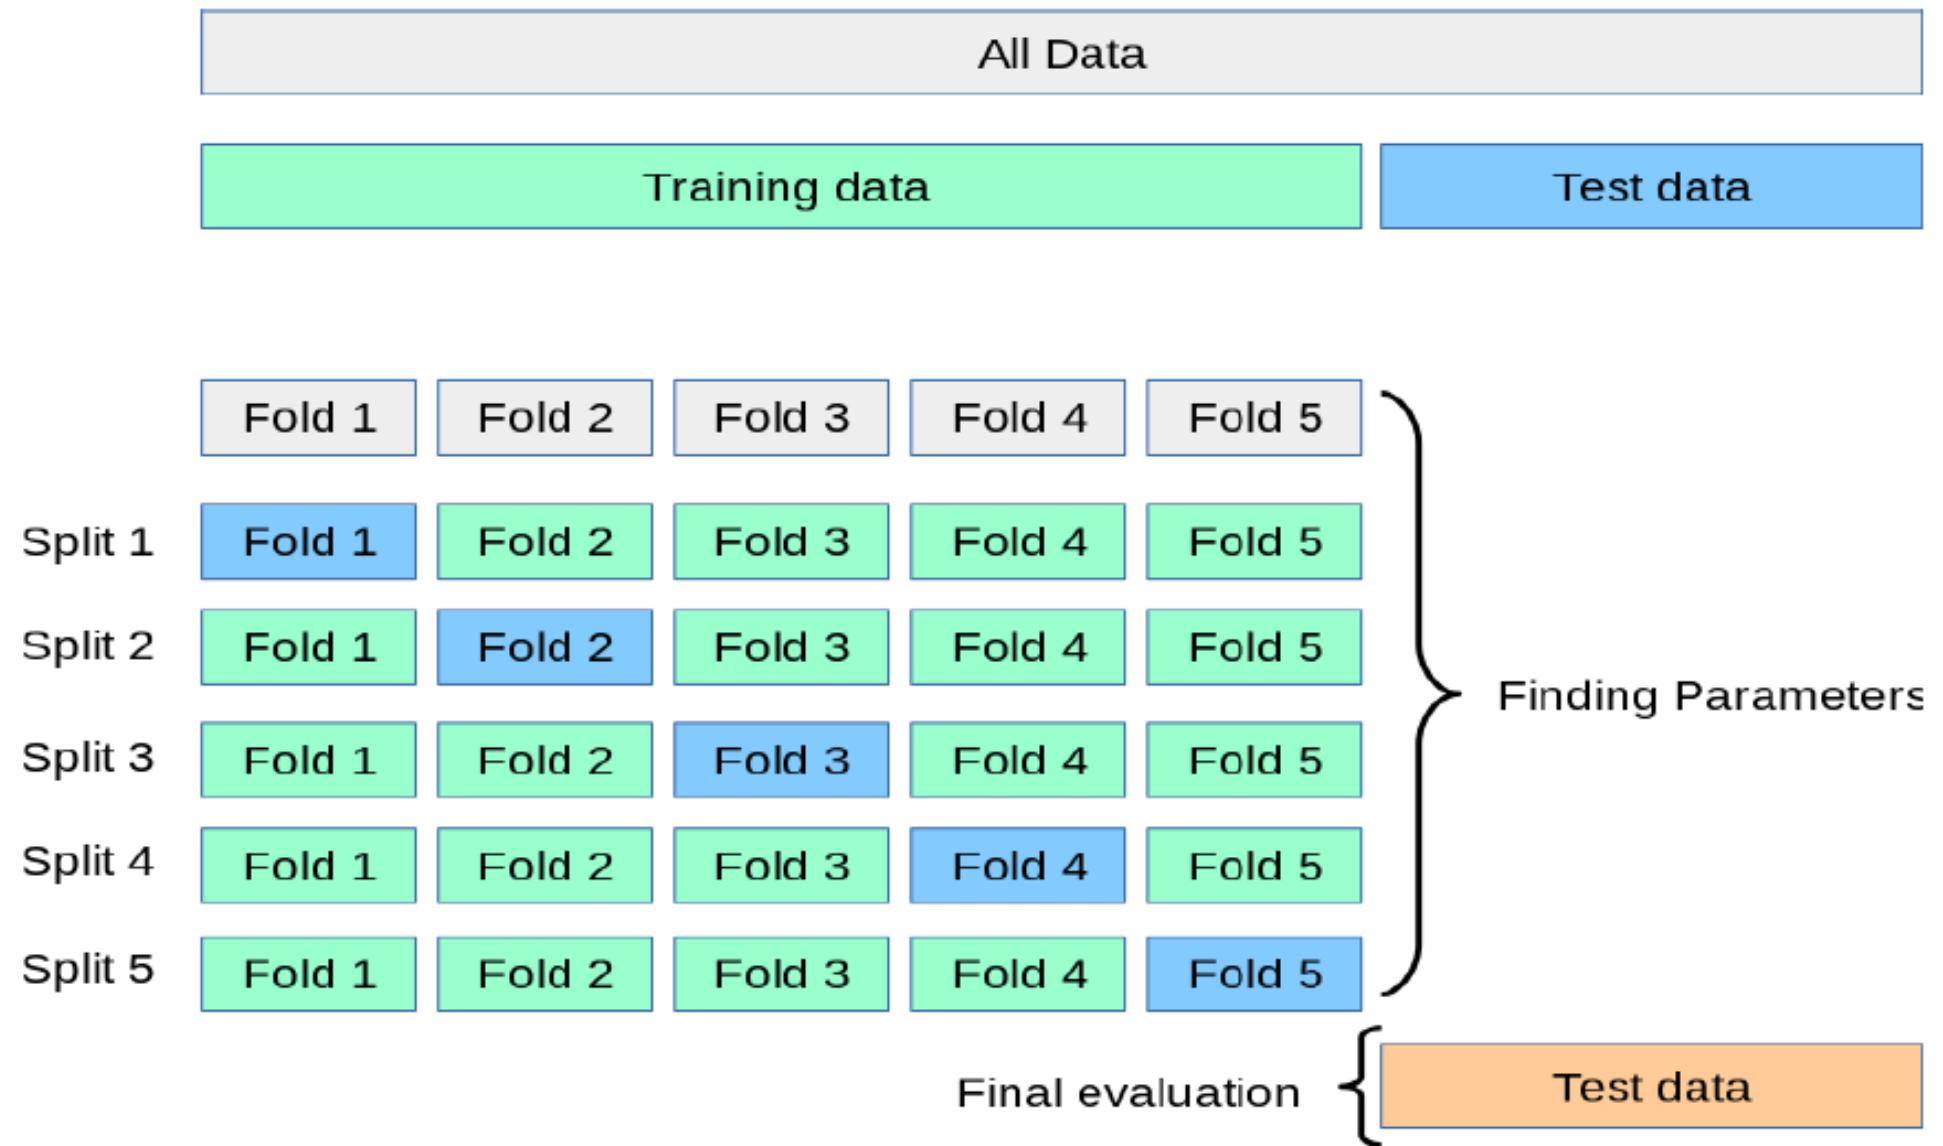
# Cross-validation

- **Cross-validation** is an alternative to having a dedicated validation set.
- ***k*-fold cross-validation** works as follows:
  1. Shuffle the training data and partition it into  $k$  equally-sized blocks.
  2. For  $i = 1, \dots, k$ 
    - Train the model on all the data except block  $i$
    - Evaluate the model (i.e., compute the validation error) using block  $i$
  3. Average the  $k$  validation errors; this is the final estimate of the true error.



## Cross-validation

- Note that this process (except for the shuffling and partitioning) must be repeated for every hyperparameter configuration we wish to test.
- This is roughly  $k$  times as much computation required using the held-out validation set.
- This can be very ***expensive*** when the model takes a long time to train.



Training Data  
(optimize the model  
parameters)

Validation Data  
(optimize the model's  
hyperparameters)

Test Data  
(evaluate the  
optimized model)

# Tuning the Hyperparameters

- **Hyperparameters are not model parameters, and they can not be directly trained from the data.**
- **The general process of tuning hyper-parameters:**
  - Define a model
  - Define the range of possible values of all hyperparameters
  - Define a method for sampling hyperparameter values
  - Define an evaluation criteria to judge the model
  - Define a cross-validation method

- **Grid Search:** an exhaustive sampling of the hyperparameter space. i.e., we simply build a model for each possible combination of all the hyperparameter values provided, evaluate each model and select the structure which produces the best results.

- For example: Consider a ***random forests classifier***, which is an ***ensemble model*** comprised of a collection of ***decision trees***. This model has two hyperparameters:

N\_trees: the number of decision trees included, with values [10,50,100]

M\_depth: maximum depth allowed for each decision tree model, with values [3,5]

Performing grid search would generate the following models to be evaluated:

Random\_forest\_classifier(N\_tree=10,M\_depth=3)

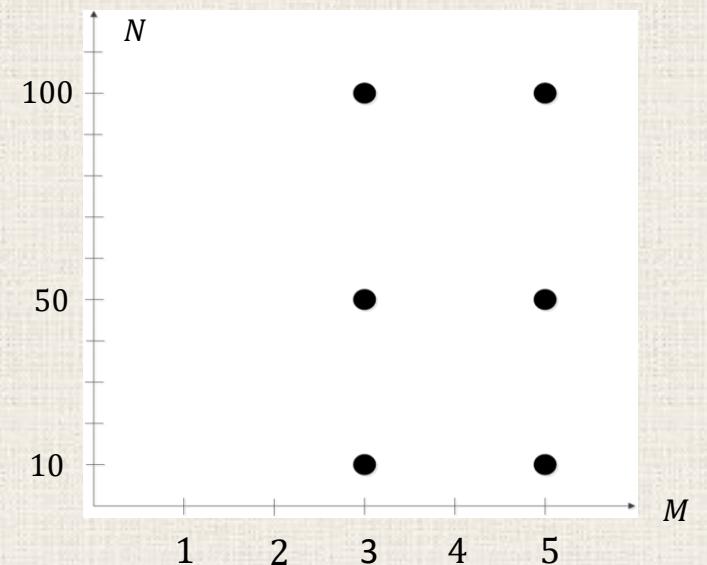
Random\_forest\_classifier(N\_tree=10,M\_depth=5)

Random\_forest\_classifier(N\_tree=50,M\_depth=3)

Random\_forest\_classifier(N\_tree=50,M\_depth=5)

Random\_forest\_classifier(N\_tree=100,M\_depth=3)

Random\_forest\_classifier(N\_tree=100,M\_depth=5)



- Each of these model will then be evaluated on a validation set or using cross-validation method.

- **Random Search:** tries random combinations of a range of values of hyperparameters.
  - It is good in testing a wide range of values and normally it reaches a very good combination very fast, but it does not guarantee to give the best parameter combination.
  - Random search has a probability of 95% of finding a combination of parameters within the 5% optima with only 60 iterations.

- Consider any distribution in a hyperparameter space with a finite maximum.
- Let's randomly sample points from this space and check if any of them lands in the 5% interval within the maximum. Each trial will have 5% chance that the sample point is within that interval.
- If the sampling within each trial is independent, then, the probability that all  $n$  trials missed the desired interval is  $(1 - 0.05)^n$  .
- So, the probability that at least one trial hit the desired interval after  $n$  trials is  $1 - (1 - 0.05)^n$ . If we want this probability to be larger than 95%, i.e.,

$$1 - (1 - 0.05)^n > 0.95$$

- The number of trials can be found as  $n \geq 60$ .

## **Goodness-of-fit (performance metrics of regression model)**

The following are the popular performance metrics used to evaluate a trained regression model:

- **Mean Squared Error (MSE):** (`mean_squared_error` in sk-learn)

Let  $\hat{y}_i$  be the predicted value of the  $i^{th}$  sample, and  $y_i$  be the corresponding target value. Then the mean squared error estimated over  $n$  samples is defined as:

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

- **Mean Absolute Error (MAE):** (`mean_absolute_error` in sk-learn)

Let  $\hat{y}_i$  be the predicted value of the  $i^{th}$  sample, and  $y_i$  be the corresponding target value. Then the mean absolute error estimated over  $n$  samples is defined as:

$$MAE = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

- **$R^2$  score: coefficient of determination (r2\_score in sk-learn)**

Let  $\hat{y}_i$  be the predicted value of the  $i^{th}$  sample, and  $y_i$  be the corresponding target value. Then the  $R^2$  score estimated over  $n$  samples is defined as:

$$R^2 = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$$

Where,

$$\bar{y} = \frac{1}{n} \sum_{i=1}^n y_i$$

$R^2$  is a statistical measure of how well the regression predictions approximate the real data points.

**For more detailed information on performance evaluation of regression models using sk-learn, visit the website:**

[https://scikit-learn.org/stable/model\\_selection.html](https://scikit-learn.org/stable/model_selection.html)

# **Bias-Variance Decomposition of Model Prediction Error**

- Recall from the regression problem, our goal is to estimate a hypothesis model  $h(\mathbf{x})$  to approximate the unknown model  $f(\mathbf{x})$  governing the sampled data set.
- One question is: ***how exactly can we measure the effectiveness of a hypothesis model(model prediction error)?***
- Let's make some assumptions:
  - We have a training data set:  $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^N$ ,  $\mathbf{x}_i$  and  $y_i$ 's are values sampled from random variables  $X_i$  and  $Y_i$ , that is, the training data set is a random sample from a random variable  $\mathcal{D}$ , consisting of random variables  $X_i$  and  $Y_i$ .
  - For some arbitrary test input  $\mathbf{x}$ , we use the trained model to estimate corresponding output,  $h(\mathbf{x}; \mathcal{D})$ , which depends on the random variable  $\mathcal{D}$  that was used to train  $h$ . Since  $\mathcal{D}$  is random, we will have a slightly different hypothesis model  $h(\mathbf{x}; \mathcal{D})$  every time we use a new training dataset.
  - Note that  $\mathbf{x}$  and  $\mathcal{D}$  are completely independent from one another:  $\mathbf{x}$  is a test point, while  $\mathcal{D}$  consists of the training data.

- Our objective is to, for a fixed test point  $\mathbf{x}$ , evaluate how closely the trained model can estimate (predict) the noisy observation  $Y$  corresponding to  $\mathbf{x}$ .
- We express our metric as the ***expected squared error*** between the prediction and the observation  $Y = f(\mathbf{x}) + Z$ :

$$\epsilon(\mathbf{x}, h) = E[(h(\mathbf{x}, \mathcal{D}) - Y)^2]$$

The expectation here is over two random variables,  $\mathcal{D}$  and  $Y$ .

- We need the following facts to decompose the error:
  - First, let's find the expectation and variance of  $Y$ , assuming  $E[Z] = 0$ :

$$E[Y] = E[f(\mathbf{x}) + Z] = f(\mathbf{x}) + E[Z] = f(\mathbf{x})$$

$$Var(Y) = Var(f(\mathbf{x}) + Z) = Var(Z)$$

- Also, notice that for any random variable  $X$ , we have,

$$Var(X) = E[X^2] - (E[X])^2 \Rightarrow E[X^2] = Var(X) + (E[X])^2$$

- Let's use these facts to decompose the error:

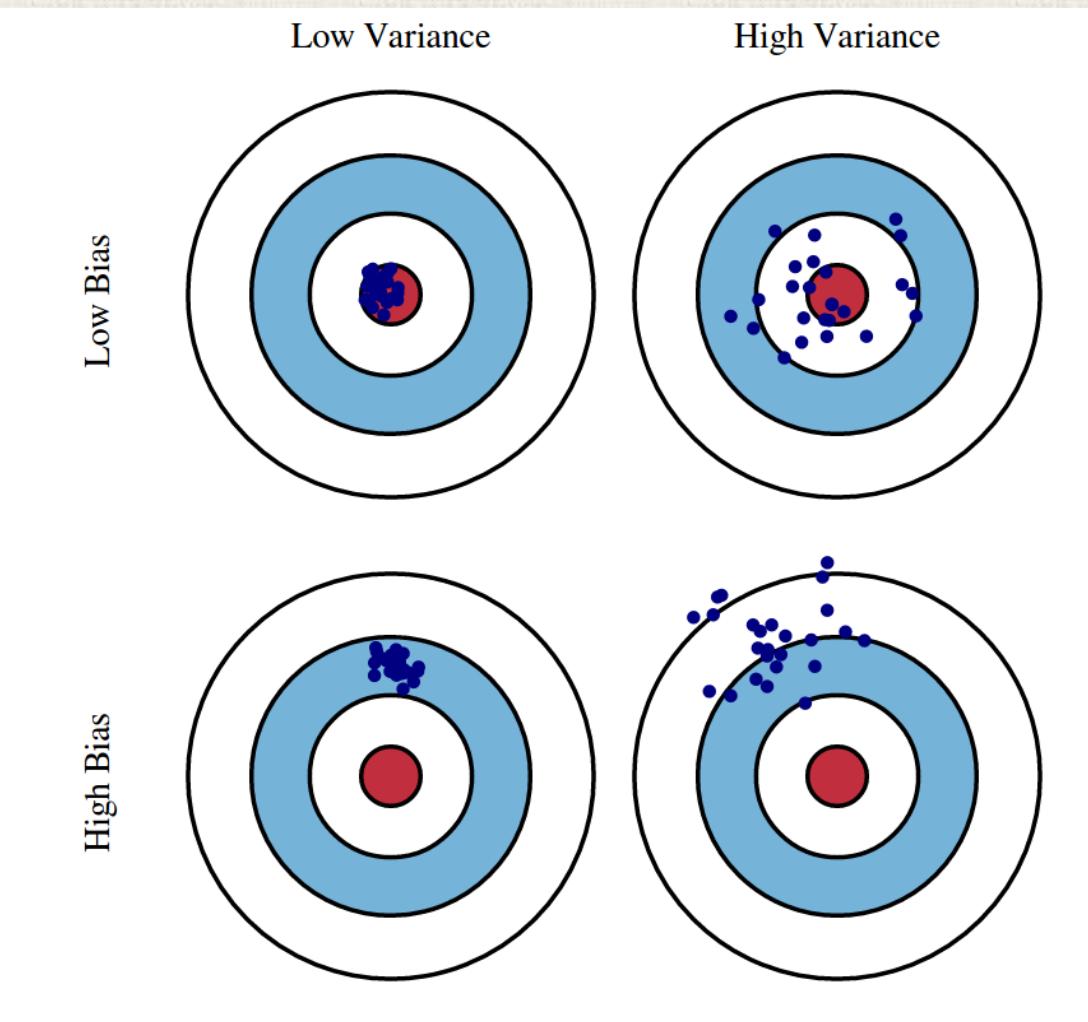
$$\begin{aligned}
 \epsilon(\mathbf{x}, h) &= E[(h(\mathbf{x}, \mathcal{D}) - Y)^2] \\
 &= E[(h(\mathbf{x}, \mathcal{D})^2 - 2h(\mathbf{x}, \mathcal{D})Y + Y^2)] = E[h(\mathbf{x}, \mathcal{D})^2] - 2E[h(\mathbf{x}, \mathcal{D})Y] + E[Y^2] \\
 &= \{Var(h(\mathbf{x}, \mathcal{D})) + (E[h(\mathbf{x}, \mathcal{D})])^2\} - 2E[h(\mathbf{x}, \mathcal{D})]E[Y] + \{Var(Y) + (E[Y])^2\} \\
 &= \{(E[h(\mathbf{x}, \mathcal{D})])^2 - 2E[h(\mathbf{x}, \mathcal{D})]E[Y] + (E[Y])^2\} + Var(h(\mathbf{x}, \mathcal{D})) + Var(Y) \\
 &= \{E[h(\mathbf{x}, \mathcal{D})] - E[Y]\}^2 + Var(h(\mathbf{x}, \mathcal{D})) + Var(Y) \\
 &= \{E[h(\mathbf{x}, \mathcal{D})] - f(\mathbf{x})\}^2 + Var(h(\mathbf{x}, \mathcal{D})) + Var(Z)
 \end{aligned}$$

|      |          |       |
|------|----------|-------|
| Bias | Variance | Noise |
|------|----------|-------|

Now that the error is decomposed into three parts, where,

- $\{E[h(\mathbf{x}, \mathcal{D})] - f(\mathbf{x})\}^2$  is called the ***bias of model***: which measures how well the average hypothesis (over all possible training sets) can come close to the true underlying value  $f(\mathbf{x})$ , for a fixed value of  $\mathbf{x}$ .
  - A low bias means that on average the fitted model  $h(\mathbf{x})$  accurately estimates  $f(\mathbf{x})$ .
- $Var(h(\mathbf{x}, \mathcal{D}))$  is the ***variance of model***: which measures the variance of the hypothesis (over all possible training sets), for a fixed value of  $\mathbf{x}$ .
  - A low variance means that the prediction does not change much as the training set varies.
- $Var(Z)$  is the ***irreducible error***: this is the error in the model that we can not control or eliminate, because it is due to errors inherent in the noisy observation  $Y$ .

*expected squared error = bias + variance + noise*



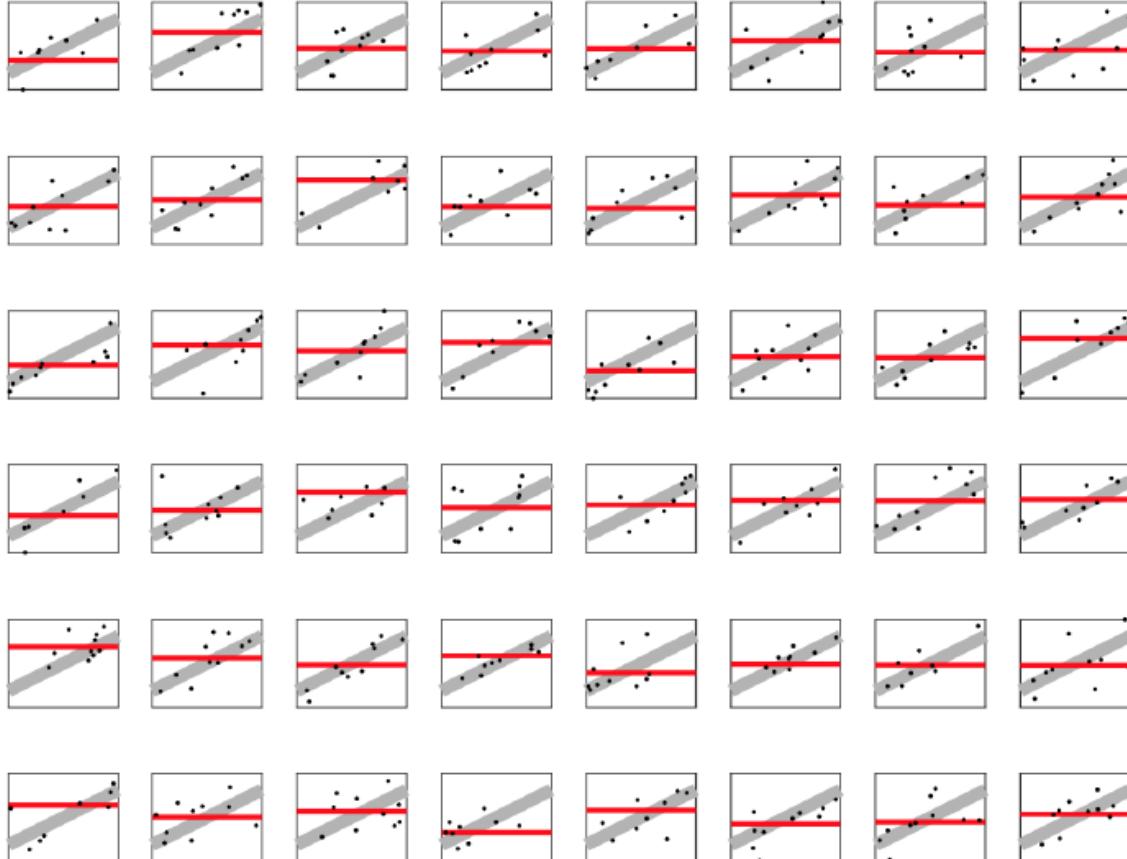
# An experiment on Bias-Variance tradeoff

*expected squared error = bias + variance + noise*

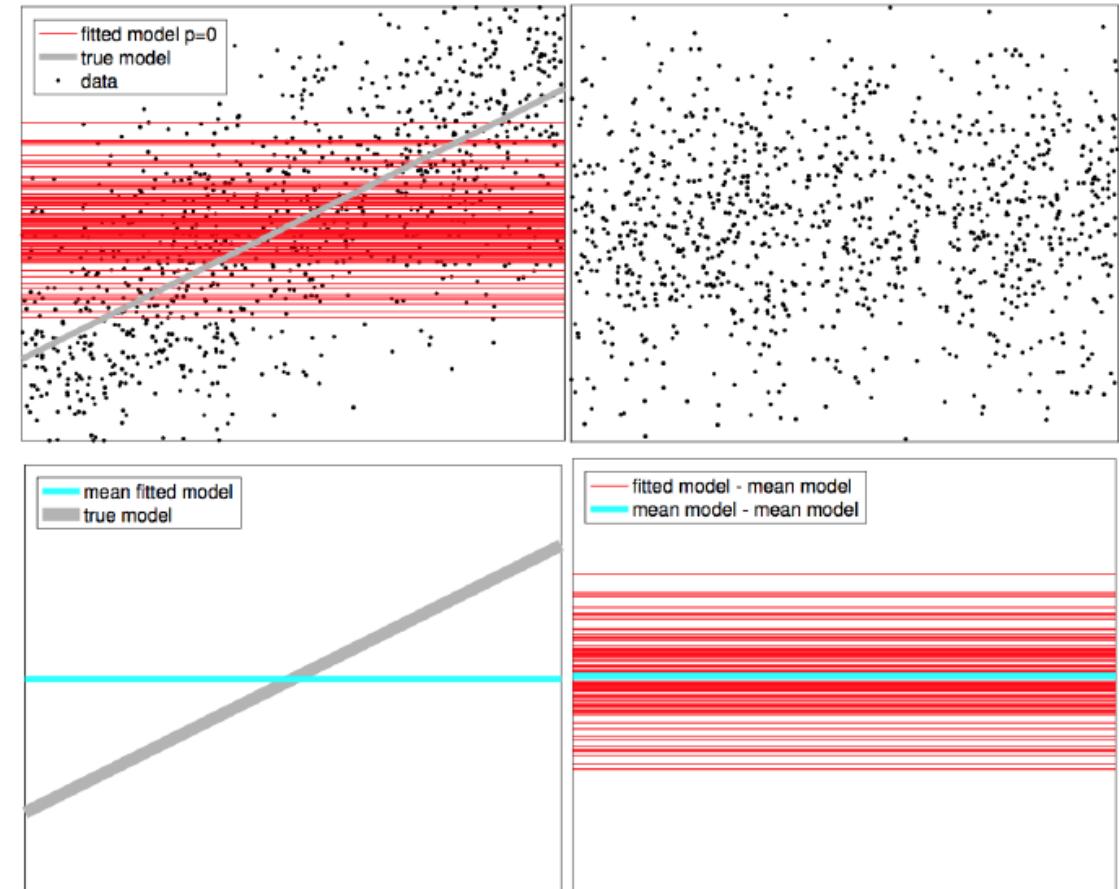
- Let's confirm the result of bias-variance decomposition with an experiment that measures the bias and variance for polynomial regression with 0 degree, 1<sup>st</sup> degree, and 2<sup>nd</sup> degree polynomials.
- In our experiment, we will repeatedly fit our hypothesis model to a random training set of 10 points drawn from a linear model  $y = a_0 + a_1x + z$ , where  $z \sim N(0, \sigma^2)$ .

# An experiment on Bias-Variance tradeoff

Fitting A Model over Multiple Datasets:  $m = 0$

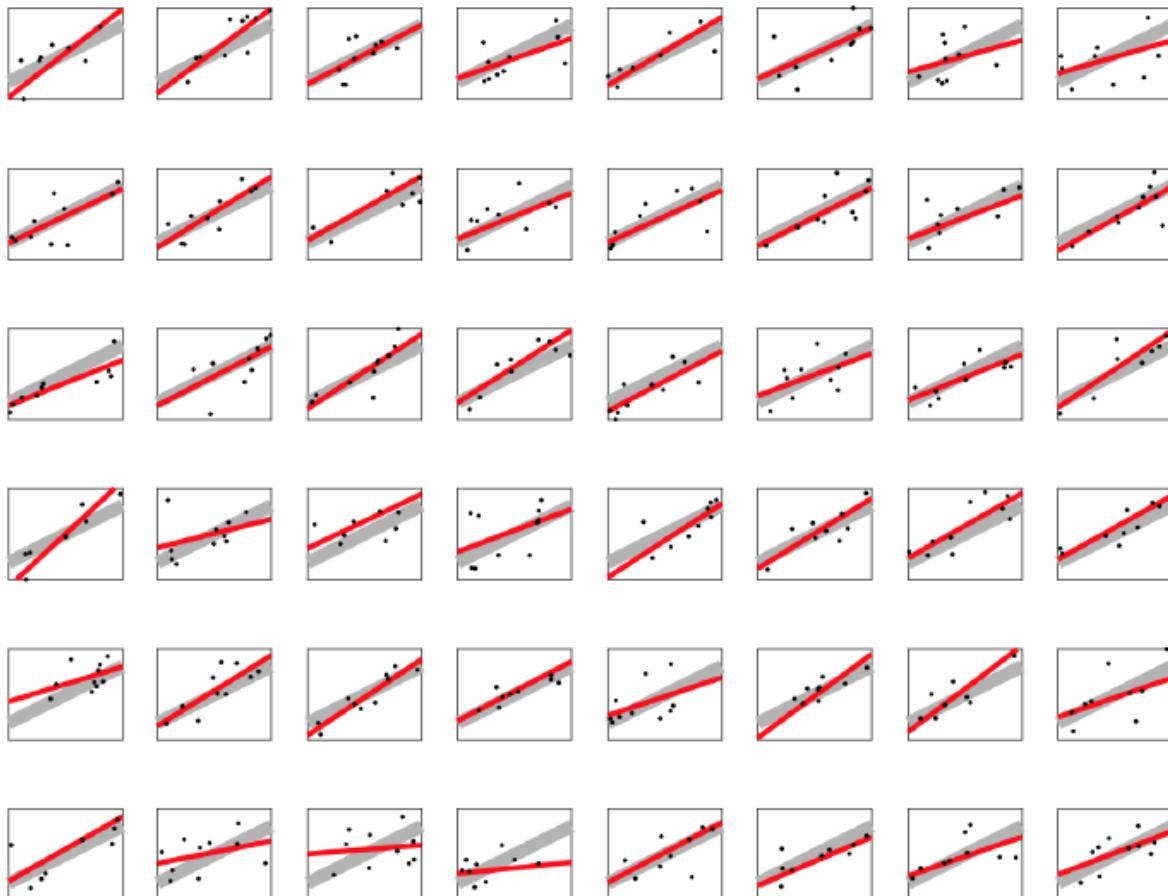


Bias and Variance in Model Selection:  $m = 0$

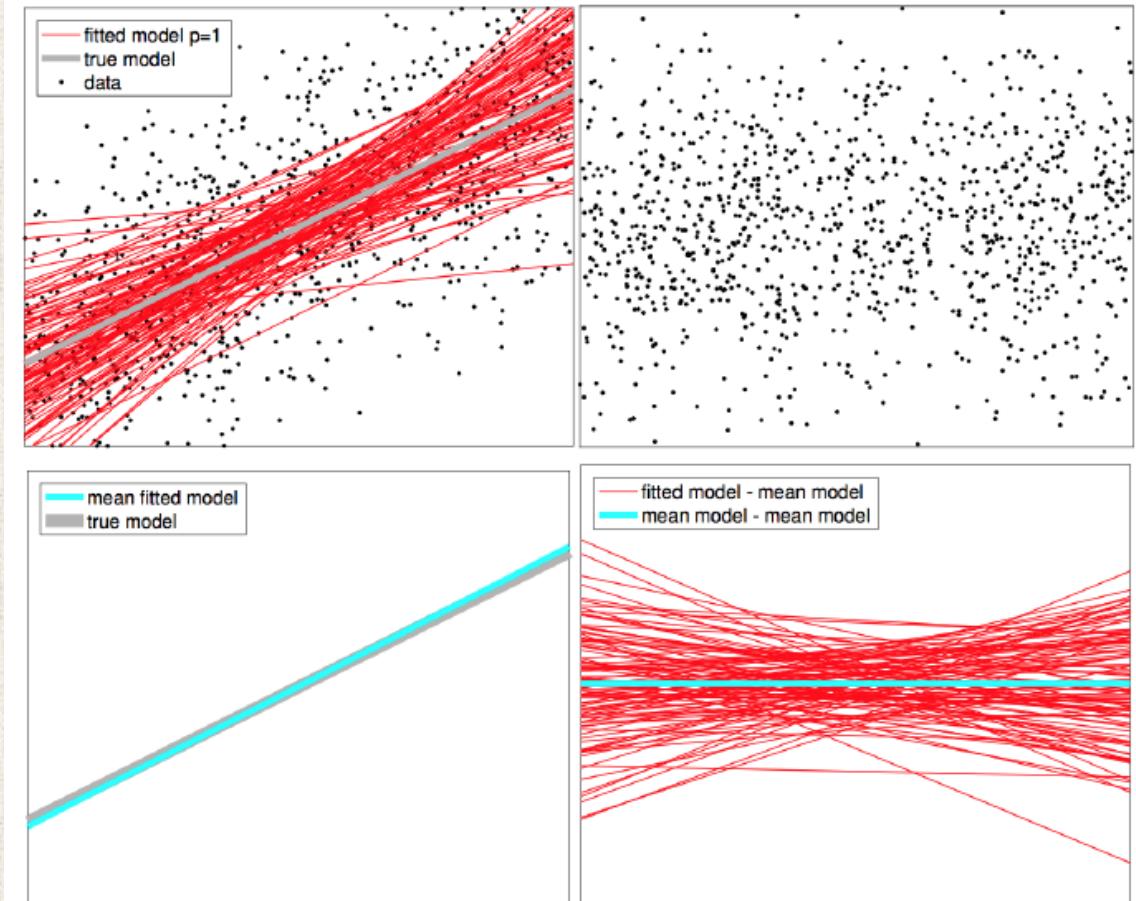


# An experiment on Bias-Variance tradeoff

Fitting A Model over Multiple Datasets:  $m = 1$

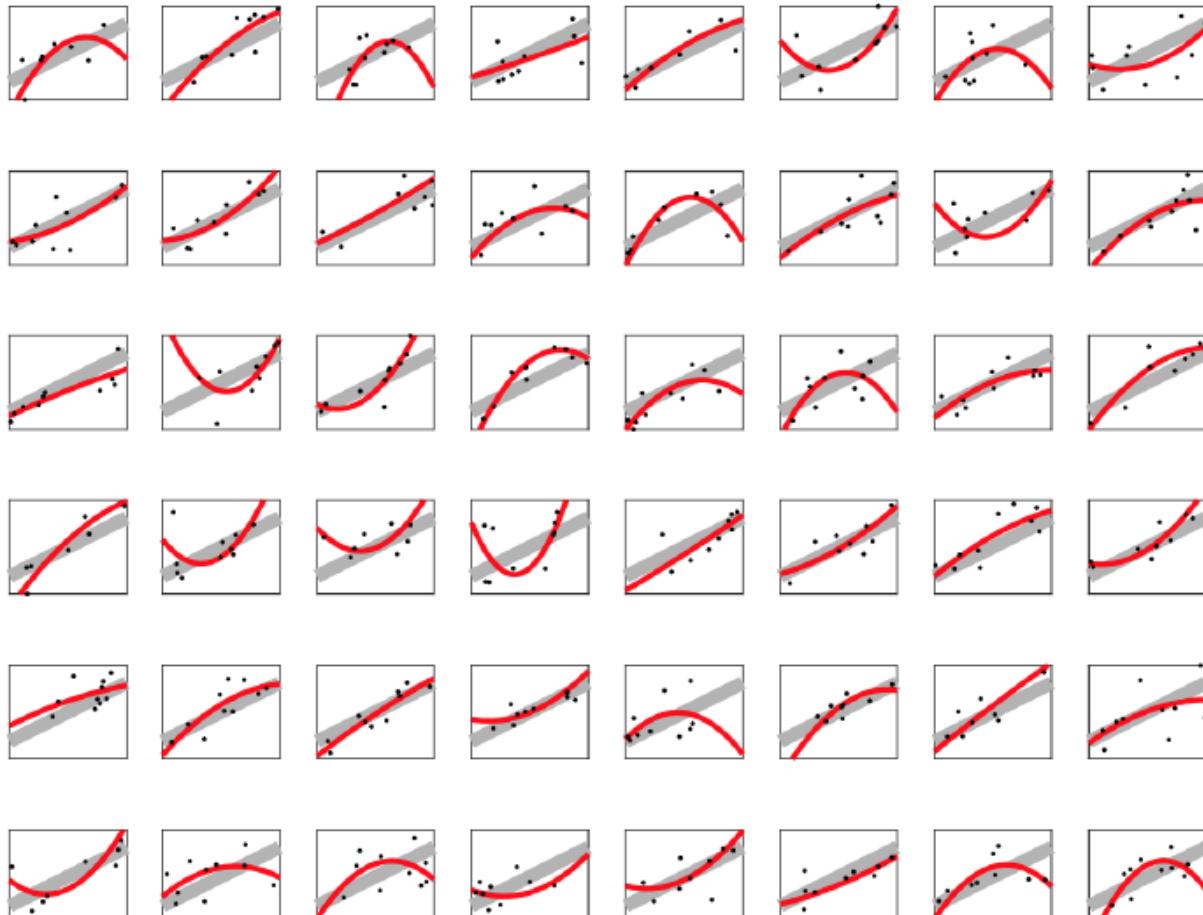


Bias and Variance in Model Selection:  $m = 1$

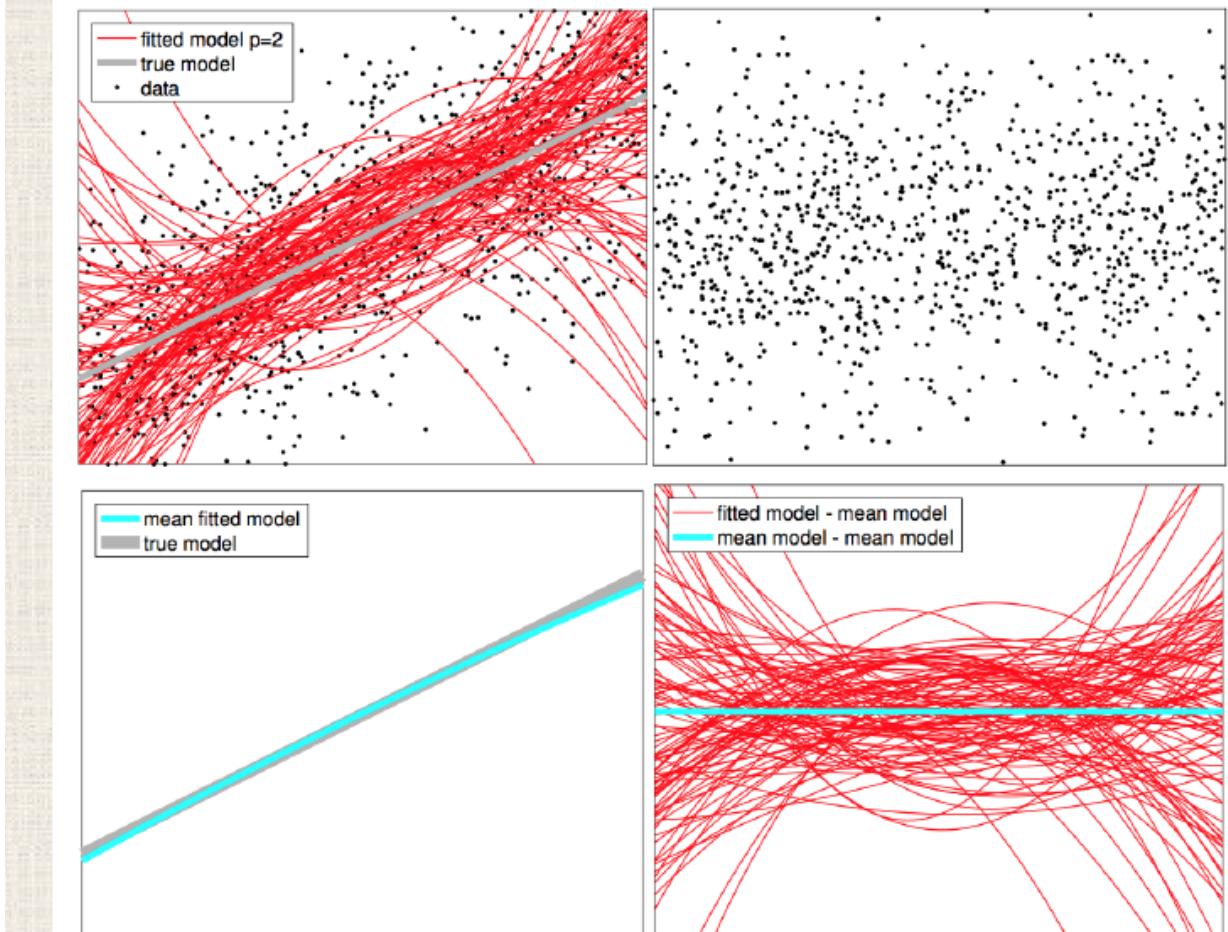


# An experiment on Bias-Variance tradeoff

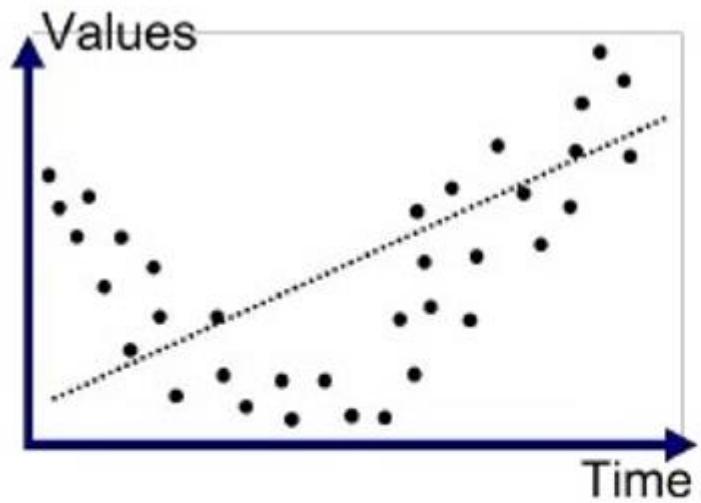
Fitting A Model over Multiple Datasets:  $m = 2$



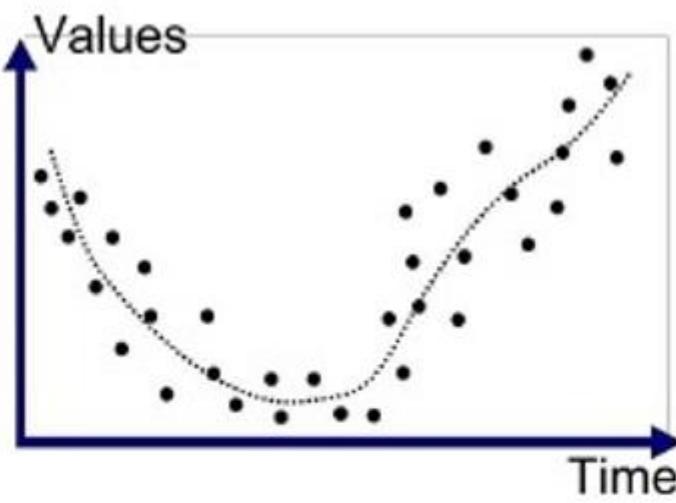
Bias and Variance in Model Selection:  $m = 2$



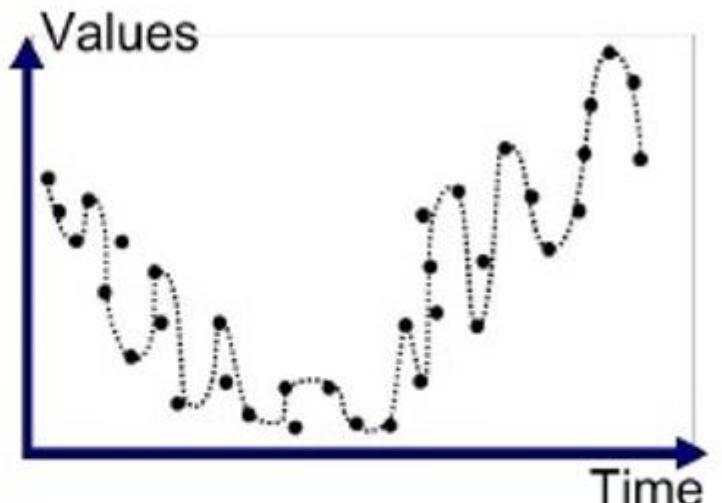
# Underfitting and Overfitting



Underfitted



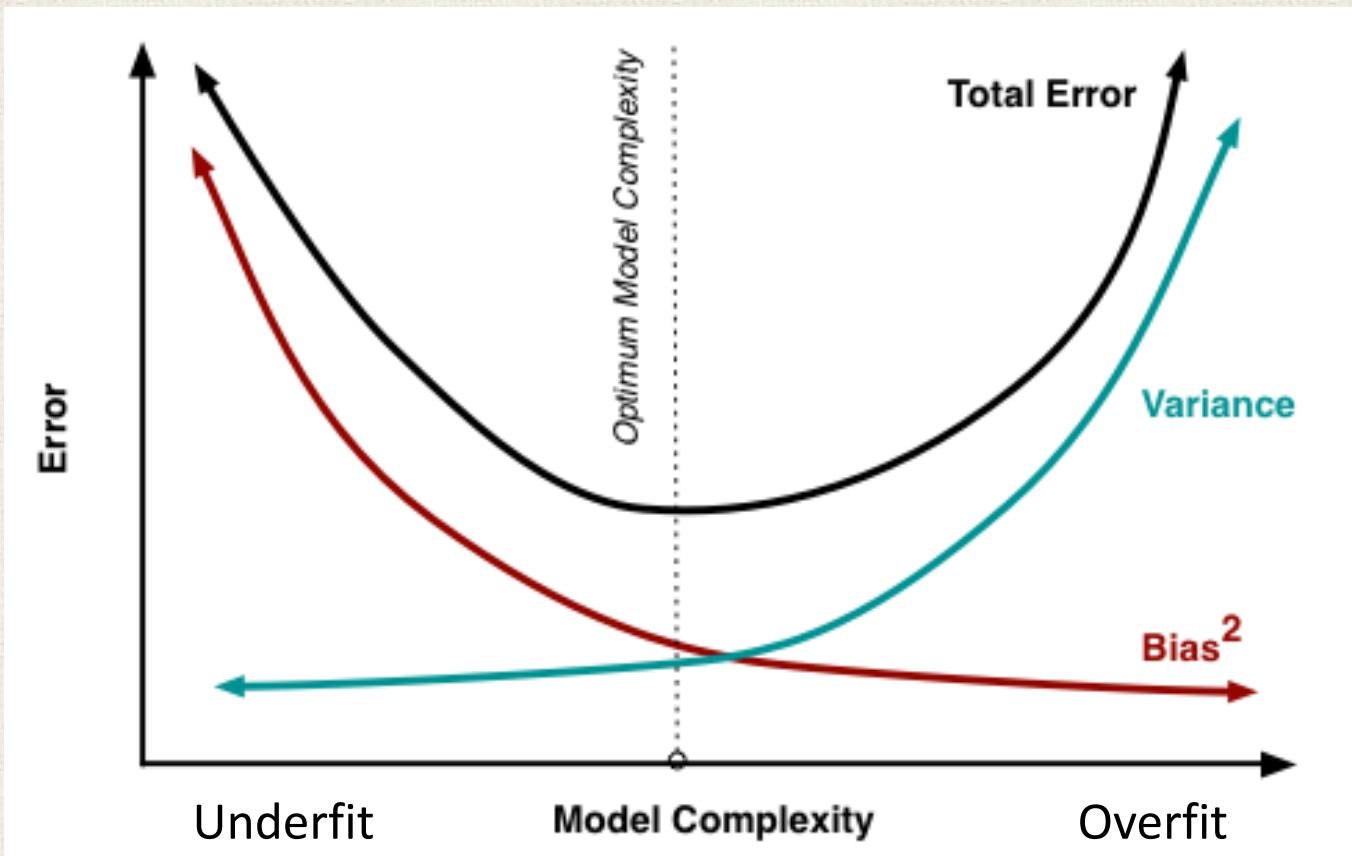
Good Fit/R robust



Overfitted

- **Overfitting** occurs when a statistical model or machine learning algorithm captures the noise of the data.
  - Intuitively, overfitting occurs when the model or algorithm fits the data too well.
  - Specifically, **overfitting model or algorithm shows low bias but high variance**
  - Overfitting is often a result of an **excessively complicated model**.
- **Underfitting** occurs when a statistical model or machine learning algorithm can not capture the underlying trend of the data.
  - Intuitively, underfitting occurs when the model or algorithm does not fit the data well enough.
  - Specifically, **underfitting model or algorithm shows low variance but high bias**.
  - Underfitting is often a result of an **excessively simple model**.
- **Both underfitting and overfitting lead to poor prediction on new data sets.**

# Bias-Variance tradeoff



- *Training error reflects bias but not variance. Test error reflects both.*
- *In practice, if the training error is much smaller than test error, then there is overfitting*
- *Decreasing the bias will increase the variance*
- *Decreasing the variance will increase the bias*
- There is *no escaping* this relationship between bias and variance in a given supervised machine learning problem.
- There is a tradeoff between these two concerns for your problems.
- Use validation or cross-validation to *tune some hyperparameters* of your algorithm to achieve the trade-off

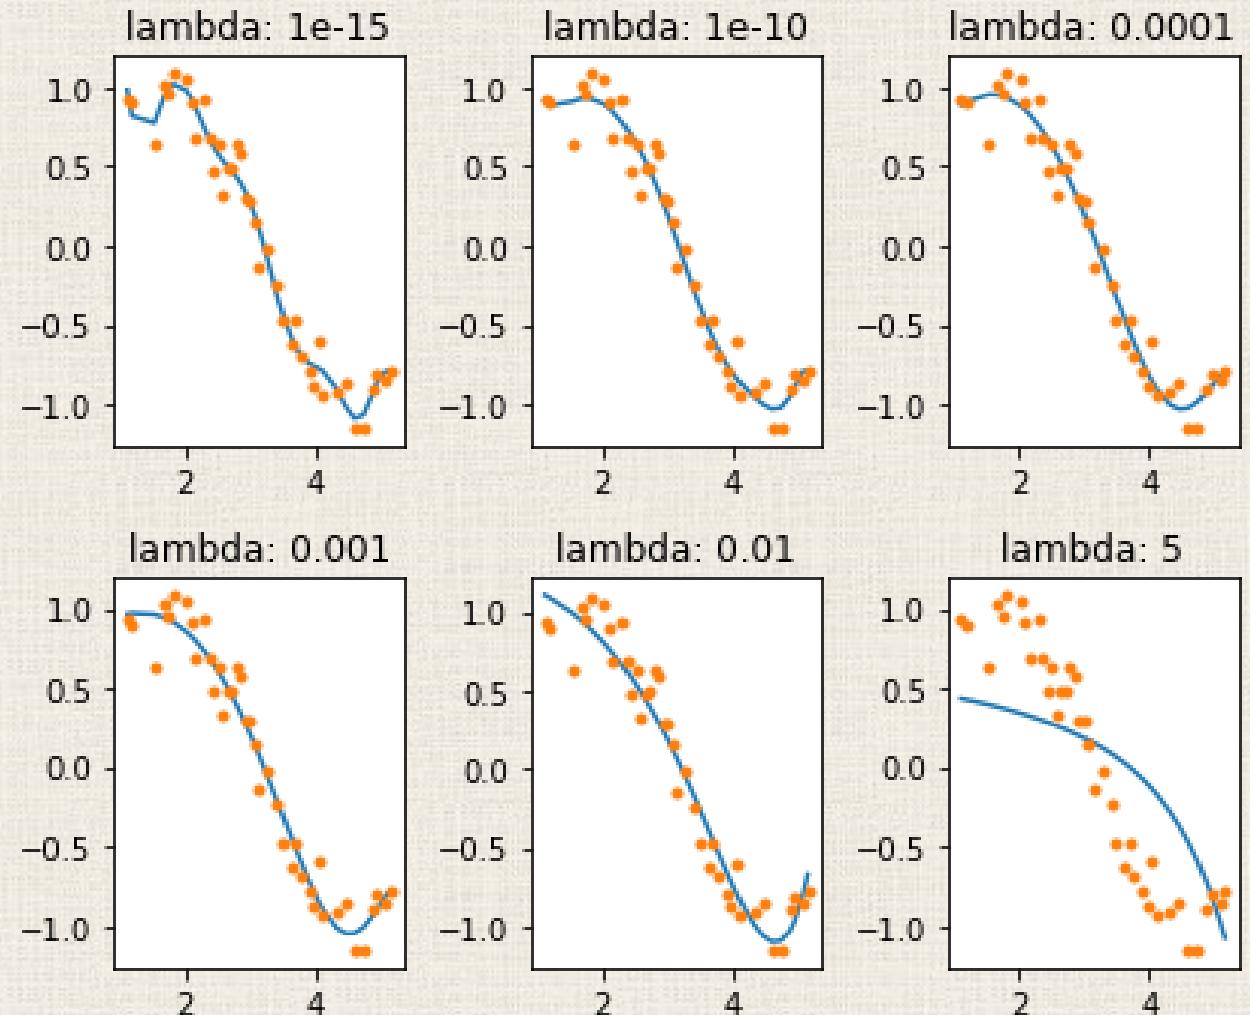
# Regularization to overcome Overfitting

---

$m = 15$

Loss function:(L2 regularization)

$$L(\mathbf{w}) = \|\Phi\mathbf{w} - \mathbf{y}\|^2 + \lambda \|\mathbf{w}\|^2$$



## **Examples of Bias and Variance**

- Consider a picture classification task. An ideal classifier (such as a human) may achieve nearly perfect performance in this task.
- Suppose your algorithm perform as following:
  - Training error = 1%
  - Test error = 11%

- Consider a picture classification task. An ideal classifier (such as a human) may achieve nearly perfect performance in this task.
- Suppose your algorithm perform as following:
  - Training error = 15%
  - Test error = 16%

- Consider a picture classification task. An ideal classifier (such as a human) may achieve nearly perfect performance in this task.
- Suppose your algorithm perform as following:
  - Training error = 15%
  - Test error = 30%

- Consider a picture classification task. An ideal classifier (such as a human) may achieve nearly perfect performance in this task.
- Suppose your algorithm perform as following:
  - Training error = 0.5%
  - Test error = 1%

- Consider a speed recognition task. Suppose that 14% of the audio clips have so much background noise that even a human can not recognize what was said. Hence the ***optimal error rate*** in this case is set as 14%.
- Suppose your algorithm perform as following:
  - Training error = 15%
  - Test error = 30%

- Unavoidable bias = 14%
- Avoidable bias = 1%
- Variance = 15%

- Consider a speed recognition task. Suppose that 14% of the audio clips have so much background noise that even a human can not recognize what was said. Hence the optimal error rate in this case is set as 14%.
- Suppose your algorithm perform as following:
  - Training error = 15%
  - Test error = 16%

## **Techniques to reduce avoidable bias**

- Increase the model size (complexity)
- Reduce or eliminate regularization

## **Techniques to reduce variance**

- Decrease the model size
- Add regularization
- Add more training data
- Feature selection to reduce number/type of input features
- Add early stop during the training process