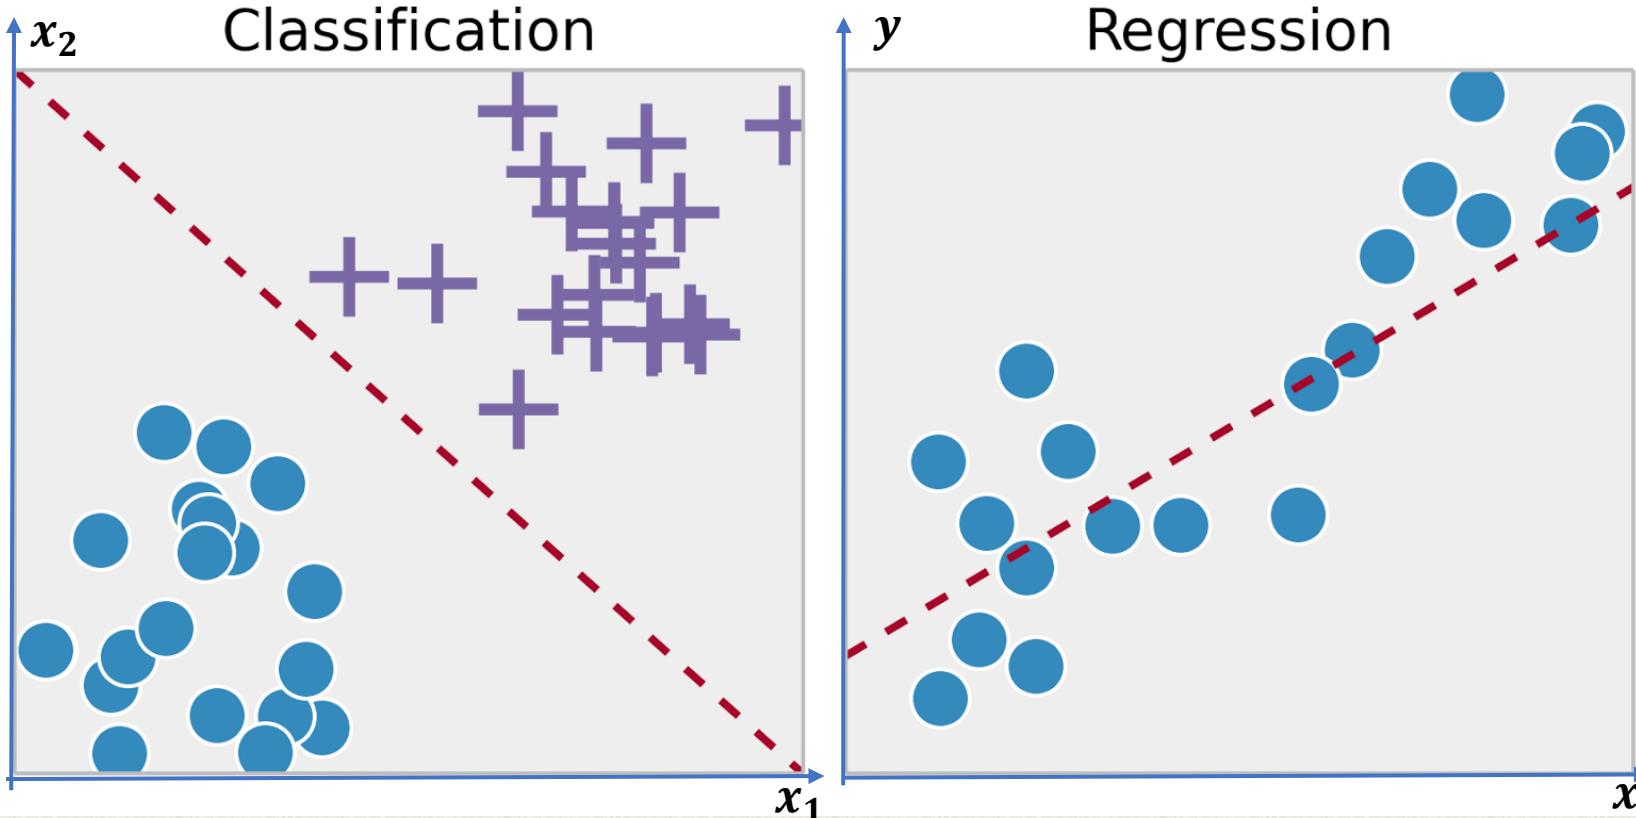


# Regression vs Classification



Predict a nominal class label

Predict a continuous quantity

# The Classification Problems

- ***The classification problems:*** given a training data set  $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$  of  $n$  points, where each data points  $x_i \in \mathcal{R}^d$  is paired with a known discrete class label  $y_i$  which represent one of  $K$  discrete classes  $\mathcal{C}_k$ , where  $k = 1, 2, \dots, K$ . The goal is to train a classifier which, when fed any arbitrary d-dimensional data point, classifies that data point as one of the  $K$  discrete classes.

- **Target values in classification problems:**

- In a binary classification problem,  $y \in \{0,1\}$  or  $\{-1,1\}$ , such that  $y = 1$  represents class  $\mathcal{C}_1$  and  $y = 0$ , or ( $y = -1$ ) represents class  $\mathcal{C}_2$
- In multi-class classification, it is convenient to use a ***1-of-K (one-hot) coding scheme*** where  $\mathbf{y}$  becomes a vector of length  $K$  such that if the class is  $\mathcal{C}_j$  then all the elements of  $\mathbf{y}$  are zero except the element  $y_j$ , which takes the value of 1. For instance, let  $K = 4$  classes, the pattern from class 2 would be given the target vector

$$\mathbf{y} = [0 \quad 1 \quad 0 \quad 0]^T$$

- The value of  $y_j$  can be interpreted as the probability that the class is  $\mathcal{C}_j$ .

## Target values in classification problems:

- We prefer one-hot encoding over representing  $y = 1, 2, 3, 4$  (**ordinal values**) for  $\mathcal{C}_1, \mathcal{C}_2, \mathcal{C}_3$  and  $\mathcal{C}_4$  (**nominal values**) is that the later implies order as  $4 > 3 > 2 > 1$ .
  - Let the target value of an input  $x$  be  $y = 4$ . Classifier#1 classifies the input as  $y_1 = 1$ , classifier#2 classifies the same input as  $y_2 = 3$ .
  - Since  $y - y_1 = 3 > y - y_2 = 1$ , this implies that classifier#2 is better than classifier#1. This is not true. Both classifiers are equally bad in the classification problem.
- **Mean squared error** is no longer a popular loss function in classification problem.

- ***Binary classification (dichotomizer)*** ( $K = 2$ ) and ***multi-class classification*** ( $K > 2$ )
  - Many applications are binary classification problems
  - Some classification methods can only handle binary classification problems (Perceptron, SVM, etc.)
  - Some classification methods can naturally handle multi-class classification problems (nearest neighbor, ANN, decision trees, etc.)
  - In many cases, a multi-class classification problem can be decomposed as multiple binary classification problems.

- ***Binary classification for multi-class classification***

- Multi-class classification problem can be turned into multiple binary classification problems.
- One-vs-Rest (OvR)

For example, a multi-class classification problem with three classes: “red”, “blue” and “green”. This could be divided into three binary classification problems as follows:

Binary classifier #1: red vs (blue and green)

Binary classifier #2: blue vs (red and green)

Binary classifier #3: green vs (red and blue)

- One-vs-One (OvO)

Using the same example, the problem is split into multiple binary classification problems as follows:

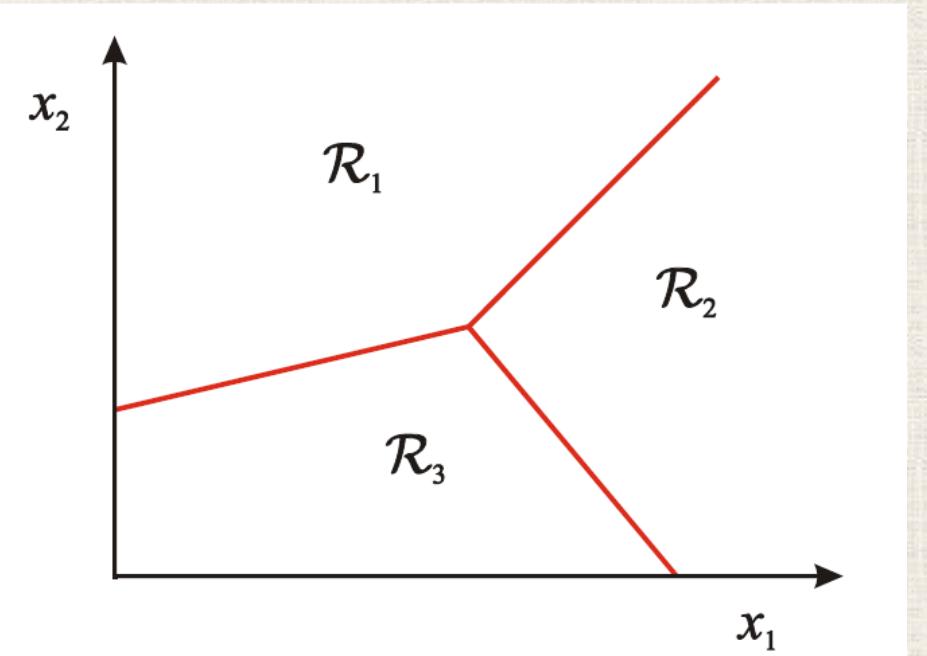
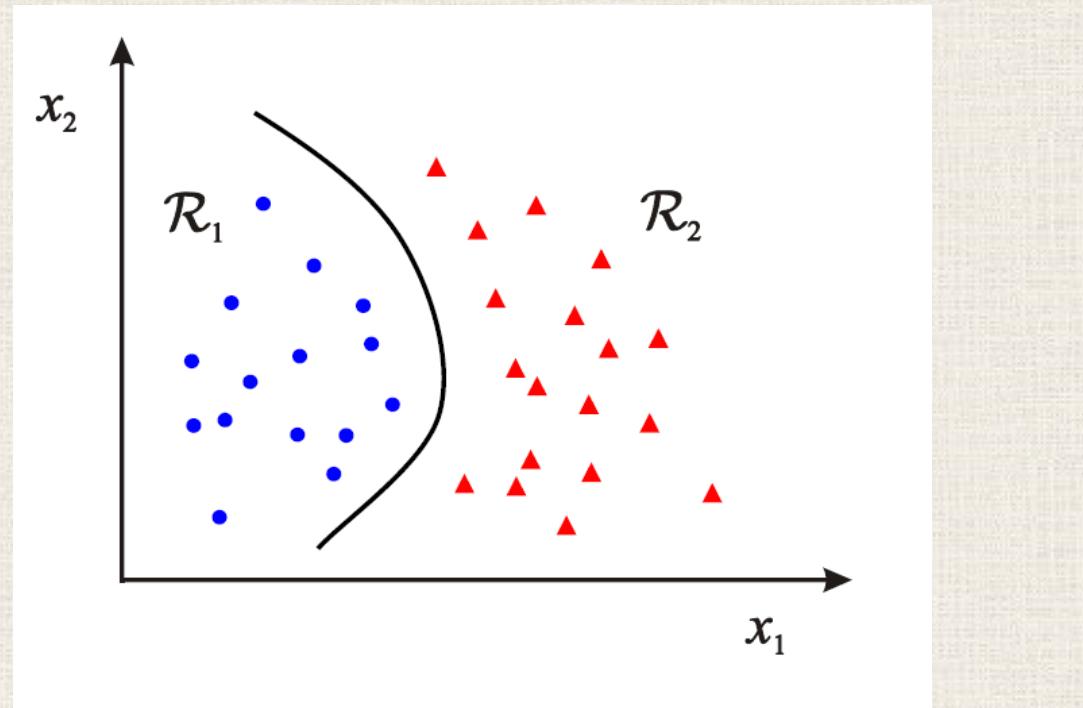
Binary classifier #1: red vs blue

Binary classifier #2: red vs green

Binary classifier#3: blue vs green

# Decision regions and Decision boundaries

- In most common scenario, the classes are taken to be disjoint so that ***each input is assigned to one and only one class (no class label ambiguity!)***
- The input space is thereby divided into ***decision regions***  $R_k$  such that a point falling in  $R_k$  is assigned to class  $\mathcal{C}_k$ . The boundaries between these decision regions are called ***decision boundaries***.



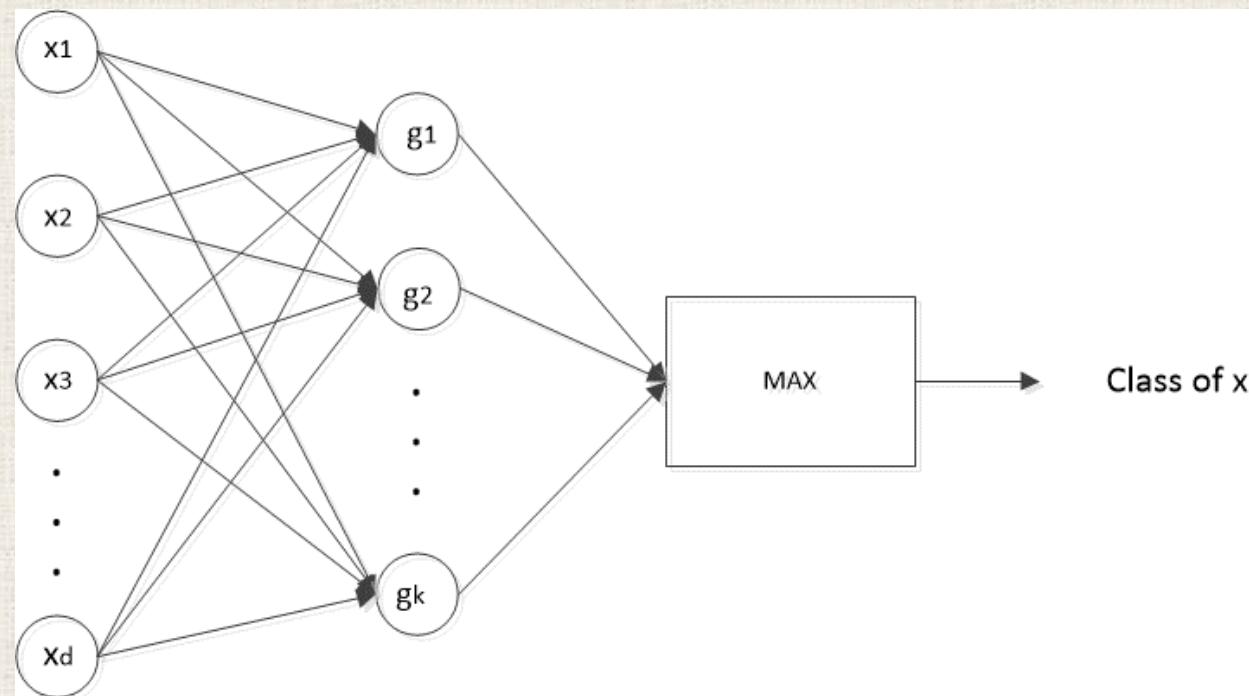
## Discriminant Functions

- One of the most useful way to represent classifiers is in terms of a set of ***discriminant functions***  $g_i(\mathbf{x}), i = 1, \dots, K$ .
- The classifier is said to assign a feature vector  $\mathbf{x}$  to class  $\mathcal{C}_i$  if

$$g_i(\mathbf{x}) > g_j(\mathbf{x}) \quad \text{for all } j \neq i$$

# Discriminant Functions

- The classifier is viewed as a network that computes  $K$  discriminant functions and selects the category corresponding to the largest discriminant.



- One popular example:

$$g_i(\mathbf{x}) = P(\mathcal{C}_i | \mathbf{x})$$

The maximum discriminant function becomes the maximum posterior probability

- ***Discriminant function is not unique!***

$g_i(\mathbf{x})$  can be replaced by  $f(g_i(\mathbf{x}))$  if  $f(\cdot)$  is ***monotonically increasing function!***

$$g_i(\mathbf{x}) = P(\mathcal{C}_i | \mathbf{x}) = \frac{P(\mathbf{x} | \mathcal{C}_i)P(\mathcal{C}_i)}{\sum_{i=1}^K P(\mathbf{x} | \mathcal{C}_i)P(\mathcal{C}_i)}$$

$$g_i(\mathbf{x}) = P(\mathbf{x} | \mathcal{C}_i)P(\mathcal{C}_i)$$

$$g_i(\mathbf{x}) = \ln P(\mathbf{x} | \mathcal{C}_i) + \ln P(\mathcal{C}_i)$$

- ***Discriminant function for two-class case:***

Instead of using two discriminant functions  $g_1$  and  $g_2$  and assigning  $x$  to  $\mathcal{C}_1$  if  $g_1(x) > g_2(x)$ , it is more common to define a single discriminant function

$$g(x) = g_1(x) - g_2(x)$$

And to use the following decision rule:

Assign  $x$  to  $\mathcal{C}_1$  if  $g(x) > 0$ , otherwise assign  $x$  to  $\mathcal{C}_2$

***Examples:***

$$g(x) = P(\mathcal{C}_1|x) - P(\mathcal{C}_2|x)$$

$$g(x) = \ln\left(\frac{P(x|\mathcal{C}_1)}{P(x|\mathcal{C}_2)}\right) + \ln\left(\frac{P(\mathcal{C}_1)}{P(\mathcal{C}_2)}\right)$$

# **Performance Metrics for Classifiers**

- **Confusion Matrix:** a natural way to represent classification results.

A binary classifier case:

classifier		Actual class	
		positive	negative
Predicted class	positive	10	2
	negative	3	5

- 13 positive instances and 7 negative instances
- 10 positive instances are classified as positive
- 3 positive instances are classified as negative
- 5 negative instances are classified as negative
- 2 negative instances are classified as positive
- Diagonal numbers represent number of instance classified correctly
- Off-diagonal numbers represent classification errors

- A multi-class classifier case: (iris data set)

classifier		Actual class		
		setosa	versicolor	virginica
Predicted class	setosa	13	0	0
	versicolor	0	10	6
	virginica	0	0	9

## A binary classifier case:

Classifier		Actual class	
		positive	negative
Predicted class	Positive	TP (hit)	FP (false alarm)
	negative	FN (miss)	TN (correct rejection)

- **Condition positive (P):** the number of real positive instances in the data.
- **Condition negative (N):** the number of real negative instances in the data.
- **True positive (TP) (hit):** the number of real positive instances that are classified as positive.
- **True negative (TN) (correct rejection):** the number of negative instances that are classified as negative.
- **False positive (FP) (false alarm) (type I error):** number of real negative instances that are classified as positive.
- **False negative (FN) (miss) (type II error):** number of positive instances that are classified as negative.

# *Accuracy and Error Rate:*

Classifier		Actual class	
		positive	negative
Predicted class	Positive	TP	FP
	negative	FN	TN

- $\text{Accuracy} = \frac{\# \text{ of instances classified correctly}}{\# \text{ of instances tested}}$

$$= \frac{TP + TN}{TP + FP + FN + TN}$$

- $\text{Error rate} = 1 - \text{Accuracy}$

$$= \frac{FP + FN}{TP + FP + FN + TN}$$

## Example: Find the accuracy of the following two classifiers

Classifier#1		Actual class	
		positive	negative
Predicted class	Positive	90	10
	negative	10	90

$$\text{Accuracy} = \frac{90+90}{200} = 90\%$$

Classifier#2		Actual class	
		positive	negative
Predicted class	Positive	10	10
	negative	10	170

$$\text{Accuracy} = \frac{10+170}{200} = 90\%$$

# Problems with Accuracy and Error rate

- ***The problem with class imbalance:***

Accuracy simply treat all instances the same. Accuracy maybe fine if you deal with balanced (or approximated balanced) data sets. It can be misleading when deal with imbalanced data sets.

- In real applications domain, imbalanced data sets are the norm!  
For example, cancer diagnose, intrusion detection, defects detection, etc.  
The positive examples to negative examples ratio can be 1:10 or less!
- Need more metrics for classifier performance evaluation!!!

# Precision

Classifier		Actual class	
		positive	negative
Predicted class	Positive	TP	FP
	negative	FN	TN

- $$\text{Precision} = \frac{TP}{TP+FP}$$
- Precision measures what portion of instances that are classified as positive are actually positive.
- Precision gives us information about its performance with respect to False Positive (false alarm).
- If the focus of the application is on **minimizing False Positive (false alarm)**, we should make precision close to 100%.
- For example, spam email detector.

Classifier#1		Actual class	
		positive	negative
Predicted class	Positive	90	10
	negative	10	90

$$\text{Accuracy} = \frac{90+90}{200} = 90\%$$

$$\text{Precision} = \frac{90}{90+10} = 90\%$$

Classifier#2		Actual class	
		positive	negative
Predicted class	Positive	10	10
	negative	10	170

$$\text{Accuracy} = \frac{10+170}{200} = 90\%$$

$$\text{Precision} = \frac{10}{10+10} = 50\%$$

# Recall or Sensitivity

Classifier		Actual class	
		positive	negative
Predicted class	Positive	TP	FP
	negative	FN	TN

- **Recall (sensitivity)** =  $\frac{TP}{TP+FN}$
- **Recall** measures what portion of positive instances that are classified as positive
- **Recall** gives us information about a classifier performance with respect to **False Negative (miss)**
- *If the application focus more on minimizing False Negative (miss), we should make recall close to 100% and without precision being too bad.*
- *For example, cancer diagnose, information retrieval*

# Performance comparison of two classifiers

Classifier#2		Actual class	
		positive	negative
Predicted class	Positive	10	10
	negative	10	170

$$\text{Recall} = \frac{10}{10+10} = 50\%$$

$$\text{precision} = \frac{10}{10+10} = 50\%$$

$$\text{Accuracy} = \frac{10+170}{200} = 90\%$$

Classifier#2		Actual class	
		positive	negative
Predicted class	Positive	20	20
	negative	0	160

$$\text{Recall} = \frac{20}{15+5} = 100\%$$

$$\text{precision} = \frac{20}{20+20} = 50\%$$

$$\text{Accuracy} = \frac{20+160}{200} = 90\%$$

## F-score: a combination of precision and recall

- $F_1$  score (balanced F score) is the **harmonic mean** of precision and recall

$$F_1 = \frac{2}{recall^{-1} + precision^{-1}} = \frac{2 * precision * recall}{precision + recall}$$

- $F_\beta$  score:  $\beta > 0$  is chosen such that recall is considered  $\beta$  times as important as precision

$$F_\beta = (1 + \beta^2) \frac{precision * recall}{(\beta^2) * precision + recall}$$

## Other metrics

- **True positive rate (TPR)(hit rate)** =  $\frac{TP}{TP+FN} = \frac{TP}{\text{total positives}} = \text{recall} = \text{sensitivity}$   
=  $1 - \text{false negative rate (FNR)(miss rate)}$
- **False positive rate (FPR)(false alarm rate)** =  $\frac{FP}{FP+TN} = \frac{FP}{\text{total negatives}}$   
=  $1 - \text{true negative rate (TNR)}$
- **Specificity** = **true negative rate (TNR)** =  $\frac{TN}{FP+TN}$

# Generative vs. Discriminative Classifier

- There are two main types of classification models: ***generative models and discriminative models***
- Posterior probability plays critical roles in classification problems
- ***Generative models*** have strong roots in probabilistic modeling. As in Bayes Theorem, generative models **explicitly** use prior probability and likelihood to calculate the posterior probabilities

$$\hat{y} = \underset{i}{\operatorname{argmax}} P(\mathcal{C}_i | x) = \underset{i}{\operatorname{argmax}} P(x | \mathcal{C}_i)P(\mathcal{C}_i)$$

e.g., Naïve Bayes, Hidden Markov Models (HMMs), Bayesian Networks, etc.

# Generative vs. Discriminative Classifier

- On the other hand, in ***discriminative models***, we bypass learning a generative model altogether and directly learn a decision boundary.
- **Discriminative models** are parameterized by weights that either form a posterior distribution without considering the prior or conditional distribution or directly form a hard decision boundary without considering any probabilities in the first place.

e.g., Perceptron, logistic regression, SVM, artificial neural networks (ANN), nearest neighbor, etc.

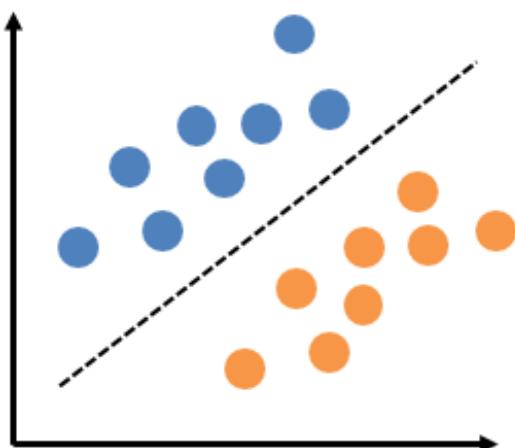
# Linearly separable data

## Binary linearly separable case:

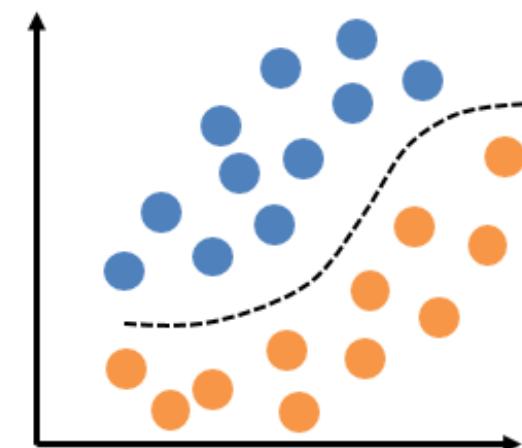
- Let's assume that we have two sets of points on ***2D Euclidian plane***, one set of points are colored blue and other colored red.
- These two sets are ***linearly separable*** if there exists at least one ***line*** in the plane with all the blue points on one side and all the red points on the other side.
- This definition generalizes to ***higher-dimensional Euclidian spaces*** if ***line*** is replaced by ***hyperplane***

## Binary Linearly separable data sets

Linear

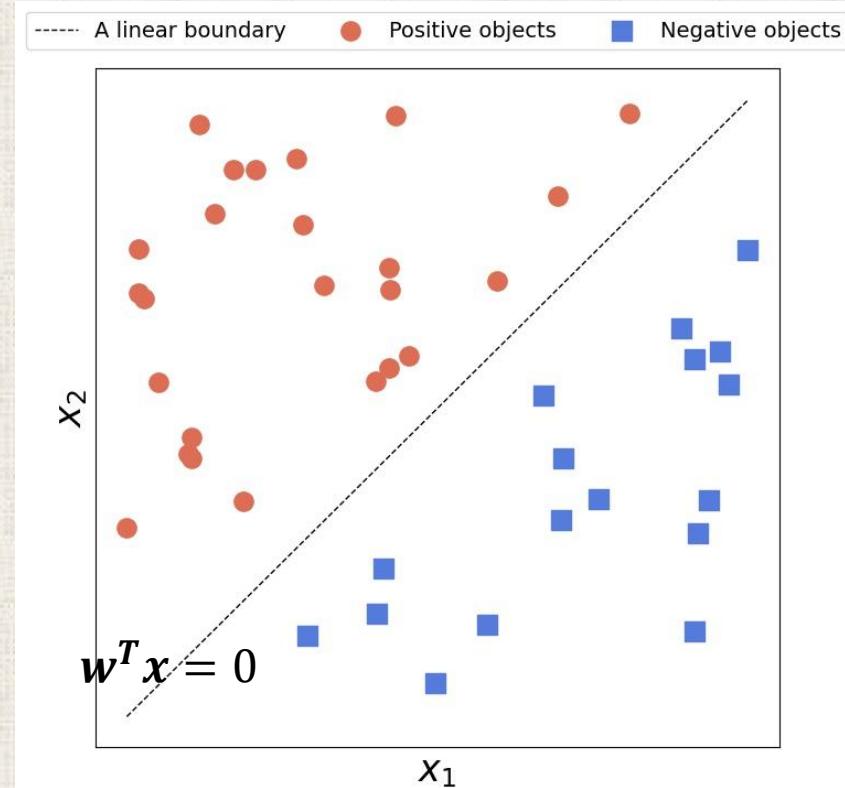


Nonlinear



# Equation of the linear boundary

- The equation of the hyperplane:  $\mathbf{w}^T \mathbf{x} = 0$ 
  - Any point  $\mathbf{x}$  on the hyperplane:  $\mathbf{w}^T \mathbf{x} = 0$
  - Any point  $\mathbf{x}$  above the hyperplane:  $\mathbf{w}^T \mathbf{x} > 0$
  - Any point  $\mathbf{x}$  below the hyperplane:  $\mathbf{w}^T \mathbf{x} < 0$



# The Perceptron Model (a discriminative model)

## Binary linearly separable case:

- The ***perceptron classifier*** classifies an input  $x$  into one of two classes:  $C_1$  and  $C_2$ . The target values used in this problem is  $y = 1$  for  $C_1$  and  $y = -1$  for  $C_2$ .
- The ***discriminant function*** is in a ***linear form***:  $g(x) = \mathbf{w}^T \mathbf{x}$ ,

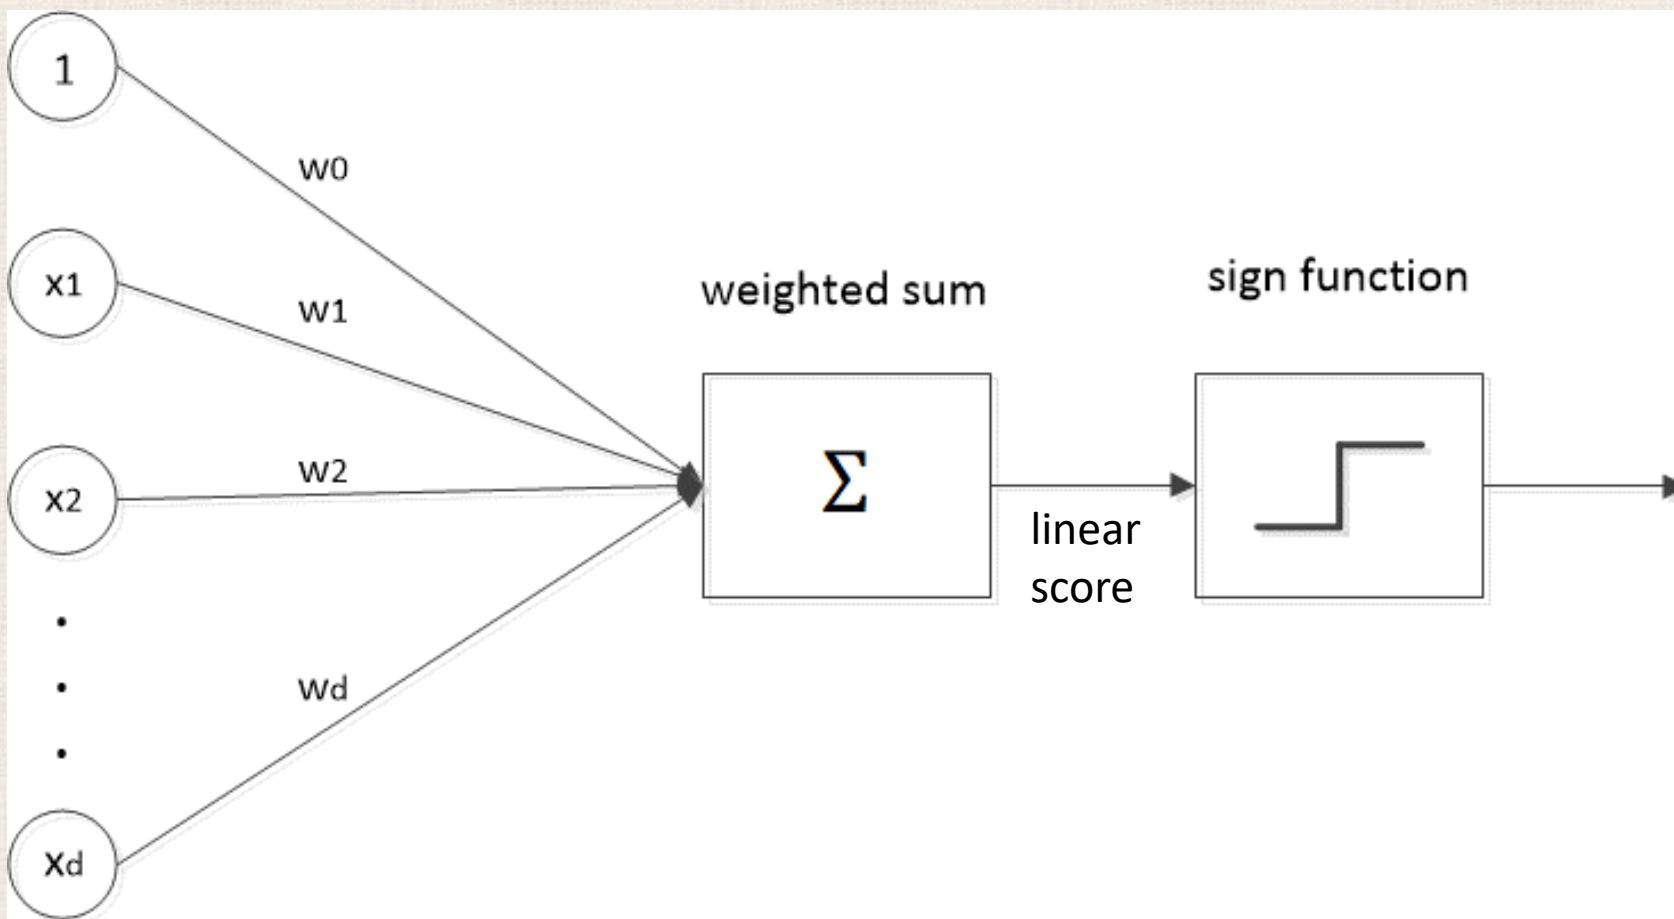
$$\mathbf{w}^T \mathbf{x}_i > 0, \quad \mathbf{x}_i \text{ in } C_1$$

$$\mathbf{w}^T \mathbf{x}_i \leq 0, \quad \mathbf{x}_i \text{ in } C_2$$

- $\mathbf{w}$  is the weight vector to be determined.
- $\mathbf{w}^T \mathbf{x}$  is called the ***linear score*** of the input  $x$

# The Perceptron Model - a graphical explanation

$$y(\mathbf{x}) = f(\mathbf{w}^T \mathbf{x})$$



## Binary linearly separable case:

- Notice that in this algorithm, we are seeking the weight vector  $\mathbf{w}$  such that for any input  $\mathbf{x}_i$  in the given data set,

$$\mathbf{w}^T \mathbf{x}_i > 0, \quad \mathbf{x}_i \text{ in } C_1 \ (y_i = 1)$$

$$\mathbf{w}^T \mathbf{x}_i \leq 0, \quad \mathbf{x}_i \text{ in } C_2 \ (y_i = -1)$$

This is equivalent to

$$\mathbf{w}^T \mathbf{x}_i y_i \geq 0, \quad \text{any } \mathbf{x}_i \in D$$

Where  $y_i \in \{1, -1\}$  is the target value of the input  $\mathbf{x}_i$ .  $\mathbf{w}^T \mathbf{x}_i y_i$  is the ***normalized linear score*** of input  $\mathbf{x}_i$ .

## ***The training algorithm of Perceptron Model***

How to find the weight vector that can completely classifies all the data points in the training data set.

### ***The Perceptron criterion : the error (loss) function:***

Hence, we have the following perceptron criterion error function:

$$E_p(\mathbf{w}) = - \sum_{i \in M} \mathbf{w}^T \mathbf{x}_i y_i$$

Where,  $M$  denote the ***set of all misclassified instances.***

The perceptron criterion associates zero error with any instance that is correctly classified; for a misclassified instance  $\mathbf{x}_i$ , it tries to minimize the quantity  $-\mathbf{w}^T \mathbf{x}_i y_i$

Show that for any **misclassified** instance,  $-\mathbf{w}^T \mathbf{x}_i y_i \geq 0$

- An instance  $\mathbf{x}_i$  is misclassified using the decision rule:

$$\mathbf{w}^T \mathbf{x}_i > 0, \quad \mathbf{x}_i \text{ in } C_1$$

$$\mathbf{w}^T \mathbf{x}_i \leq 0, \quad \mathbf{x}_i \text{ in } C_2$$

- Case 1:  $\mathbf{x}_i$  is a positive instance ( $y_i = +1$ ) but classified as a negative instance.  $\mathbf{w}^T \mathbf{x}_i \leq 0$ , hence,  $\mathbf{w}^T \mathbf{x}_i y_i \leq 0$
- Case 2:  $\mathbf{x}_i$  is a negative instance ( $y_i = -1$ ) but classified as a positive instance.  $\mathbf{w}^T \mathbf{x}_i > 0$ , hence,  $\mathbf{w}^T \mathbf{x}_i y_i < 0$

## The Gradient Descent Procedure

To find the weight vector  $w$  that minimize the *perceptron criterion error function*  $E_p(w)$ , we use the following ***gradient descent procedure***:

$$w^{k+1} = w^k - \eta * \nabla E_p(w^k)$$

Where,

- $k$  is an integer that indexes the steps of the algorithm;
- $\eta$  is a positive scale factor or learning rate that set the step size;
- $\nabla E_p(w)$  is the ***gradient*** vector of the *perceptron criterion error function*.

- To find the weight vector  $\mathbf{w}$  that minimize the *perceptron criterion error function*  $E_p(\mathbf{w})$ , we use the following ***gradient descent procedure***:

$$\mathbf{w}^{k+1} = \mathbf{w}^k - \eta \nabla E_p(\mathbf{w}^k)$$

- Gradient descent is based on the observation that if a multi-variable function  $E_p(\mathbf{w})$  is defined and differentiable in a neighborhood of a point  $\mathbf{w}$ , then  $E_p(\mathbf{w})$  decreases fastest if one goes from  $\mathbf{w}$  in the direction of the negative gradient of  $E_p$  at  $\mathbf{w}$ ,  $-\nabla E_p(\mathbf{w})$ .
- It follows that if  $\mathbf{w}^{k+1} = \mathbf{w}^k - \eta \nabla E_p(\mathbf{w}^k)$ , and for  $\eta$  ***small enough***, one starts with a guess  $\mathbf{w}^0$  for the local minimum of  $E_p$ , and consider the sequence  $\mathbf{w}^0, \mathbf{w}^1, \mathbf{w}^2, \dots$ , we will have a monotonic sequence

$$E_p(\mathbf{w}^0) \geq E_p(\mathbf{w}^1) \geq E_p(\mathbf{w}^2) \geq \dots,$$

so the sequence  $\mathbf{w}^k, k \geq 0$  converges to the desired local minimum.

- To find the weight vector  $\mathbf{w}$  that minimize the *perceptron criterion error function*  $E_p(\mathbf{w})$ , we use the following ***gradient descent procedure***:

$$\mathbf{w}^{k+1} = \mathbf{w}^k - \eta \nabla E_p(\mathbf{w}^k)$$

Where,  $k$  is an integer that indexes the steps of the algorithm;  $\eta^k$  is a positive scale factor or learning rate that set the step size;  $\nabla E_p(\mathbf{w})$  is the gradient vector of the *perceptron criterion error function*.

$$\nabla E_p(\mathbf{w}) = - \sum_{i \in M} \mathbf{x}_i y_i$$

- Substitute this into weight update expression, we obtain:

$$\mathbf{w}^{k+1} = \mathbf{w}^k + \eta \sum_{i \in M} \mathbf{x}_i y_i$$

# The Perceptron training Algorithms

*When should the weight vector be updated?*

- *Batch Training Algorithm*
- *Sequential Training Algorithm*

# Some terminologies about gradient descent algorithm

- ***Batch size*** is the number of training instances to work through before updating the weight vector
- Training dataset can be divided into one or more batches
- An ***epoch*** is one cycle through the entire training dataset
  - It usually takes more than a few epochs for a gradient descent algorithm to converge
- ***Batch gradient descent:*** batch size == N
- ***Sequential (stochastic) gradient descent:*** batch size == 1 + shuffling the training data set every epoch

## **Batch Training Algorithm:**

**Begin** initialize  $\mathbf{w}$  with random values, and  $k(\# \text{ of epoch}) \leftarrow 0$

Batch size =  $N$

**do**       $k \leftarrow k + 1$

$M_k = \{\}, j = 0$  (index of data points)

**do**  $j \leftarrow j + 1$

                Classify  $x_j$  using current  $\mathbf{w}$

                if  $x_j$  is misclassified then append  $x_j, y_j$  to  $M_k$

**until**  $j = N$

$\mathbf{w} \leftarrow \mathbf{w} + \eta \sum_{i \in M_k} \mathbf{x}_i y_i$

**until**     $M_k == \{\}$

**return**  $\mathbf{w}$

**End**

## ***Sequential Training Algorithm:***

***Begin*** initialize  $\mathbf{w}$  with random values, and  $k \leftarrow 0$  (index of points in the training dataset)

***do***       $k \leftarrow k \bmod N$

    Classify  $\mathbf{x}_k$  using current  $\mathbf{w}$

    if  $\mathbf{x}_k$  is misclassified,  $\mathbf{w} = \mathbf{w} + \eta \mathbf{x}_k y_k$

$k \leftarrow k + 1$

***until*** all instances in the training data set are properly classified

return  $\mathbf{w}$

**End**

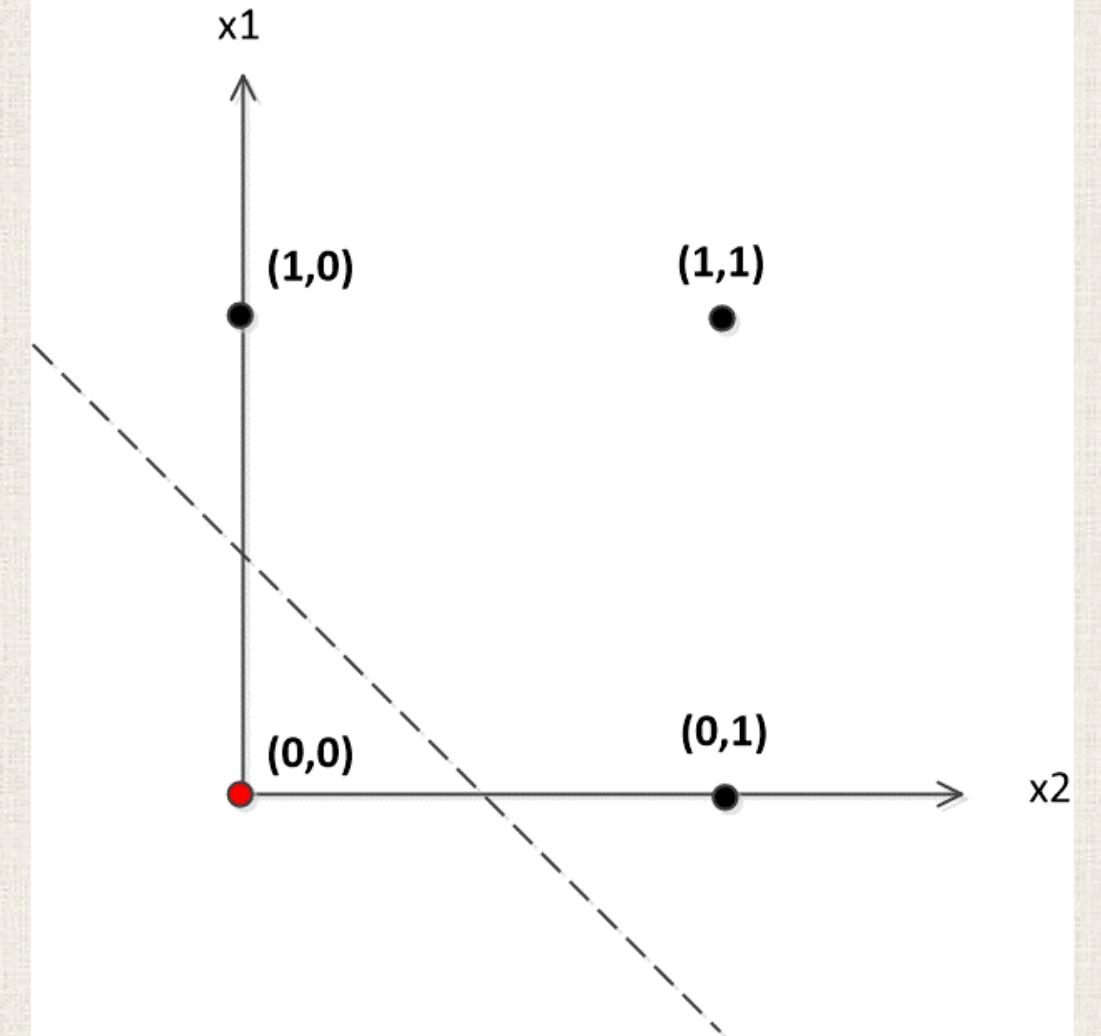
$N$  is the size of the training dataset

The logic ***OR*** problem:

$$y = x_1 \text{ or } x_2$$

$$x_1, x_2, y \in \{0,1\}$$

$x_1$	$x_2$	$y$
0	0	0
0	1	1
1	0	1
1	1	1

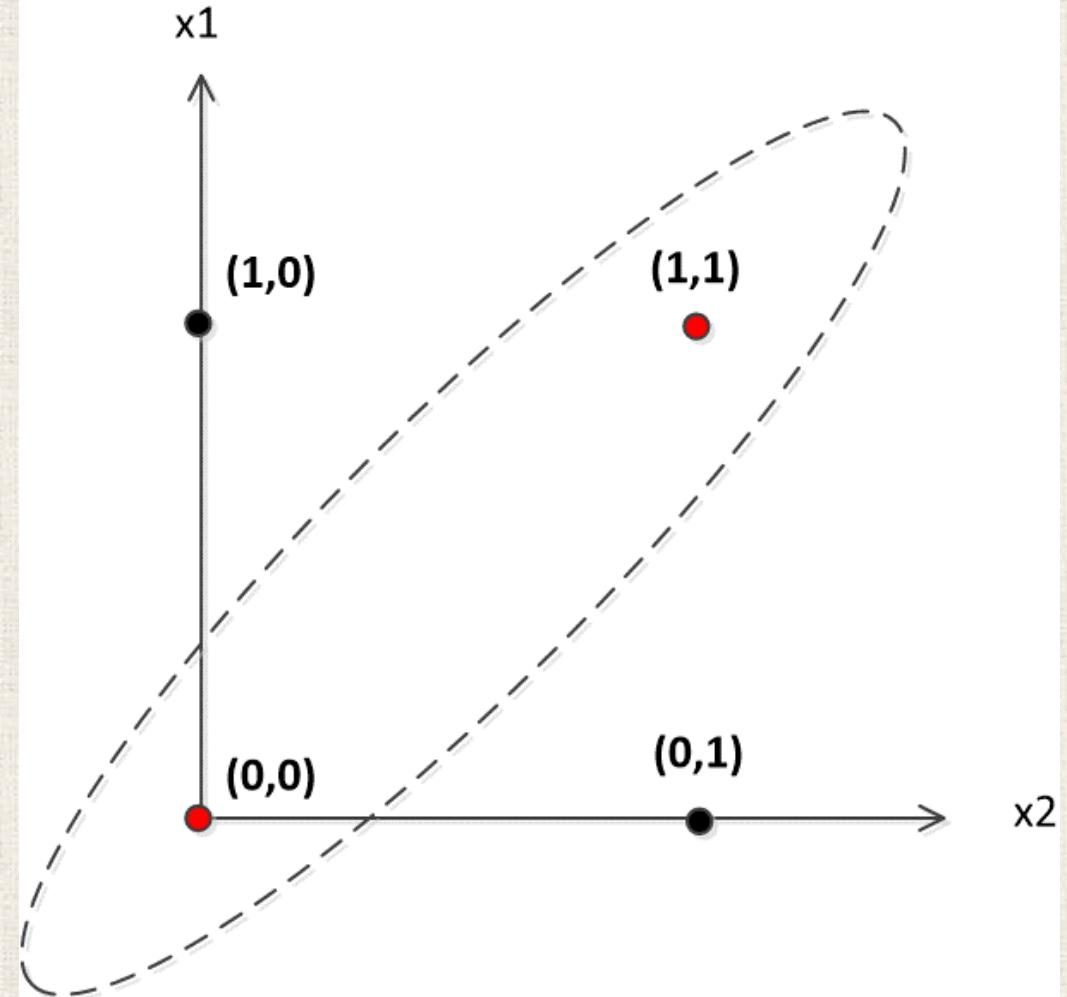


The logic ***XOR*** (exclusive ***OR***) problem:

$$y = x_1 \oplus x_2$$

$$x_1, x_2, y \in \{0,1\}$$

$x_1$	$x_2$	$y$
0	0	0
0	1	1
1	0	1
1	1	0



# Training a Perceptron for the logic OR problem (Sequential training)

The weight vector is  $\mathbf{w} = [w_0 \quad w_1 \quad w_2]^T$

the extended feature vector is  $\mathbf{x} = [1 \quad x_1 \quad x_2]^T$

Let's set the initial value of  $\mathbf{w}$  as  $\mathbf{w}^0 = [0 \quad 0 \quad 0]^T$

The learning rate is chosen as  $\eta = 1$

***Sequential weight update*** rule:

$$\mathbf{w}^{k+1} = \mathbf{w}^k + x_i y_i \quad \text{if } x_i \text{ is misclassified}$$

$x_1$	$x_2$	$y$
0	0	0 (-1)
0	1	1
1	0	1
1	1	1

$x_1$	$x_2$	$y$
0	0	0 (-1)
0	1	1
1	0	1
1	1	1

**Step 1:** feed the first training instance  $\mathbf{x}_1 = [1 \quad 0 \quad 0]^T, \quad y_1 = -1$

$$(\mathbf{w}^0)^T \mathbf{x}_1 = [0 \quad 0 \quad 0] \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} = 0, \quad \hat{y}_1 = sign((\mathbf{w}^0)^T \mathbf{x}_1) = -1, \quad correct$$

Then, no weight update

$$\mathbf{w}^1 = \mathbf{w}^0$$

$x_1$	$x_2$	$y$
0	0	0 (-1)
0	1	1
1	0	1
1	1	1

**Step2:** feed the second training instance  $\mathbf{x}_2 = [1 \ 0 \ 1]^T$ ,  $y_2 = 1$

$$(\mathbf{w}^1)^T \mathbf{x}_2 = [0 \ 0 \ 0] \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix} = 0, \quad \hat{y}_2 = sign((\mathbf{w}^1)^T \mathbf{x}_2) = -1, \quad missclassified$$

Then,

$$\mathbf{w}^2 = \mathbf{w}^1 + \mathbf{x}_2 y_2 = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix} + \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix}$$

$x_1$	$x_2$	$y$
0	0	0 (-1)
0	1	1
1	0	1
1	1	1

**Step3:** feed the third training instance  $\mathbf{x}_3 = [1 \quad 1 \quad 0]^T$ ,  $y_3 = 1$

$$(\mathbf{w}^2)^T \mathbf{x}_3 = [1 \quad 0 \quad 1] \begin{bmatrix} 1 \\ 1 \\ 0 \end{bmatrix} = 1, \quad \hat{y}_3 = \text{sign}((\mathbf{w}^2)^T \mathbf{x}_3) = 1, \quad \text{correct}$$

Then, no weight update

$$\mathbf{w}^3 = \mathbf{w}^2$$

$x_1$	$x_2$	$y$
0	0	0 (-1)
0	1	1
1	0	1
1	1	1

**Step4:** feed the fourth training instance  $\mathbf{x}_4 = [1 \ 1 \ 1]^T$ ,  $y_4 = 1$

$$\hat{y}_4(\mathbf{w}^3)^T \mathbf{x}_4 = [1 \ 0 \ 1] \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} = 2, \quad \hat{y}_4 = sign((\mathbf{w}^3)^T \mathbf{x}_4) = 1, \quad correct$$

Then,

$$\mathbf{w}^4 = \mathbf{w}^3$$

$x_1$	$x_2$	$y$
0	0	0 (-1)
0	1	1
1	0	1
1	1	1

**Step5:** feed the first training instance  $\mathbf{x}_1 = [1 \ 0 \ 0]^T$ ,  $y_1 = -1$

$$(\mathbf{w}^4)^T \mathbf{x}_1 = [1 \ 0 \ 1] \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} = 1, \quad \hat{y}_1 = sign((\mathbf{w}^4)^T \mathbf{x}_1) = 1, \quad missclassified$$

Then,

$$\mathbf{w}^5 = \mathbf{w}^4 + \mathbf{x}_1 y_1 = \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix} - \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$$

$x_1$	$x_2$	$y$
0	0	0 (-1)
0	1	1
1	0	1
1	1	1

**Step6:** feed the second training instance  $\mathbf{x}_2 = [1 \ 0 \ 1]^T$ ,  $y_2 = 1$

$$(\mathbf{w}^5)^T \mathbf{x}_2 = [0 \ 0 \ 1] \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix} = 1, \quad \hat{y}_2 = \text{sign}\left((\mathbf{w}^5)^T \mathbf{x}_2\right) = 1, \quad \text{correct}$$

Then, no weight update

$$\mathbf{w}^6 = \mathbf{w}^5$$

$x_1$	$x_2$	$y$
0	0	0 (-1)
0	1	1
1	0	1
1	1	1

**Step7:** feed the third training instance  $\mathbf{x}_3 = [1 \quad 1 \quad 0]^T, \quad y_3 = 1$

$$(\mathbf{w}^6)^T \mathbf{x}_3 = [0 \quad 0 \quad 1] \begin{bmatrix} 1 \\ 1 \\ 0 \end{bmatrix} = 0, \quad \hat{y}_3 = sign((\mathbf{w}^6)^T \mathbf{x}_3) = -1, \quad missclassified$$

Then,

$$\mathbf{w}^7 = \mathbf{w}^6 + \mathbf{x}_3 y_3 = \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix} + \begin{bmatrix} 1 \\ 1 \\ 0 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix}$$

$x_1$	$x_2$	$y$
0	0	0 (-1)
0	1	1
1	0	1
1	1	1

**Step8:** feed the fourth training instance  $\mathbf{x}_4 = [1 \quad 1 \quad 1]^T$ ,  $y_4 = 1$

$$(\mathbf{w}^7)^T \mathbf{x}_4 = [1 \quad 1 \quad 1] \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} = 3, \quad \hat{y}_4 = sign((\mathbf{w}^7)^T \mathbf{x}_4) = 1, \quad correct$$

Then, no weight update

$$\mathbf{w}^8 = \mathbf{w}^7$$

$x_1$	$x_2$	$y$
0	0	0 (-1)
0	1	1
1	0	1
1	1	1

**Step9:** feed the first training instance  $\mathbf{x}_1 = [1 \quad 0 \quad 0]^T, \quad y_1 = -1$

$$(\mathbf{w}^8)^T \mathbf{x}_1 = [1 \quad 1 \quad 1] \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} = 1, \quad \hat{y}_1 = sign((\mathbf{w}^8)^T \mathbf{x}_1) = 1, \quad missclassified$$

Then,

$$\mathbf{w}^9 = \mathbf{w}^8 + \mathbf{x}_1 y_1 = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} - \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} = \begin{bmatrix} 0 \\ 1 \\ 1 \end{bmatrix}$$

$x_1$	$x_2$	$y$
0	0	0 (-1)
0	1	1
1	0	1
1	1	1

**Step10:** feed the second training instance  $\mathbf{x}_2 = [1 \quad 0 \quad 1]^T, y_2 = 1$

$$(\mathbf{w}^9)^T \mathbf{x}_2 = [0 \quad 1 \quad 1] \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix} = 1, \quad \hat{y}_2 = sign((\mathbf{w}^9)^T \mathbf{x}_2) = 1, \quad correct$$

Then, no weight update

$$\mathbf{w}^{10} = \mathbf{w}^9$$

$x_1$	$x_2$	$y$
0	0	0 (-1)
0	1	1
1	0	1
1	1	1

**Step11:** feed the third training instance  $\mathbf{x}_3 = [1 \ 1 \ 0]^T$ ,  $y_3 = 1$

$$(\mathbf{w}^{10})^T \mathbf{x}_3 = [0 \ 1 \ 1] \begin{bmatrix} 1 \\ 1 \\ 0 \end{bmatrix} = 1, \quad \hat{y}_3 = \text{sign}((\mathbf{w}^{10})^T \mathbf{x}_3) = 1, \quad \text{correct}$$

Then, no weight update

$$\mathbf{w}^{11} = \mathbf{w}^{10}$$

$x_1$	$x_2$	$y$
0	0	0 (-1)
0	1	1
1	0	1
1	1	1

**Step12:** feed the fourth training instance  $\mathbf{x}_4 = [1 \ 1 \ 1]^T, y_4 = 1$

$$(\mathbf{w}^{11})^T \mathbf{x}_4 = [0 \ 1 \ 1] \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} = 2, \quad \hat{y}_4 = \text{sign}((\mathbf{w}^{11})^T \mathbf{x}_4) = 1, \quad \text{correct}$$

Then, no weight update

$$\mathbf{w}^{12} = \mathbf{w}^{11}$$

x1	x2	y
0	0	0 (-1)
0	1	1
1	0	1
1	1	1

**Step13:** feed the first training instance  $\mathbf{x}_1 = [1 \ 0 \ 0]^T$ ,  $y_1 = -1$

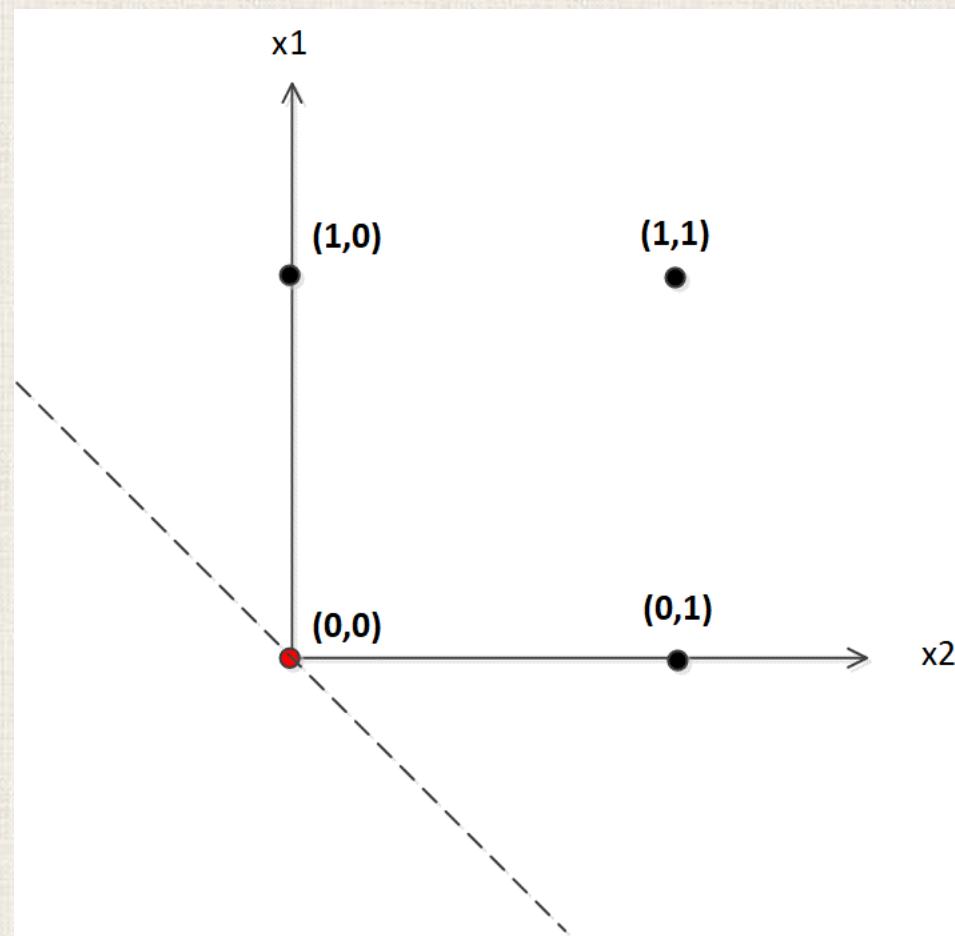
$$(\mathbf{w}^{12})^T \mathbf{x}_1 = [0 \ 1 \ 1] \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} = 0, \quad \hat{y}_1 = sign((\mathbf{w}^{12})^T \mathbf{x}_1) = -1, \quad correct$$

Now, We have all the instances classified correctly using the weight:

$$\mathbf{w} = [0 \ 1 \ 1]^T$$

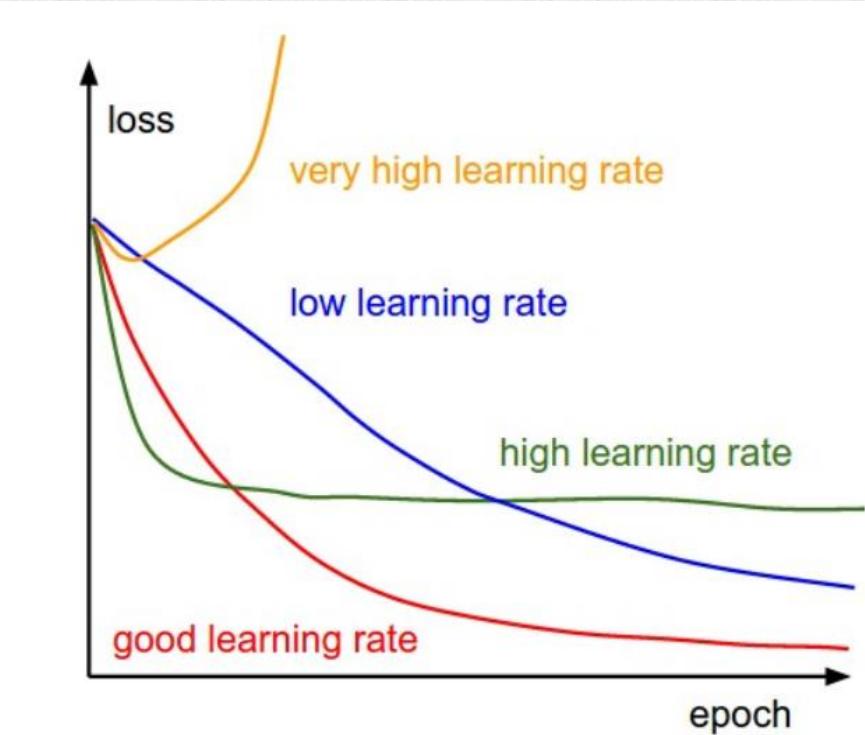
i.e.,

$$\mathbf{w}^T \mathbf{x} = x_1 + x_2$$



# Learning rate in the gradient descent algorithm

- if the learning rate  $\eta$  is too small, it would take long time for the weights to converge, and the algorithm becomes computationally expensive
- If the learning rate  $\eta$  is too large, the weights may fail to converge and overshoot the minimum
- The most commonly used learning rates are: 0.001, 0.01, 0.1
- Fixed learning rate vs changing learning rate



## Binary linearly separable case:

- Notice that in this algorithm, we are seeking the weight vector  $\mathbf{w}$  such that for an instance  $\mathbf{x}_i$

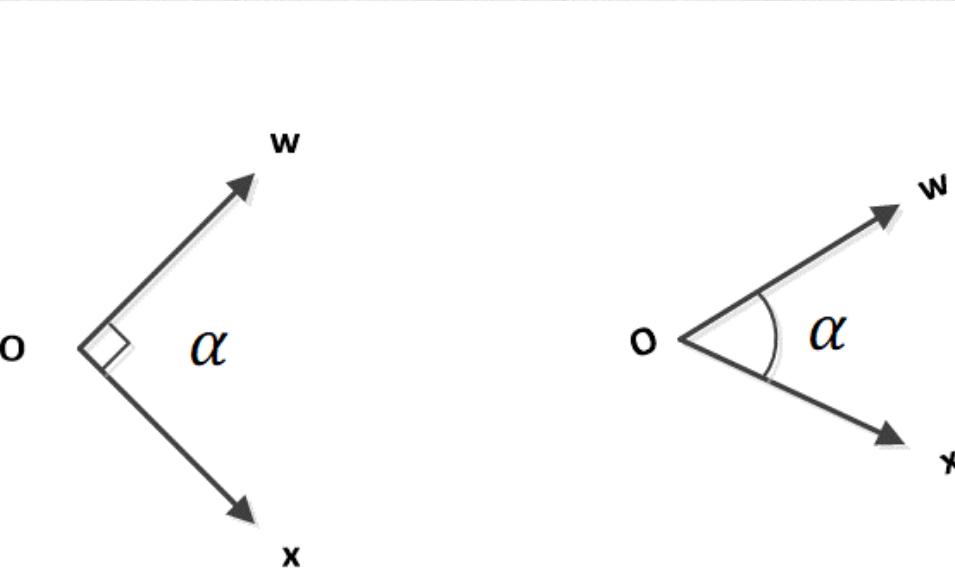
$$\mathbf{w}^T \mathbf{x}_i > 0, \quad \mathbf{x}_i \text{ in } C_1 \quad (y_i = 1)$$

$$\mathbf{w}^T \mathbf{x}_i \leq 0, \quad \mathbf{x}_i \text{ in } C_2 \quad (y_i = -1)$$

This is equivalent to

$$\mathbf{w}^T \mathbf{x}_i y_i \geq 0, \quad \text{any } \mathbf{x}_i$$

Where  $y_i \in \{1, -1\}$  is the target value of the input  $\mathbf{x}_i$ .  $\mathbf{w}^T \mathbf{x}_i y_i$  is the ***normalized linear score*** of input  $\mathbf{x}_i$ .



- Notice that

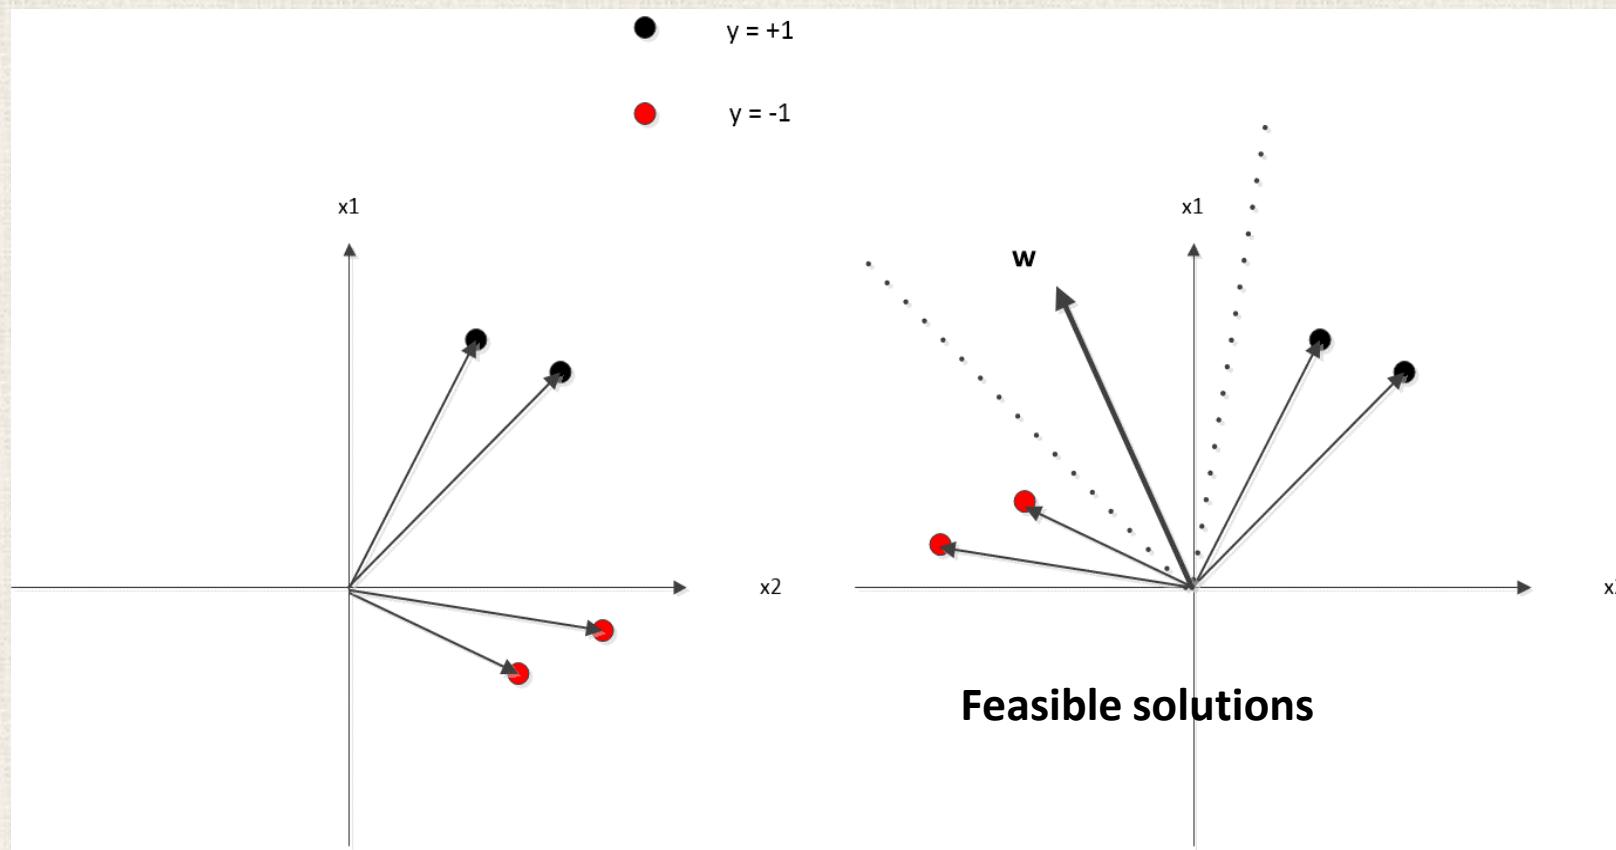
$$\mathbf{w}^T \mathbf{x} = \|\mathbf{w}\| \|\mathbf{x}\| \cos \alpha$$

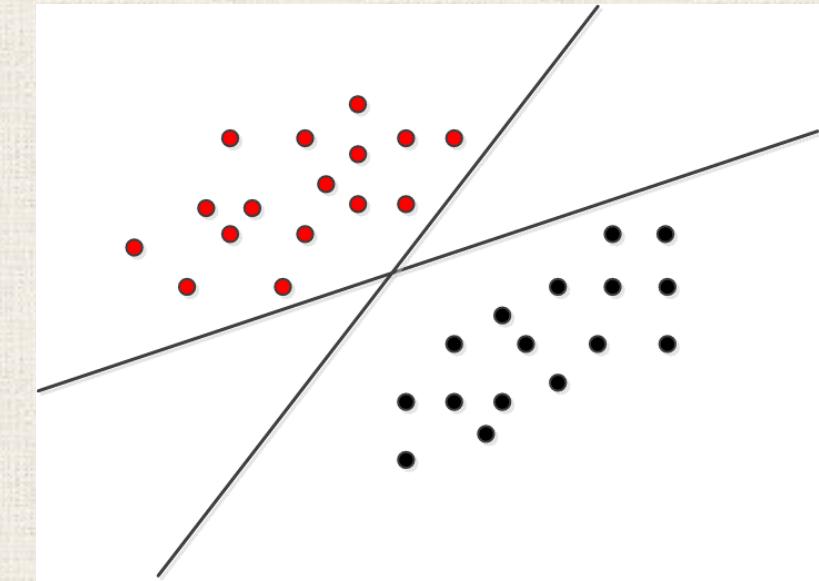
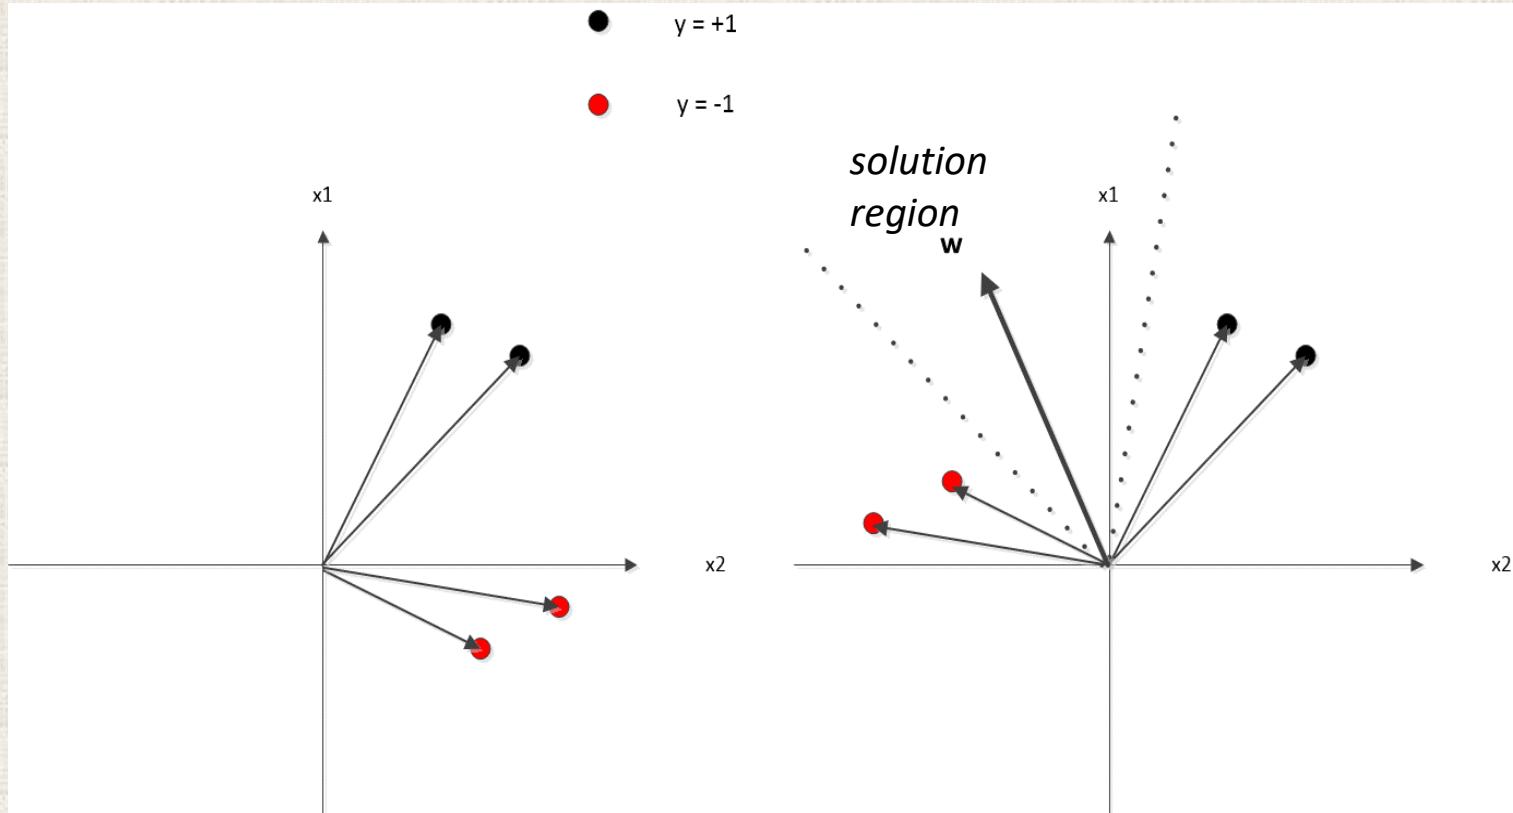
- When  $\alpha = \frac{\pi}{2}$ , we have  $\mathbf{w}^T \mathbf{x} = 0$ ,  $\mathbf{w}$  and  $\mathbf{x}$  are orthogonal to each other.
- When  $|\alpha| < \frac{\pi}{2}$ , we have  $\mathbf{w}^T \mathbf{x} > 0$

- we are seeking the weight vector  $w$  such that for an instance  $x_i$

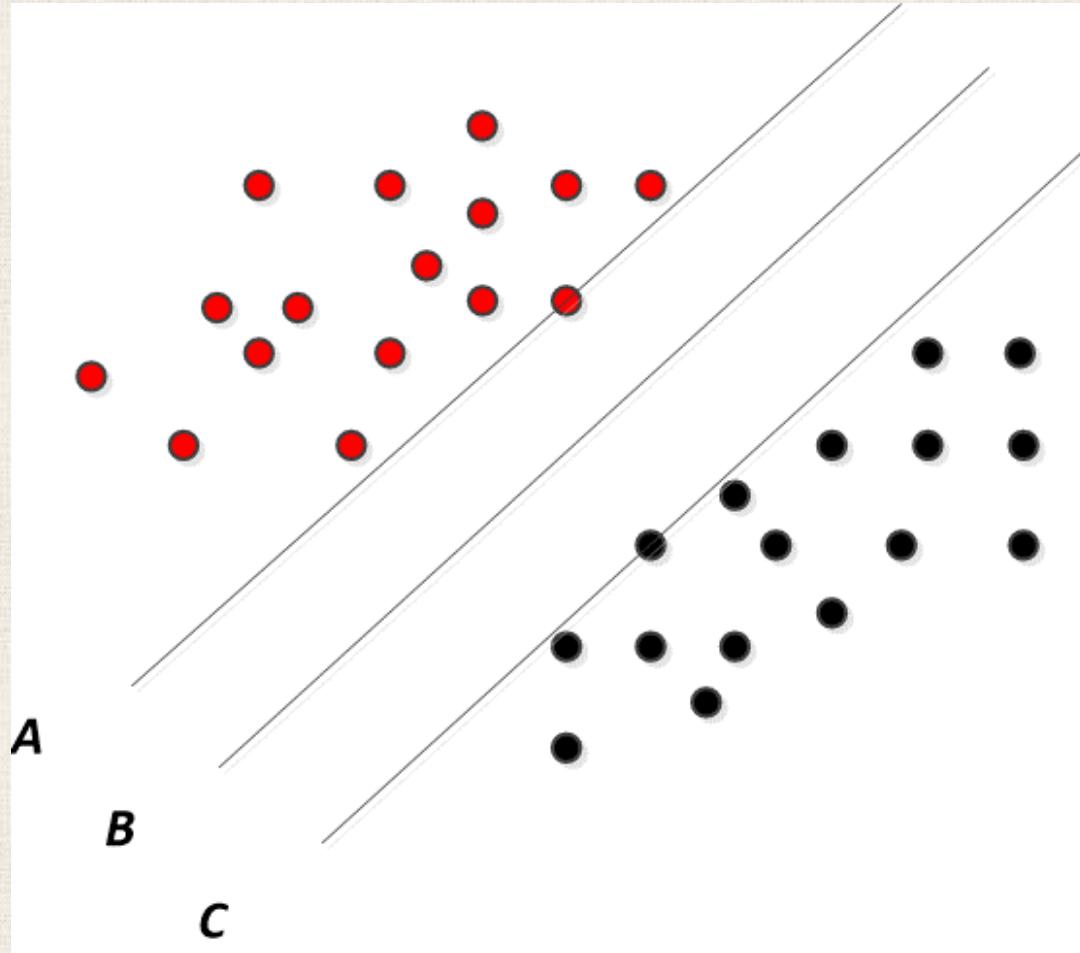
$$w^T x_i y_i \geq 0, \quad \text{any } x_i, i = 1, \dots, N$$

- Any weight vector  $w$  within the ***solution region*** can classify all the instances correctly.  $w$  is called a ***solution vector***.





- Any weight vector  $w$  within the *solution region* can classify all the instances correctly.  $w$  is called a *solution vector*.
- Solution vectors are not unique!

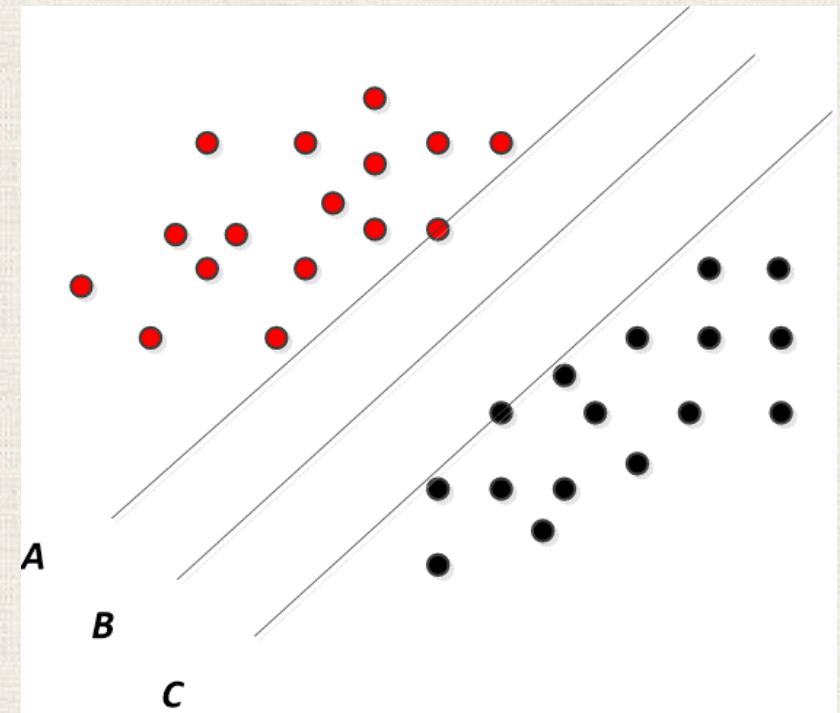


- There are several ways to impose additional requirements to constrain the solution vector.
- One possibility is to seek minimum length weight vector satisfying

$$\mathbf{w}^T \mathbf{x}_i y_i \geq b, \quad \text{for all } i$$

Where  $b > 0$  is a constant called ***margin***.

- ***The result solution is more likely to classify new test instances correctly***

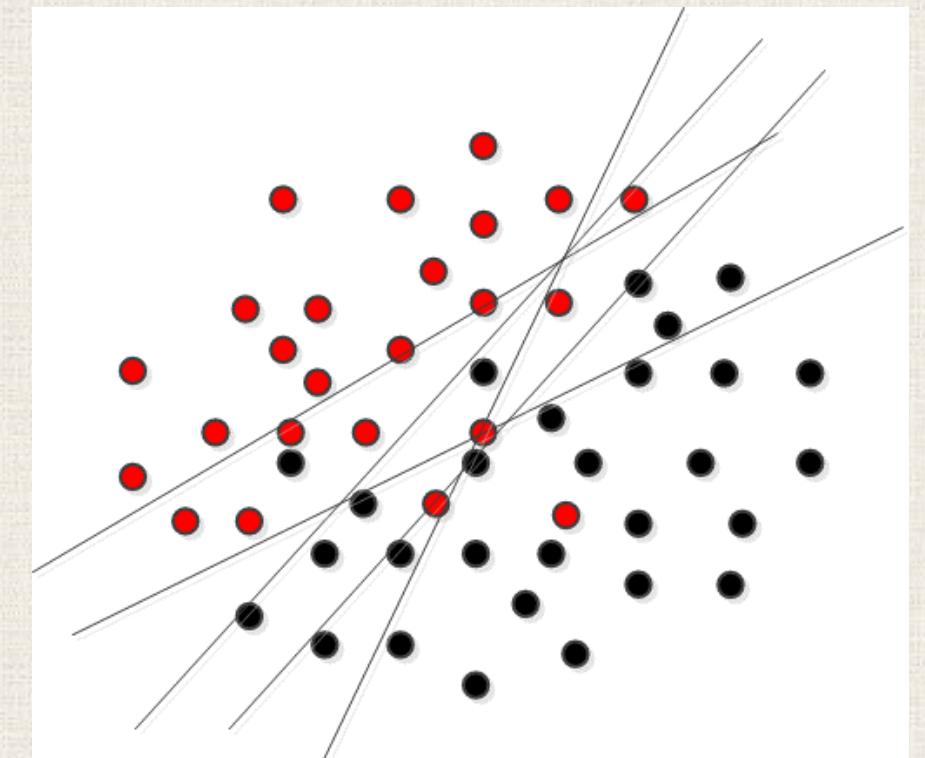


# Non-linearly separable Behavior

When apply the Perceptron model to a not linearly separable data set, what is going to happen?

- No guarantee of the convergence
- No guarantee of the performance of the trained model

**What should we do?**



# The Voted and Averaged Perceptron

**Training:**

*Input:* a labeled training set  $\{(x_1, y_1), \dots, (x_m, y_m)\}$

number of epochs  $T$

*Output:* a list of weighted Perceptrons  $\{(\mathbf{w}_1, c_1), \dots, (\mathbf{w}_k, c_k)\}$

Initialize:  $k \leftarrow 0, \mathbf{w}_1 = 0, c_1 = 0$

**Repeat**  $T$  times: { %training stops after  $T$  epochs

**for**  $i = 1, \dots, m$

        compute prediction  $\hat{y} = \text{sign}(\mathbf{w}_k^T \mathbf{x}_i)$

**if**  $y_i == \hat{y}$

$c_k = c_k + 1$  (times of “survive”)

**else**

$\mathbf{w}_{k+1} = \mathbf{w}_k + \mathbf{x}_i y_i$

$c_{k+1} = 1$

$k = k + 1$

}

**Return**  $\{(\mathbf{w}_1, c_1), \dots, (\mathbf{w}_k, c_k)\}$

**Test:**

*Given:* a list of weighted Perceptrons  $\{(\mathbf{w}_1, c_1), \dots, (\mathbf{w}_k, c_k)\}$   
an unseen test instance  $\mathbf{x}$

The **voted perceptron** computes a predicted label  $\hat{y}$  as follows:

$$s = \sum_{i=1}^k c_i \text{sign}(\mathbf{w}_i^T \mathbf{x}); \quad \hat{y} = \text{sign}(s)$$

The **averaged perceptron** computes a predicted label  $\hat{y}$  as:

$$\hat{y} = \text{sign} \left( \sum_{i=1}^k c_i (\mathbf{w}_i^T \mathbf{x}) \right)$$

# **How to enhance the Perceptron model to handle non-linearly separable data sets?**

Can nonlinear feature mapping be used in Perceptron model?

# Logistic Regression

A discriminative classification model

Let's consider the ***binary case***: no need to be linearly separable

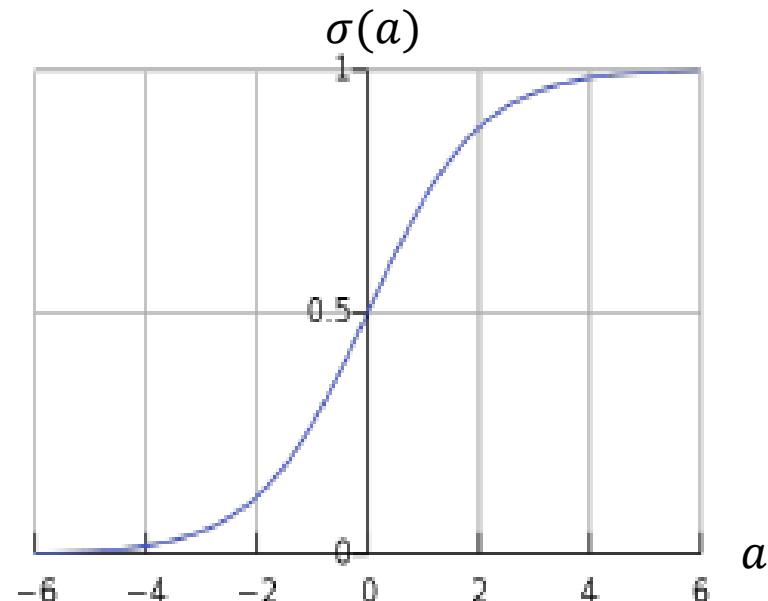
- Given a training data set  $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^N$ , where  $y_i \in \{0,1\}, i = 1, \dots, N$ .
- *Logistic regression method directly construct the posterior probability that a data point is in class  $C_1$  or  $C_2$ .*

- Logistic regression method directly construct the **posterior probability** that a data point is in class  $C_1$  or  $C_2$ .
- It converts the raw linear “score”  $\mathbf{w}^T \mathbf{x}$  to a probability between 0 and 1 by applying a **sigmoid (logistic)** function  $\sigma(\mathbf{w}^T \mathbf{x})$ , where,

$$\sigma(a) = \frac{1}{1 + e^{-a}}$$

- A nice property of the **sigmoid** function

$$\frac{d\sigma(a)}{da} = \sigma(a)(1 - \sigma(a))$$



- To classify an arbitrary point  $x$ , we use the **sigmoid (logistic)** function to output a posterior probability over the class  $C_1$  and  $C_2$ :

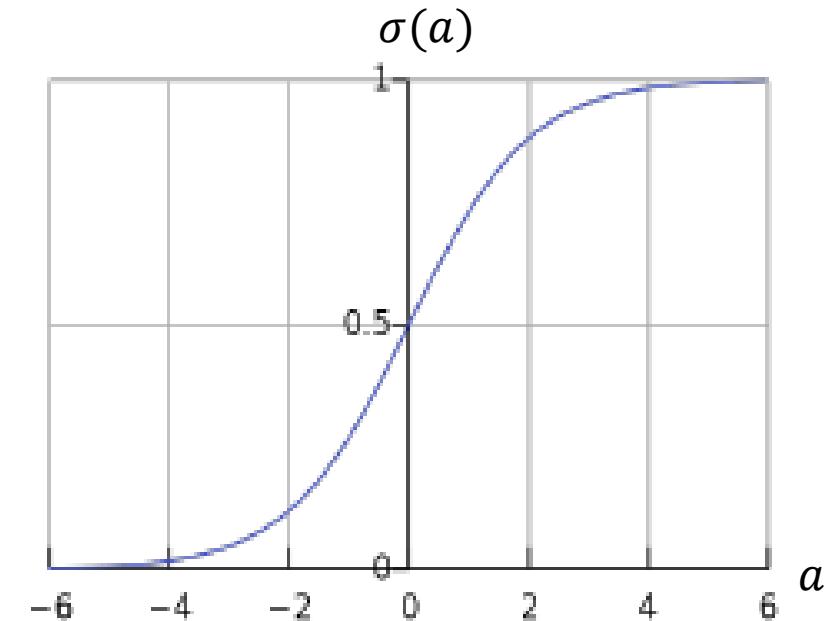
$$P(C_1|x) = \sigma(\mathbf{w}^T \mathbf{x}), \quad P(C_2|x) = 1 - \sigma(\mathbf{w}^T \mathbf{x})$$

- We classify  $x$  as the class with the maximum **posterior probability**:

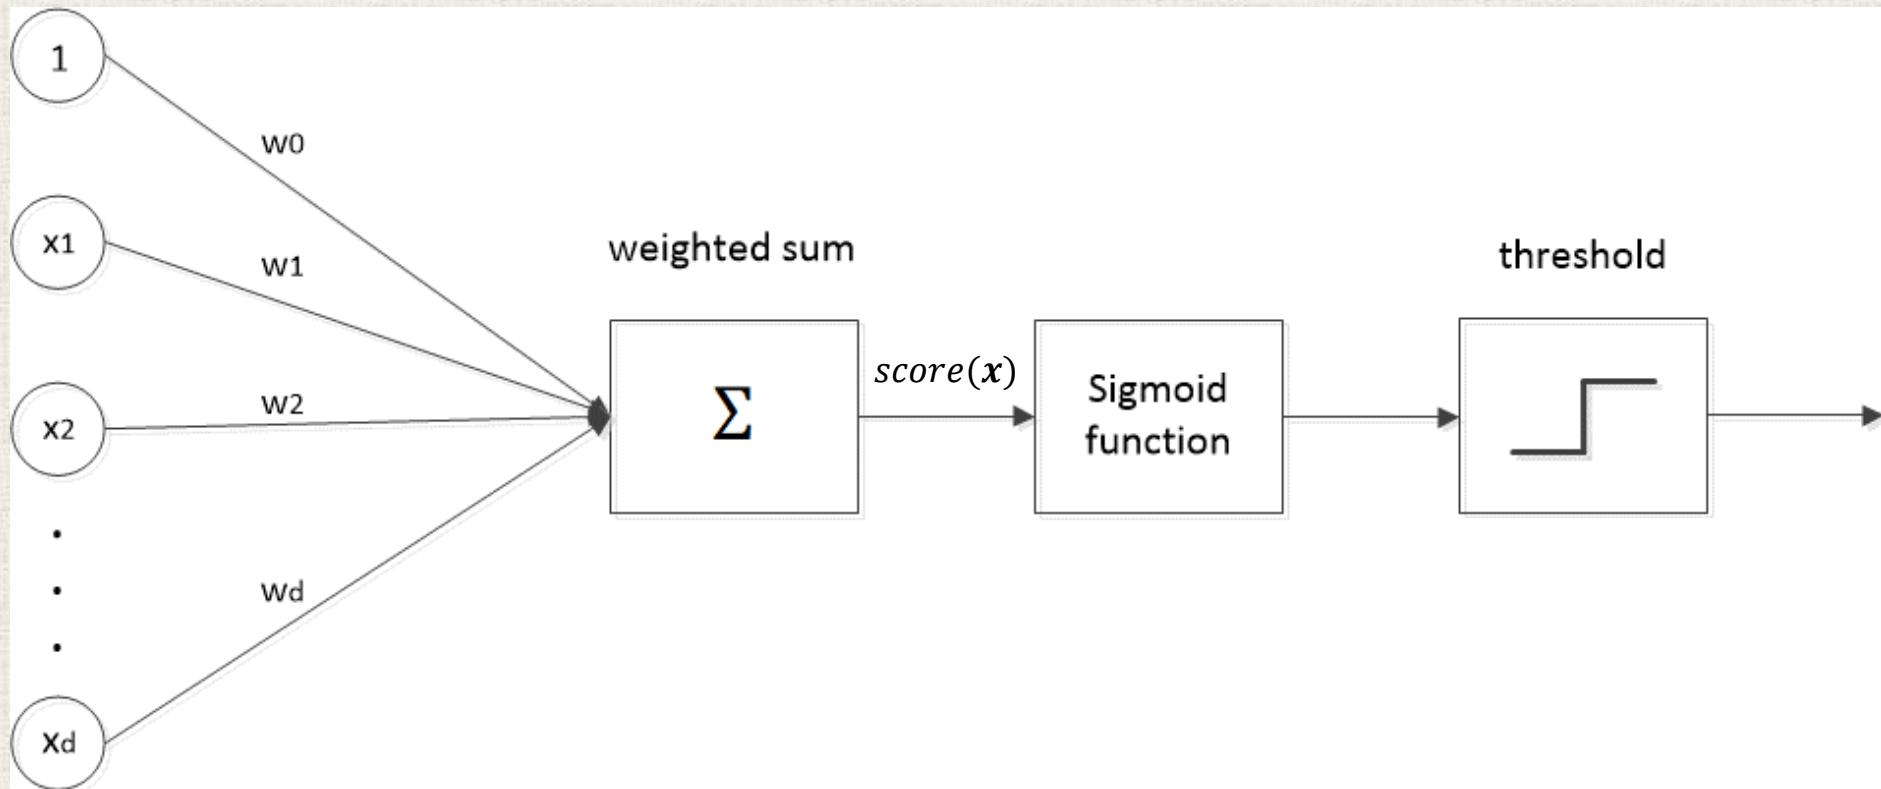
$$\hat{y} = \max_{k=1,2} P(C_k|x) = \begin{cases} 1, & \text{if } \sigma(\mathbf{w}^T \mathbf{x}) \geq 0.5 \\ 0, & \text{otherwise} \end{cases}$$

- Equivalently, we classify  $x$  as (the decision rule):

$$\hat{y} = \begin{cases} 1, & \text{if } \mathbf{w}^T \mathbf{x} \geq 0 \\ 0, & \text{otherwise} \end{cases}$$



# Graphical model of Binary Logistic Regression



## Logistic Regression - A discriminative method

- Perceptron and Logistic regression are both discriminative method
- Perceptron directly work on the discriminant function to find the decision boundary
- Logistic Regression model use posterior probability  $P(C_1|x)$  to make classification decision
- Logistic Regression does not calculate the posterior probability using likelihood and prior like in Bayes Theorem. Instead, it constructs the posterior probability by learning from the training data.
- Logistic Regression is a ***discriminative method***

$$P(C_1|x) = \sigma(\mathbf{w}^T \mathbf{x}), \quad P(C_2|x) = 1 - \sigma(\mathbf{w}^T \mathbf{x}), \quad \sigma(a) = \frac{1}{1+e^{-a}}$$

## **The error (loss) function: (with a maximum likelihood explanation)**

Given a training data set  $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^N$ , where  $y_i \in \{0,1\}$ ,  $i = 1, \dots, N$ . We view each target value  $y_i$  as an independent sample from a **Bernoulli distribution**  $Y_i \sim \text{Bern}(p_i)$ , where  $p_i$  is a function of  $x_i$ . Thus, we have the following probability:

$$P(Y_i = y_i) = \begin{cases} p_i, & \text{if } y_i = 1 \\ 1 - p_i, & \text{if } y_i = 0 \end{cases}$$

Which can also be written as follows:

$$P(Y_i = y_i) = p_i^{y_i} (1 - p_i)^{1-y_i}$$

Then the ***likelihood function*** can be written as: the probability that the specific sample (the training data set) happens given the hypothesis model (represented by the weight vector  $\mathbf{w}$ ).

$$p(\mathbf{y}|\mathbf{w}) = P[(Y_1 = y_1) \text{ and } (Y_2 = y_2) \text{ and } \dots \text{ and } (Y_N = y_N)] = \prod_{i=1}^N p_i^{y_i} (1 - p_i)^{1-y_i}$$

Where,  $\mathbf{y} = (y_1, \dots, y_N)^T$  and  $p_i = P(C_1 | \mathbf{x}_i) = \sigma(\mathbf{w}^T \mathbf{x}_i)$ .

- ***The error function:***

We can define an error function by taking the **negative** logarithm of the likelihood, which gives the ***cross-entropy*** error function in the form:

$$E(\mathbf{w}) = -\ln p(\mathbf{y}|\mathbf{w}) = -\sum_{i=1}^N \{y_i \ln p_i + (1 - y_i) \ln(1 - p_i)\}$$

- ***This is the maximum likelihood explanation of the error function!***
- ***Minimizing  $E(\mathbf{w})$  is equivalent to Maximizing the likelihood function  $p(\mathbf{y}|\mathbf{w})$***

## Logistic Regression – the training method

The error function:

$$E(\mathbf{w}) = -\ln p(\mathbf{y}|\mathbf{w}) = -\sum_{i=1}^N \{y_i \ln p_i + (1 - y_i) \ln(1 - p_i)\}$$

Where,

$$p_i = P(C_1|x_i) = \sigma(\mathbf{w}^T \mathbf{x}_i) = \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}_i}}$$

The gradient is:

$$\nabla_{\mathbf{w}} E(\mathbf{w}) = \nabla_{\mathbf{w}} \left( - \sum_{i=1}^N \{y_i \ln p_i + (1 - y_i) \ln(1 - p_i)\} \right)$$

$$= - \sum_{i=1}^N \{y_i \nabla_{\mathbf{w}} \ln p_i + (1 - y_i) \nabla_{\mathbf{w}} \ln(1 - p_i)\}$$

$$= - \sum_{i=1}^N \left\{ \frac{y_i}{p_i} \nabla_{\mathbf{w}} p_i - \frac{(1 - y_i)}{1 - p_i} \nabla_{\mathbf{w}} p_i \right\} = - \sum_{i=1}^N \left\{ \frac{y_i}{p_i} - \frac{(1 - y_i)}{1 - p_i} \right\} \nabla_{\mathbf{w}} p_i$$

Let's find out  $\nabla_{\mathbf{w}} p_i$ , with  $p_i = P(C_1 | \mathbf{x}_i) = \sigma(\mathbf{w}^T \mathbf{x}_i) = \frac{1}{1+e^{-\mathbf{w}^T \mathbf{x}_i}}$ .

Note that  $\nabla_a \sigma(a) = \sigma(a)(1 - \sigma(a))$ , and from the chain rule of derivative, we have,

$$\begin{aligned}\nabla_{\mathbf{w}} p_i &= \nabla_{\mathbf{w}} \sigma(\mathbf{w}^T \mathbf{x}_i) \\ &= \sigma(\mathbf{w}^T \mathbf{x}_i)(1 - \sigma(\mathbf{w}^T \mathbf{x}_i)) \mathbf{x}_i \\ &= p_i(1 - p_i) \mathbf{x}_i\end{aligned}$$

Hence, we have,

$$\begin{aligned}\nabla_{\mathbf{w}} E(\mathbf{w}) &= - \sum_{i=1}^N \left\{ \frac{y_i}{p_i} - \frac{(1-y_i)}{1-p_i} \right\} p_i(1-p_i) \mathbf{x}_i \\ &= - \sum_{i=1}^N \{y_i(1-p_i) - (1-y_i)p_i\} \mathbf{x}_i \\ &= - \sum_{i=1}^N (y_i - p_i) \mathbf{x}_i\end{aligned}$$

The gradient descent update is:

$$\mathbf{w}_{k+1} = \mathbf{w}_k - \eta \nabla_{\mathbf{w}} E(\mathbf{w}_k) = \mathbf{w}_k + \eta \sum_{i=1}^N (y_i - p_i) \mathbf{x}_i$$

Where,  $p_i = \sigma(\mathbf{w}_k^T \mathbf{x}_i) = \frac{1}{1+e^{-\mathbf{w}_k^T \mathbf{x}_i}}$

**Q1:** What are the differences between Perceptron training and Logistic regression training algorithms?

**Q2:** Do we have batch training and stochastic training for Logistic regression?

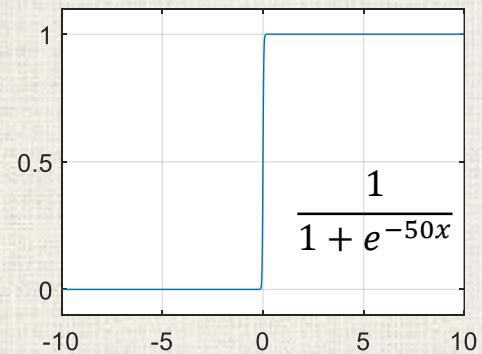
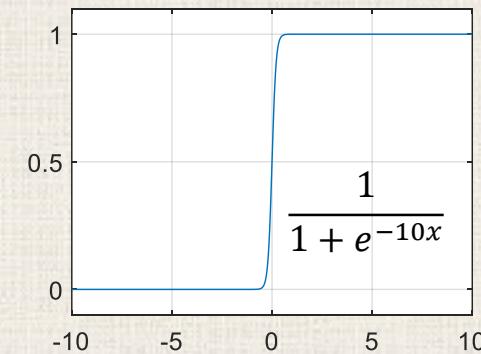
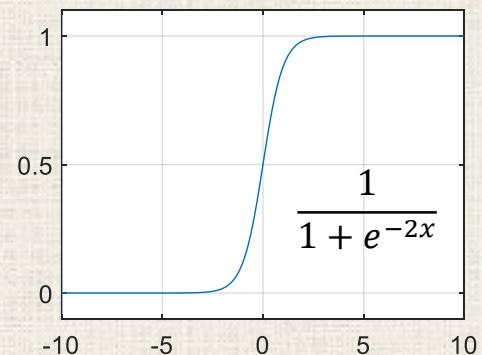
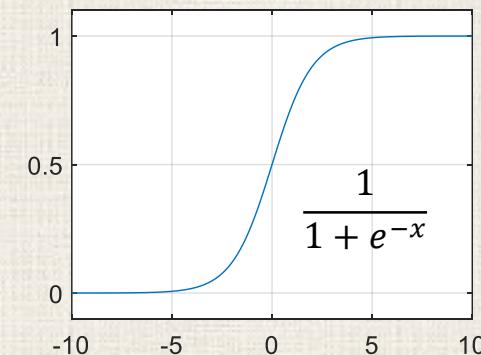
**Q3:** Can we use feature mapping  $\phi(\mathbf{x})$  in Logistic Regression?

# Regularized Logistic Regression

- ***When data set is linearly separable, logistic regression weights go to infinite.***

when data set is linearly separable, the conditional likelihood function is perfectly 0 and 1, the error function  $E(\mathbf{w})$  goes to zero.

This is corresponding to the steepest sigmoid function  $\frac{1}{1+e^{\mathbf{w}^T \mathbf{x}}}$  where the weights go to infinite.



# Regularized Logistic Regression

- *Also, using feature mapping  $\phi(x)$  may cause overfitting*
- Large number of features in the model trained using a small data set may cause overfitting
- To avoid overfitting, a regularization penalty term can be included into the error function:

$$E(\mathbf{w}) = -\ln p(\mathbf{y}|\mathbf{w}) + \lambda \|\mathbf{w}\|_2^2 = -\sum_{i=1}^N \{y_i \ln p_i + (1 - y_i) \ln(1 - p_i)\} + \lambda \|\mathbf{w}\|_2^2$$

# Regularized Logistic Regression

This new error function is equivalent to:

$$E(\mathbf{w}) = 1/\lambda \left\{ - \sum_{i=1}^N \{y_i \ln p_i + (1 - y_i) \ln(1 - p_i)\} \right\} + \|\mathbf{w}\|_2^2$$

$$E(\mathbf{w}) = C \left\{ - \sum_{i=1}^N \{y_i \ln p_i + (1 - y_i) \ln(1 - p_i)\} \right\} + \|\mathbf{w}\|_2^2$$

Where,  $C = 1/\lambda$  is the inverse of the regularization strength (*sk-learn use regularized LR by default*).

# Multinomial Logistic Regression

## *Multiclass Logistic Regression:*

- Let's generalize logistic regression to the case where there are  $K$  classes
  - The training data set  $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^N$
  - The one-hot vector encoding:
- If the  $i^{th}$  observation belongs to class  $C_k$ , instead of  $y_i = k$ , we use the representation  $\mathbf{y}_i = \mathbf{e}_k$ , the  $k^{th}$  canonical basis vector. For example, if an observation belong to the class  $C_2$ , its label representation would be: ( $K = 3, k = 2$ )

$$\mathbf{y}_i = \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}$$

Now, the target value becomes a  $K$  dimensional vector.

### **Multiclass Logistic Regression:**

- Instead of a weight vector, now we have a weight matrix to be determined:

$$\mathbf{W} = \begin{bmatrix} \mathbf{w}_1^T \\ \vdots \\ \mathbf{w}_K^T \end{bmatrix}$$

Where  $\mathbf{w}_k \in R^{d+1}$  is the weight vector for the  $k^{th}$  class.

- For each input  $\mathbf{x}_i \in R^{d+1}$ , each class  $k$  is given a “score”  $z_k = \mathbf{w}_k^T \mathbf{x}_i$
- In total there are  $K$  raw linear scores for an arbitrary input  $\mathbf{x}$ :

$$\mathbf{Wx} = \begin{bmatrix} \mathbf{w}_1^T \mathbf{x} \\ \vdots \\ \mathbf{w}_K^T \mathbf{x} \end{bmatrix} = \begin{bmatrix} z_1 \\ \vdots \\ z_K \end{bmatrix}$$

# Multinomial Logistic Regression

**Multiclass Logistic Regression:**

$$W\mathbf{x} = \begin{bmatrix} \mathbf{w}_1^T \mathbf{x} \\ \vdots \\ \mathbf{w}_K^T \mathbf{x} \end{bmatrix} = \begin{bmatrix} z_1 \\ \vdots \\ z_K \end{bmatrix}$$

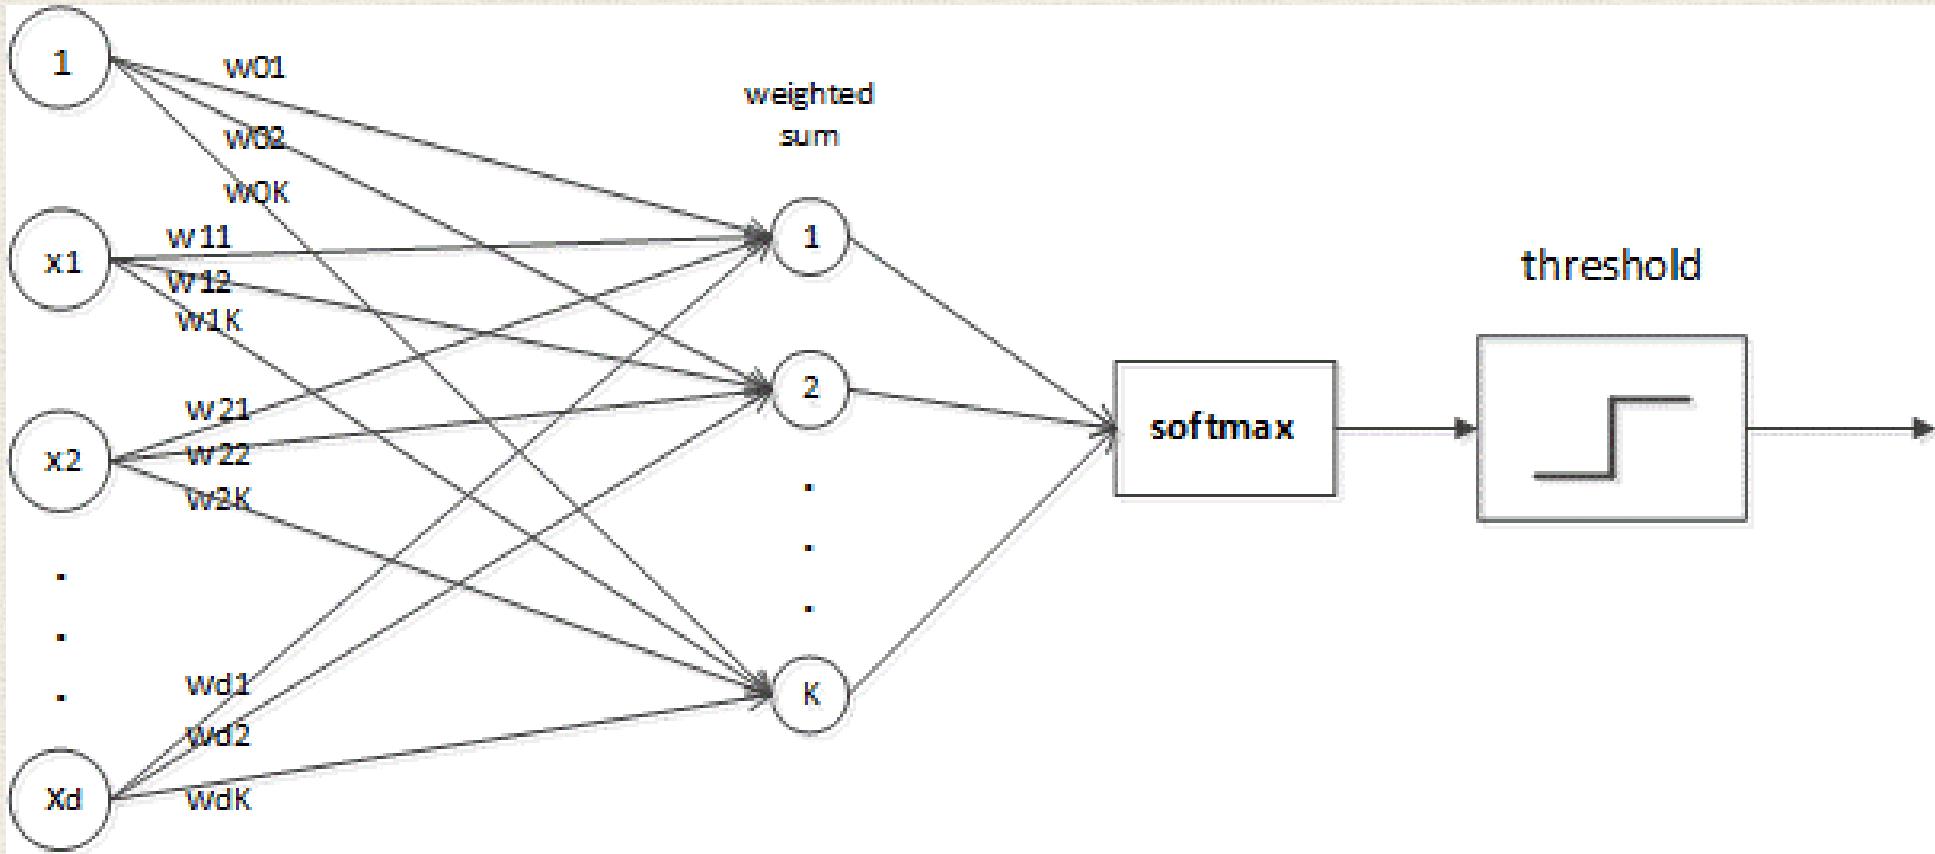
- We now must transform all of these scores into a posterior probability distribution  $P(C_k|\mathbf{x})$ . For the multiclass case, We use **softmax** function:

$$P(C_k|\mathbf{x}) = \sigma(z_k) = \frac{e^{z_k}}{\sum_{j=1}^K e^{z_j}} = \frac{e^{\mathbf{w}_k^T \mathbf{x}}}{\sum_{j=1}^K e^{\mathbf{w}_j^T \mathbf{x}}}, \quad k = 1, 2, \dots, K$$

- We then predict the class with the maximum probability:

$$\hat{y} = \max_k P(C_k|\mathbf{x})$$

# Graphical model of multinomial Logistic Regression



### **The error function for multi-class logistic regression:**

- The *likelihood function* is given by:

$$p(\mathbf{y}|\mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_k) = \prod_{i=1}^N \prod_{k=1}^K [P(C_k | \mathbf{x}_i)]^{y_{ik}} = \prod_{i=1}^N \prod_{k=1}^K [p_{ik}]^{y_{ik}}$$

Where  $\mathbf{y}$  is an  $K \times N$  matrix of target values with elements  $y_{ik}$  and

$$p_{ik} = P(C_k | \mathbf{x}_i) = \frac{e^{\mathbf{w}_k^T \mathbf{x}_i}}{\sum_{j=1}^K e^{\mathbf{w}_j^T \mathbf{x}_i}}$$

- Taking the negative logarithm then gives:

$$E(\mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_K) = -\ln p(\mathbf{y}|\mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_k) = -\sum_{i=1}^N \sum_{k=1}^K (y_{ik} \ln p_{ik}) = -\sum_{i=1}^N \sum_{k=1}^K \left( y_{ik} \ln \left( \frac{e^{\mathbf{w}_k^T \mathbf{x}_i}}{\sum_{j=1}^K e^{\mathbf{w}_j^T \mathbf{x}_i}} \right) \right)$$

Which is known as the ***cross-entropy*** error function for the multiclass logistic regression problem.

- Instead of finding the gradient with respect to all of the parameters of the matrix  $\mathbf{W}$ , let's find the gradient with respect to one row of  $\mathbf{W}$  at a time:

$$\begin{aligned}
 \nabla_{\mathbf{w}_l} E(\mathbf{w}_1, \dots, \mathbf{w}_K) &= \nabla_{\mathbf{w}_l} \left\{ - \sum_{i=1}^N \sum_{k=1}^K \left( y_{ik} \ln \left( \frac{e^{\mathbf{w}_k^T \mathbf{x}_i}}{\sum_{j=1}^K e^{\mathbf{w}_j^T \mathbf{x}_i}} \right) \right) \right\} \\
 &= - \sum_{i=1}^N \sum_{k=1}^K \left( y_{ik} \nabla_{\mathbf{w}_l} \left\{ \ln \left( \frac{e^{\mathbf{w}_k^T \mathbf{x}_i}}{\sum_{j=1}^K e^{\mathbf{w}_j^T \mathbf{x}_i}} \right) \right\} \right) \\
 &= - \sum_{i=1}^N \sum_{k=1}^K \left( y_{ik} \nabla_{\mathbf{w}_l} \left\{ \ln \left( \frac{e^{\mathbf{w}_k^T \mathbf{x}_i}}{\sum_{j=1}^K e^{\mathbf{w}_j^T \mathbf{x}_i}} \right) \right\} \right) \\
 &= - \sum_{i=1}^N (y_{il} - p_{il}) \mathbf{x}_i
 \end{aligned}$$

## Multinomial Logistic Regression

- The gradient descent update for  $\mathbf{w}_l$  is then, ( $l = 1, \dots, K$ )

$$\mathbf{w}_l^{t+1} = \mathbf{w}_l^t + \eta \sum_{i=1}^N (y_{il} - p_{il}) \mathbf{x}_i$$