






# **K - Nearest Neighbor Method – An Instance based method**






# The Nearest Neighbor Method



- ***Observation:*** similar examples tend to belong to the same class.
- ***Idea:*** assign a test input with the class label from the example in the training data set that is the most similar to the test input.
- ***Implementation:*** calculate the distances between the test input and the examples in the data set and find the one with the shortest distance.

# The K - Nearest Neighbor Method



- ***Observation:*** similar examples tend to belong to the same class.
- ***Idea:*** assign a test input with the dominating class label from the examples in the training data set that are within the neighborhood of the test input.
- ***Implementation:*** calculate the distances between the test input and the examples in the training data set and find the K ones that are the most closest to the test input and find the dominating class label within these K neighboring examples.


## **Histogram** density model

In the case of a single continuous variable  $x$ . Given  $N$  sample values from  $x$ .

- Partition  $x$  into distinct bins of width  $\Delta_i$
- Count the number of observations of  $x$  falling in bin  $i$  as  $n_i$
- Probability density value for bin  $i$  is given by

$$p_i = \frac{n_i}{N\Delta_i}$$

Hence, we have,  $\sum_{i=1}^m p_i \Delta_i = 1$



- Let's suppose we have a data set comprising  $N_k$  points in class  $\mathcal{C}_k$  with  $N$  points in total so that  $\sum_{k=1}^C N_k = N$ .
- If we wish to classify a new point  $x$ , we draw a sphere centered on  $x$  containing precisely  $K$  points irrespective of their class. Suppose this sphere has volume  $V$  and contains  $K_k$  points from class  $\mathcal{C}_k$ .
- We want to classify the new point  $x$  to class  $\mathcal{C}_k$  that has the highest ***posterior probability of class membership:***

$$p(\mathcal{C}_k|x) = \frac{p(x|\mathcal{C}_k)p(\mathcal{C}_k)}{p(x)}$$

- Then the likelihood can be estimated as:

$$p(x|\mathcal{C}_k) = \frac{K_k}{N_k V}$$

- The unconditional density can be estimated as:

$$p(x) = \frac{K}{NV}$$

- The class prior can be calculated as:

$$p(\mathcal{C}_k) = \frac{N_k}{N}$$

- Applying ***Bayes' theorem*** gives the ***posterior probability of class membership***:

$$p(\mathcal{C}_k | \mathbf{x}) = \frac{p(\mathbf{x} | \mathcal{C}_k) p(\mathcal{C}_k)}{p(\mathbf{x})} = \frac{K_k}{K}$$

- To minimize the probability of misclassification, we assign the test point  $\mathbf{x}$  to the class having the largest posterior probability, i.e., the largest value of  $K_k/K$
- ***To classify a new point, we identify the K nearest points from the training data set and then assign the new point to the class having the greatest number of representatives amongst the K nearest neighbor set.***

# The KNN Classification Algorithm

**Begin**

load the training data set  $\{x_i, y_i\}$ , and the test data point  $x$

Choose the value of  $K$ , set  $\mathcal{D} = \{\}$

**For**  $i = 1$  to  $N$

    calculate  $d_i = \|x - x_i\|$

    append  $\{y_i, d_i\}$  to  $\mathcal{D}$

**End**


sort  $\mathcal{D}$  in the ascend order with respect to  $d$  value

pick the first  $K$  entries  $\{y_n, d_n\}, n = 1, \dots, K$  from  $\mathcal{D}$

$y = mode(\{y_n, d_n\}, n = 1, \dots, K)$

**Return**  $y$

**End**



# **Metrics used in the KNN method**

- ***Distance functions*** are used in KNN method as measure of similarity of instances
- Commonly used distance functions in machine learning include the following:

Let  $\mathbf{x} = [x_1 \quad \cdots \quad x_m]^T$  and  $\mathbf{y} = [y_1 \quad \cdots \quad y_m]^T$  be two vectors, then the distances between  $\mathbf{x}$  and  $\mathbf{y}$  are:

- *Euclidean distance*

$$d(\mathbf{x}, \mathbf{y}) = \left( \sum_{i=1}^m (x_i - y_i)^2 \right)^{\frac{1}{2}}$$

- *Minkowsky distance*

$$d(\mathbf{x}, \mathbf{y}) = \left( \sum_{i=1}^m |x_i - y_i|^p \right)^{\frac{1}{p}}$$

- *Manhattan* distance

$$d(\mathbf{x}, \mathbf{y}) = \sum_{i=1}^m |x_i - y_i|$$

- *Chebyshev* distance

$$d(\mathbf{x}, \mathbf{y}) = \max_i |x_i - y_i|$$

- *Hamming* distance:

$$d_i = \begin{cases} 1, & \text{if } x_i \neq y_i \\ 0, & \text{otherwise} \end{cases}; \quad d(\mathbf{x}, \mathbf{y}) = \sum_{i=1}^m d_i$$

Let  $\mathbf{x} = [x_1 \quad \cdots \quad x_m]^T$  and  $\mathbf{y} = [y_1 \quad \cdots \quad y_m]^T$  be two vectors, we have,

- Cosine similarity measure:

$$sim(\mathbf{x}, \mathbf{y}) = \frac{\langle \mathbf{x}, \mathbf{y} \rangle}{\|\mathbf{x}\| \|\mathbf{y}\|} = \frac{\sum_{i=1}^m x_i y_i}{\sqrt{\sum_{i=1}^m x_i^2} \sqrt{\sum_{i=1}^m y_i^2}}$$

- Correlation:

$$corr(\mathbf{x}, \mathbf{y}) = \frac{\sum_{i=1}^m (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum (x_i - \bar{x})^2} \sqrt{\sum (y_i - \bar{y})^2}} = \frac{\langle \mathbf{x} - \bar{\mathbf{x}}, \mathbf{y} - \bar{\mathbf{y}} \rangle}{\|\mathbf{x} - \bar{\mathbf{x}}\| \|\mathbf{y} - \bar{\mathbf{y}}\|} = sim(\mathbf{x} - \bar{\mathbf{x}}, \mathbf{y} - \bar{\mathbf{y}})$$

## **Questions about the KNN method**

**Q1:** How do we train a KNN algorithm?

**Q2:** Does a KNN algorithm need a training set/test set?

**Q3:** Can the KNN method be applied to the regression problem?

**Q4:** Is KNN a good candidate for online (real-time) classification?


# **The curse of dimensionality**

- The loss of locality in *higher dimensions*

- Consider  $N$  data points uniformly distributed in a  $d$  dimensional unit hypercube.
- Let's find out a hypercube of edge  $l$  that contains  $K$  nearest neighbors of a test point. That is, we want to find a cube of edge  $l$  to contain a fraction  $r = \frac{K}{N}$  of the entire data set.
- Since the data points are uniformly distributed, this means the volume of the cube is fraction  $r$  of the unit cube, i.e.,  

$$l^d = r \Rightarrow l = r^{\frac{1}{d}}$$
- When  $d \rightarrow \infty$ , we have  $l \rightarrow 1$ . i.e., the hypercube is close to the entire unit hypercube where the data is uniformly distributed!!! This is true even for very small  $r$  value.


$$0 < l < 1$$



For example, let  $N = 1000, K = 10$ , we have,  $r = 0.01$

$$l^d = r \Rightarrow l = r^{\frac{1}{d}}$$

$d$	$l$
1	0.01
2	0.1
3	0.21
10	0.63
100	0.96
1000	0.995



To find the 10 – nearest neighbor, we need almost the entire input space in high dimensions!!!


- To find the 10 – nearest neighbor, we need almost the entire input space in high dimensions!!!
- This breaks down the KNN assumptions because in high dimensions, the points in  $K$  nearest neighbor are not particularly closer (similar) than any other data points in the training set.
- Can we increase the number of training data points until the  $K$  nearest neighbors are truly closer to the test point?
- How many data points do we need such that  $l$ , the size of the neighborhood become truly small?

For example,

$$\text{let } l = 0.1. \text{ then, } r = \frac{K}{N} = l^d \Rightarrow N = \frac{K}{l^d} = K \times 10^d$$

When  $d = 10$ ,  $N$  becomes a huge number!!!

Dimension reduction (PCA) may provide a solution to alleviate the problem!



# **The Naïve Bayes Classifier – A generative approach**

- The training data set is given as  $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^N$ , where, an instance  $x_i$  belongs to class  $\mathcal{C}_k$  if  $y_i = k$ ,  $k = 1, \dots, K$ .
- Given a new instance,  $x = [x_1 \quad x_2 \quad \cdots \quad x_d]^T$ , we wish to predict the class of this instance.
- The ***Bayesian approach*** to classify the new instance is to assign the most probable target value,  $y_{MAP}$ , given the attribute values  $x_1, \dots, x_d$ . That is,

$$y_{MAP} = \underset{k}{\operatorname{argmax}} P(\mathcal{C}_k | x_1, \dots, x_d)$$

- Using ***Bayes' theorem***, this probability can be written as:

$$\begin{aligned} y_{MAP} &= \underset{k}{\operatorname{argmax}} \frac{P(x_1, \dots, x_d | \mathcal{C}_k) P(\mathcal{C}_k)}{P(x_1, \dots, x_d)} \\ &= \underset{k}{\operatorname{argmax}} P(x_1, \dots, x_d | \mathcal{C}_k) P(\mathcal{C}_k) \end{aligned}$$

- ***Assume the attributes are conditionally independent given the target value***, we have

$$P(x_1, \dots, x_d | \mathcal{C}_k) = P(x_1 | \mathcal{C}_k) P(x_2 | \mathcal{C}_k) \cdots P(x_d | \mathcal{C}_k) = \prod_{j=1}^d P(x_j | \mathcal{C}_k)$$


- Then we have,

$$y_{MAP} = \underset{k}{\operatorname{argmax}} P(\mathcal{C}_k) \prod_{j=1}^d P(x_j | \mathcal{C}_k)$$


- This is why the method is called “***naïve***” Baye’s method

# The Naïve Bayes Classifier – A Toy Example

- The training data set is given as:



- $y_i \in \{+, -\}, i = 1, 2, \dots, 12$  are the class labels.
- $\mathbf{x}_i = [x_{i1} \quad x_{i2}]^T, x_{i1} \in \{\text{blue, green, red, yellow}\}, x_{i2} \in \{\text{square, circle}\}, i = 1, \dots, 12$  are the feature vectors.



- Given the following new input, let's use *naïve Bayes method* to predict its class:



- The feature vector is  $x = [x_1, x_2]^T = [\text{blue}, \text{square}]^T$

- Under the assumption that the training samples are *i.i.d*, we have the class priors as:

$$P(+) = \frac{7}{12} = 0.58; \quad P(-) = \frac{5}{12} = 0.42$$

- Assuming features are conditionally independent for given class**, the class conditional probabilities can be calculated as:

$$P(x|+) = P(x_1 = \text{blue}, x_2 = \text{square}|+) = P(\text{blue}|+)P(\text{square}|+) = \frac{3}{7} \times \frac{5}{7} = 0.31$$

$$P(x|-) = P(x_1 = \text{blue}, x_2 = \text{square}|-) = P(\text{blue}|-)P(\text{square}|-) = \frac{3}{5} \times \frac{3}{5} = 0.36$$


- Then the posterior can be calculated as:

$$P(+|x) = P(x|+)P(+) = 0.58 \times 0.31 = 0.18$$

$$P(-|x) = P(x|-)P(-) = 0.42 \times 0.36 = 0.15$$

- Since  $P(+|x) = 0.18 > P(-|x) = 0.15$ , the sample point  $x = \boxed{\phantom{00}}$  is classified as +.

- Suppose a new sample as the following is given:



- The feature vector is  $\mathbf{x} = [x_1, x_2]^T = [\text{yellow}, \text{square}]^T$ , the color attribute value “yellow” is not present in the training data set. This causes the class conditional probability to become:

$$P(\mathbf{x}|+) = P(x_1 = \text{yellow}, x_2 = \text{square}|+) = P(\text{yellow}|+)P(\text{square}|+) = 0 \times \frac{5}{7} = 0$$

$$P(\mathbf{x}|-) = P(x_1 = \text{yellow}, x_2 = \text{square}|-) = P(\text{yellow}|-)P(\text{square}|-) = 0 \times \frac{3}{5} = 0$$

And the posteriors will hence be all calculated as 0s.

- This problem can be solved by using the **Laplace smoothing (additive smoothing)** technique to calculate the class conditional probability:

$$P(x_i | \mathcal{C}_k) = \frac{N_{x_i, \mathcal{C}_k} + 1}{N_{\mathcal{C}_k} + d}$$

where,  $N_{x_i, \mathcal{C}_k}$  is the number of time feature  $x_i$  appears in training data set from class  $\mathcal{C}_k$ ;

$N_{\mathcal{C}_k}$  is the total count of all features appears in training data set from class  $\mathcal{C}_k$ ;  $d$  is the dimensionality of the feature vector

- Then, using this new technique, the class conditional probability of the new sample  $x$  can be calculated as:

$$P(x|+) = P(x_1 = \text{yellow}, x_2 = \text{square}|+) = P(\text{yellow}|+)P(\text{square}|+)$$

$$= \frac{1}{7+2} \times \frac{5}{7} = 0.08$$

$$P(x|-) = P(x_1 = \text{yellow}, x_2 = \text{square}|-) = P(\text{yellow}|-)P(\text{square}|-)$$


$$= \frac{1}{5+2} \times \frac{3}{5} = 0.09$$

- Note that the Laplace smoothing formula is used to calculate the likelihood of the features that do not present in the training data set.

- The posterior probability of the classes can be calculated as:

$$P(+|x) = P(x|+)P(+) = 0.08 \times 0.58 = 0.0464$$

$$P(-|x) = P(x|-)P(-) = 0.09 \times 0.42 = 0.0378$$

Therefore, the new sample  should be classified as +

# The Naïve Bayes Classifier – Text Documents Classification

- ***Inputs:***

- a document  $x$
- A fixed set of classes  $\{\mathcal{C}_1, \dots, \mathcal{C}_K\}$  e.g., {"spam", "ham"}, {"computer science", "biology", "statistics", "economics", "politics"}
- A training set of  $N$  hand-labeled documents  $\{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$  (training corpus)

- ***Output:***

- The class of the document  $x$

# Preprocessing of text documents

- **Tokenization:** breaks down a text corpus into individual words and removes punctuations

A swimmer likes swimming; thus, he swims.



a	swimmer	likes	swimming	thus	he	swims
---	---------	-------	----------	------	----	-------

# Preprocessing of text documents

- **Remove stop words:** stop words are particularly common in a text corpus and thus considered as rather uninformative (e.g., words such as *so*, *and*, *or*, *the*, etc.) a stop list can be created and used for stop words removal.

A swimmer likes swimming; thus, he swims.



swimmer	likes	swimming	he	swims
---------	-------	----------	----	-------

## The Bag of Words Model

- ***Creation of the vocabulary:*** after preprocessing, we can create the collection of all different words that occur in the training data set and each word is associated with a count of how it occurs. ***This is a set of non-redundant items where the order does not matter (bag of words).***

Let  $D_1$  and  $D_2$  be two documents in a training dataset (corpus):

- $D_1$ : “Each state has its own laws.”
- $D_2$ : “Every country has its own culture.”

Based on these two documents, after proper preprocessing, the vocabulary can be created as:

$$V = \{each: 1, state: 1, has: 2, its: 2, own: 2, laws: 1, every: 1, country: 1, culture: 1\}$$

- **Vectorization:** The vocabulary can then be used to construct the  $d$ -dimensional feature vector for the individual documents.  $d = |V|$  is the number of different words in the vocabulary.
- The following are the ***bag of words*** representation of two sample documents  $D_1$  and  $D_2$ : (word position information does not matter)

$V$	each	state	has	its	own	laws	every	country	culture
$x_{D_1}$	1	1	1	1	1	1	0	0	0
$x_{D_2}$	0	0	1	1	1	0	1	1	1
$\Sigma$	1	1	2	2	2	1	1	1	1

- After transforming documents into numbers, popular supervised learning methods can be applied to do document classification.

# Text Documents Classification using The Naïve Bayes Classifier

- Given a test document  $x = [x_1, x_2, \dots, x_d]^T$

- Prior probability estimate:**

$$P(\mathcal{C}_k) = \frac{\text{\# of documents with } \mathcal{C}_k}{\text{Total \# of documents in the training data set}}, \quad k = 1, 2, \dots, K$$

- Class conditional probability estimate: (Laplace smoothing if necessary)**

$$P(x_i | \mathcal{C}_k) = \frac{N(x_i, \mathcal{C}_k) + 1}{N(\mathcal{C}_k) + |V|}$$

Where,  $N(x_i, \mathcal{C}_k)$  is the total number of occurrences of term  $x_i$  in all documents in the training data set with class  $\mathcal{C}_k$ ;  $N(\mathcal{C}_k)$  is the number of occurrences of all the terms in all documents in the training data set with class  $\mathcal{C}_k$ .

- ***Class conditional probability estimate:***

$$P(x_i|\mathcal{C}_k) = \frac{N(x_i, \mathcal{C}_k) + 1}{N(\mathcal{C}_k) + |V|}$$

Where,  $N(x_i, \mathcal{C}_k)$  is the total number of occurrences of term  $x_i$  in all documents in the training data set with class  $\mathcal{C}_k$ ;  $N(\mathcal{C}_k)$  is the number of occurrences of all the terms in all documents in the training data set with class  $\mathcal{C}_k$ .

- Then, assuming features are conditionally independent given class, we have, (“naïve”)

$$P(\mathbf{x}|\mathcal{C}_k) = P(x_1, x_2, \dots, x_d|\mathcal{C}_k) = \prod_{i=1}^d P(x_i|\mathcal{C}_k), k = 1, \dots, K$$

***The chance of occurrence of value  $x_i$  is not affected by the values of other attributes.***

- ***Calculation of the posterior probability:***

$$P(\mathcal{C}_k|x) \propto P(\mathbf{x}|\mathcal{C}_k)P(\mathcal{C}_k)$$